



HashiCorp Terraform

The Ultimate Beginner's Guide

Bryan Krausen

Why Learn Terraform?



66% of Organizations are Increasing Cloud Spend YoY

Organizations are continuing to migrate workloads to one or more public cloud providers.



79% of Organizations have or are Planning Multi-Cloud Deployments

Management of multi-cloud environments requires cloud-agnostic solutions.



64% of Organizations are Experiencing a Shortage of Skilled Staff

Don't have the staff expertise they need to support their cloud infrastructure strategy.

®Copyright Bryan Krausen – DO NOT DISTRIBUTE

How This Course Will Transform Your Skills

Gain Real Skills, Build Confidence, and Open New Doors in Cloud Infrastructure



Hands-On Learning, Fast Results

Get practical experience with real-world Terraform scenarios from the start. You'll build and manage infrastructure right away, gaining skills you can apply immediately.



Confidence in Infrastructure Automation

By the end, you'll feel confident using Terraform to automate and manage cloud infrastructure, laying a solid foundation for advanced cloud projects.



Unlock Opportunities in Your Cloud Career

Terraform is a critical skill in the cloud ecosystem, and this course will give you the edge to stand out, whether you're looking to advance or break into cloud roles.



👉 HASHICONF 2024

INSTRUCTOR

Bryan Krausen

Welcome! I'm a Premier Instructor, a HashiCorp Ambassador, and a seasoned technologist who is passionate about cloud infrastructure and security. I help organizations realize the potential of HashiCorp's tools.



btk.me/btk



Why Trust Me?



I Consult with VERY Large Organizations

During my (many) years of consulting on HashiCorp, I've designed, implemented, fixed, and managed it all 😊



I'm VERY Involved with the Community

I've spoken at 10+ HashiCorp User Groups (HUGs), 5 x HashiConf conferences, and I am one of three people who has been a HashiCorp Ambassador since the start



I Help Write HashiCorp Certification Exams

I've contributed to EVERY single certification exam that HashiCorp offers and continue to do so...



How Do We Get There?

A Variety of Content to Teach, Show, and Get you
Hands-On Experience



Lectures to Build Understanding



Demos for Real-World Examples



Hands-On Labs to Apply Skills





Course Topics and Timeline





Access All Course Code on GitHub

All code used in this course is published in a GitHub repository, so you can follow along, review, and experiment with every example provided.

github.com/btkrausen/terraform



HashiCorp Terraform Hands-On Labs

Get access to FREE hands-on labs to practice with Vault configurations, settings, and components...all right in your browser using GitHub Codespaces.

github.com/btkrausen/terraform-codespaces



Learning Terraform: Tips for Your Journey



Start Small and Build Gradually

Learning Terraform can feel overwhelming but remember that every small step adds up.



Mistakes Are Part of the Process

Errors are expected and even helpful. Don't be discouraged if things don't work right away. Debugging is a great way to deepen your understanding.



Ask Questions and Seek Community

You're not alone on this journey. Reach out for help when you need it, whether in the Q&A, forums, or with peers.





LET'S GET! STARTED!

LECTURE

Introduction to HashiCorp Terraform





What is Terraform?



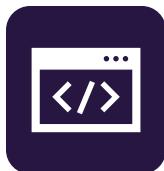
Created by HashiCorp

Often more famous by its tools rather than its name. It also offers many industry-standard solutions



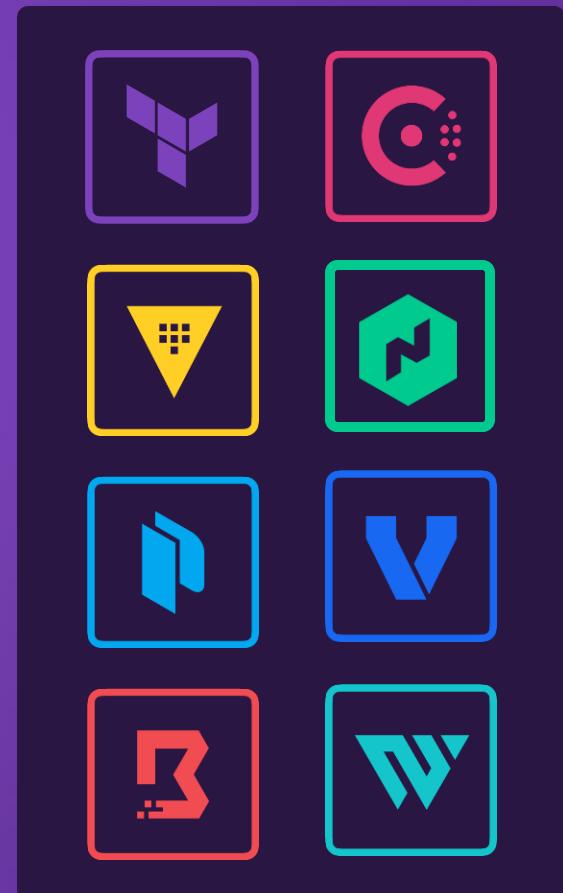
Declarative Approach

Focuses on describing "what" needs to be done rather than explicitly stating "how" to do it, allowing the system to determine the best execution steps



Based on HashiCorp Configuration Language (HCL)

A human-readable language used in Terraform to define infrastructure resources in a clear, structured, and reusable way





What is Terraform?

The industry standard Infrastructure as Code (IaC) tool that automates the provisioning, management, and versioning of infrastructure.

IoIO
IoIO

Codify Your Infrastructure

Rather than clicking around in a console, create a repeatable, versioned workflow to provision and manage resources



Standardize Workflows

Use a single tool to deploy and manage resources required for application workloads. Quickly scale up using modules and other features.



Platform Agnostic

Deploy infrastructure as code using HCL to provision resources from any cloud or datacenter.

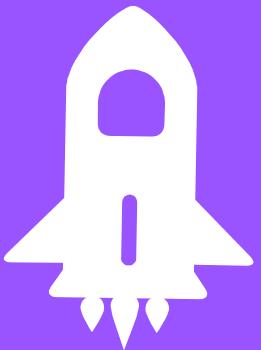


Benefits of Terraform



VERSION CONTROL AND AUDITABILITY

Infrastructure changes are tracked in version control systems, providing a clear history and ability to revert



SCALABILITY & FLEXIBILITY

Quickly scale infrastructure up or down to adapt to changing requirements by modifying code instead of manually changing resources.



AUTOMATION AND EFFICIENCY

Automate the provisioning and management of infrastructure, reducing manual tasks and accelerating deployment times



COLLABORATION

Teams can collaborate on infrastructure changes just like application code, improving transparency and reducing silos in operations



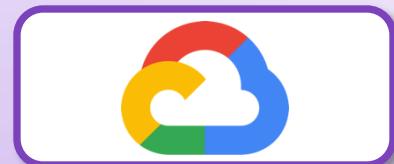
CONSISTENCY & REPEATABILITY

Ensures that infrastructure is consistently deployed across environments, reducing configuration drift and human errors

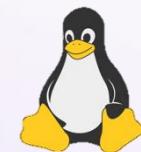
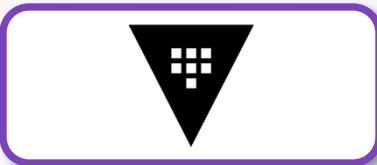


Terraform is Platform Agnostic

Cloud Platforms



Other Services



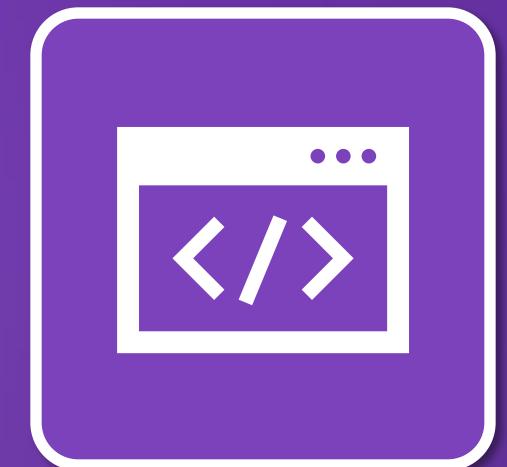
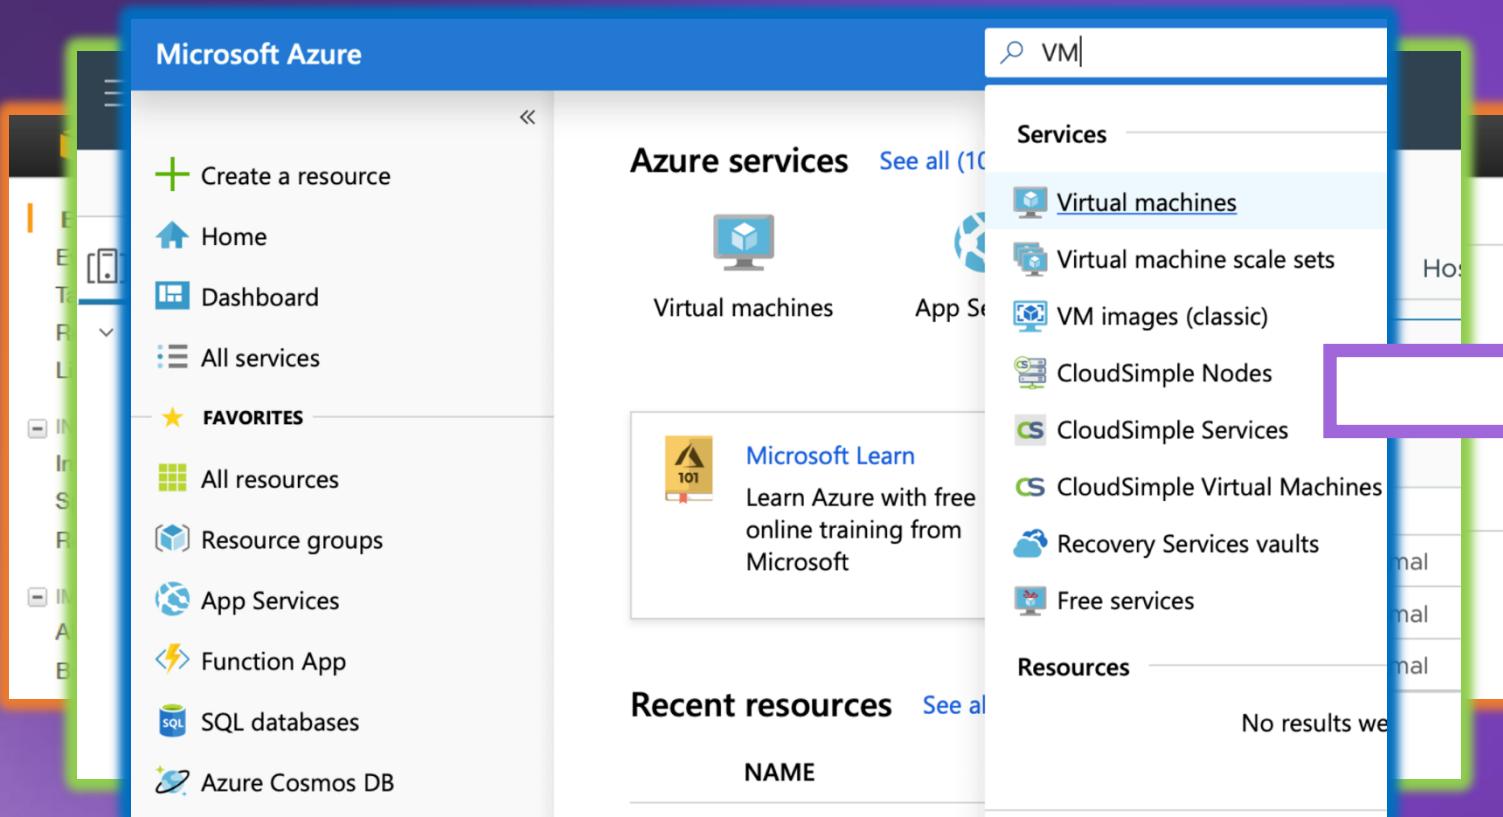


Before Using Terraform, You Need to First Be Proficient on the Target Platform

While it might feel like it, Terraform isn't magic. You need to understand the target platform before using Terraform.



Traditional Process of Provisioning Resources

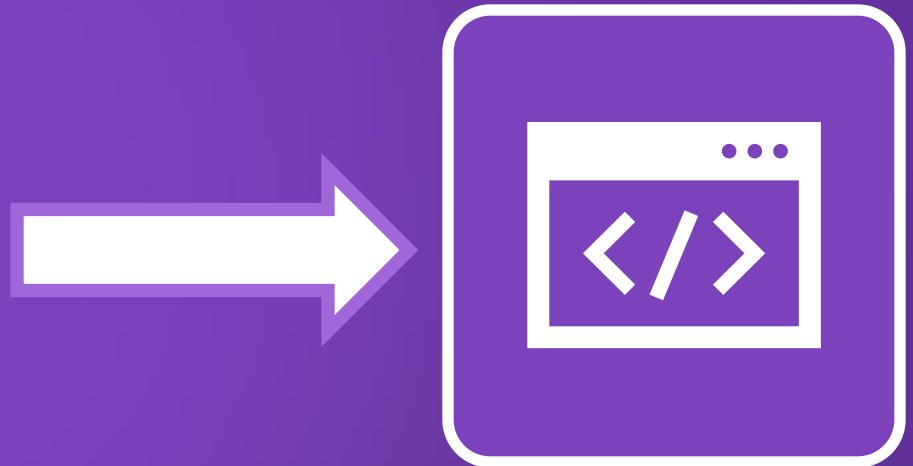


- ✓ Virtual Machine
- ✓ Container
- ✓ Network



Provisioning with Terraform

```
1 resource "aws_instance" "ubuntu_server" {  
2   ami           = data.aws_ami.ubuntu.id  
3   instance_type = "t3.micro"  
4   subnet_id     = aws_subnet.public_subnets["public_subnet_1"].id  
5   security_groups = [aws_security_group.vpc-ping.id]  
6   associate_public_ip_address = true  
7   key_name       = aws_key_pair.generated.key_name  
8   connection {  
9     user      = "ubuntu"  
10    private_key = tls_private_key.generated.private_key_pem  
11    host       = self.public_ip  
12  }
```



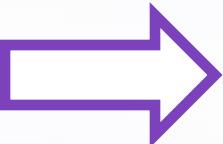
Application
Workload



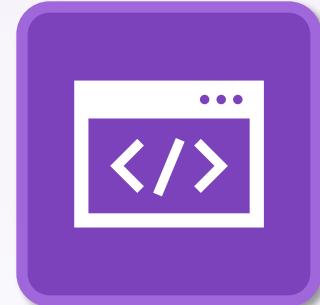
Reusing Code with Terraform

Input Variables

AMI = Ubuntu
Subnet = 10.16.7.0
App = Web Server



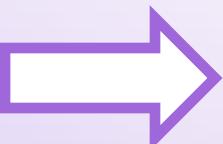
```
1 resource "aws_instance" "basic_server" {  
2   ami                         = data.aws_ami.image.id  
3   instance_type                = var.instance_size  
4   subnet_id                   = aws_subnet.public_subnets.id  
5   security_groups              = [aws_security_group.ingress-ssh.id]  
6   associate_public_ip_address = true  
7   key_name                    = aws_key_pair.generated.key_name  
8 }
```



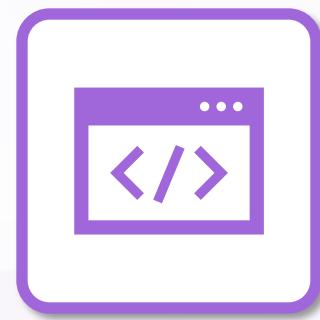
Web Application

Input Variables

AMI = Windows
Subnet = 10.0.3.0
App = Data Processing



```
1 resource "aws_instance" "basic_server" {  
2   ami                         = data.aws_ami.image.id  
3   instance_type                = var.instance_size  
4   subnet_id                   = aws_subnet.public_subnets.id  
5   security_groups              = [aws_security_group.ingress-ssh.id]  
6   associate_public_ip_address = true  
7   key_name                    = aws_key_pair.generated.key_name  
8 }
```

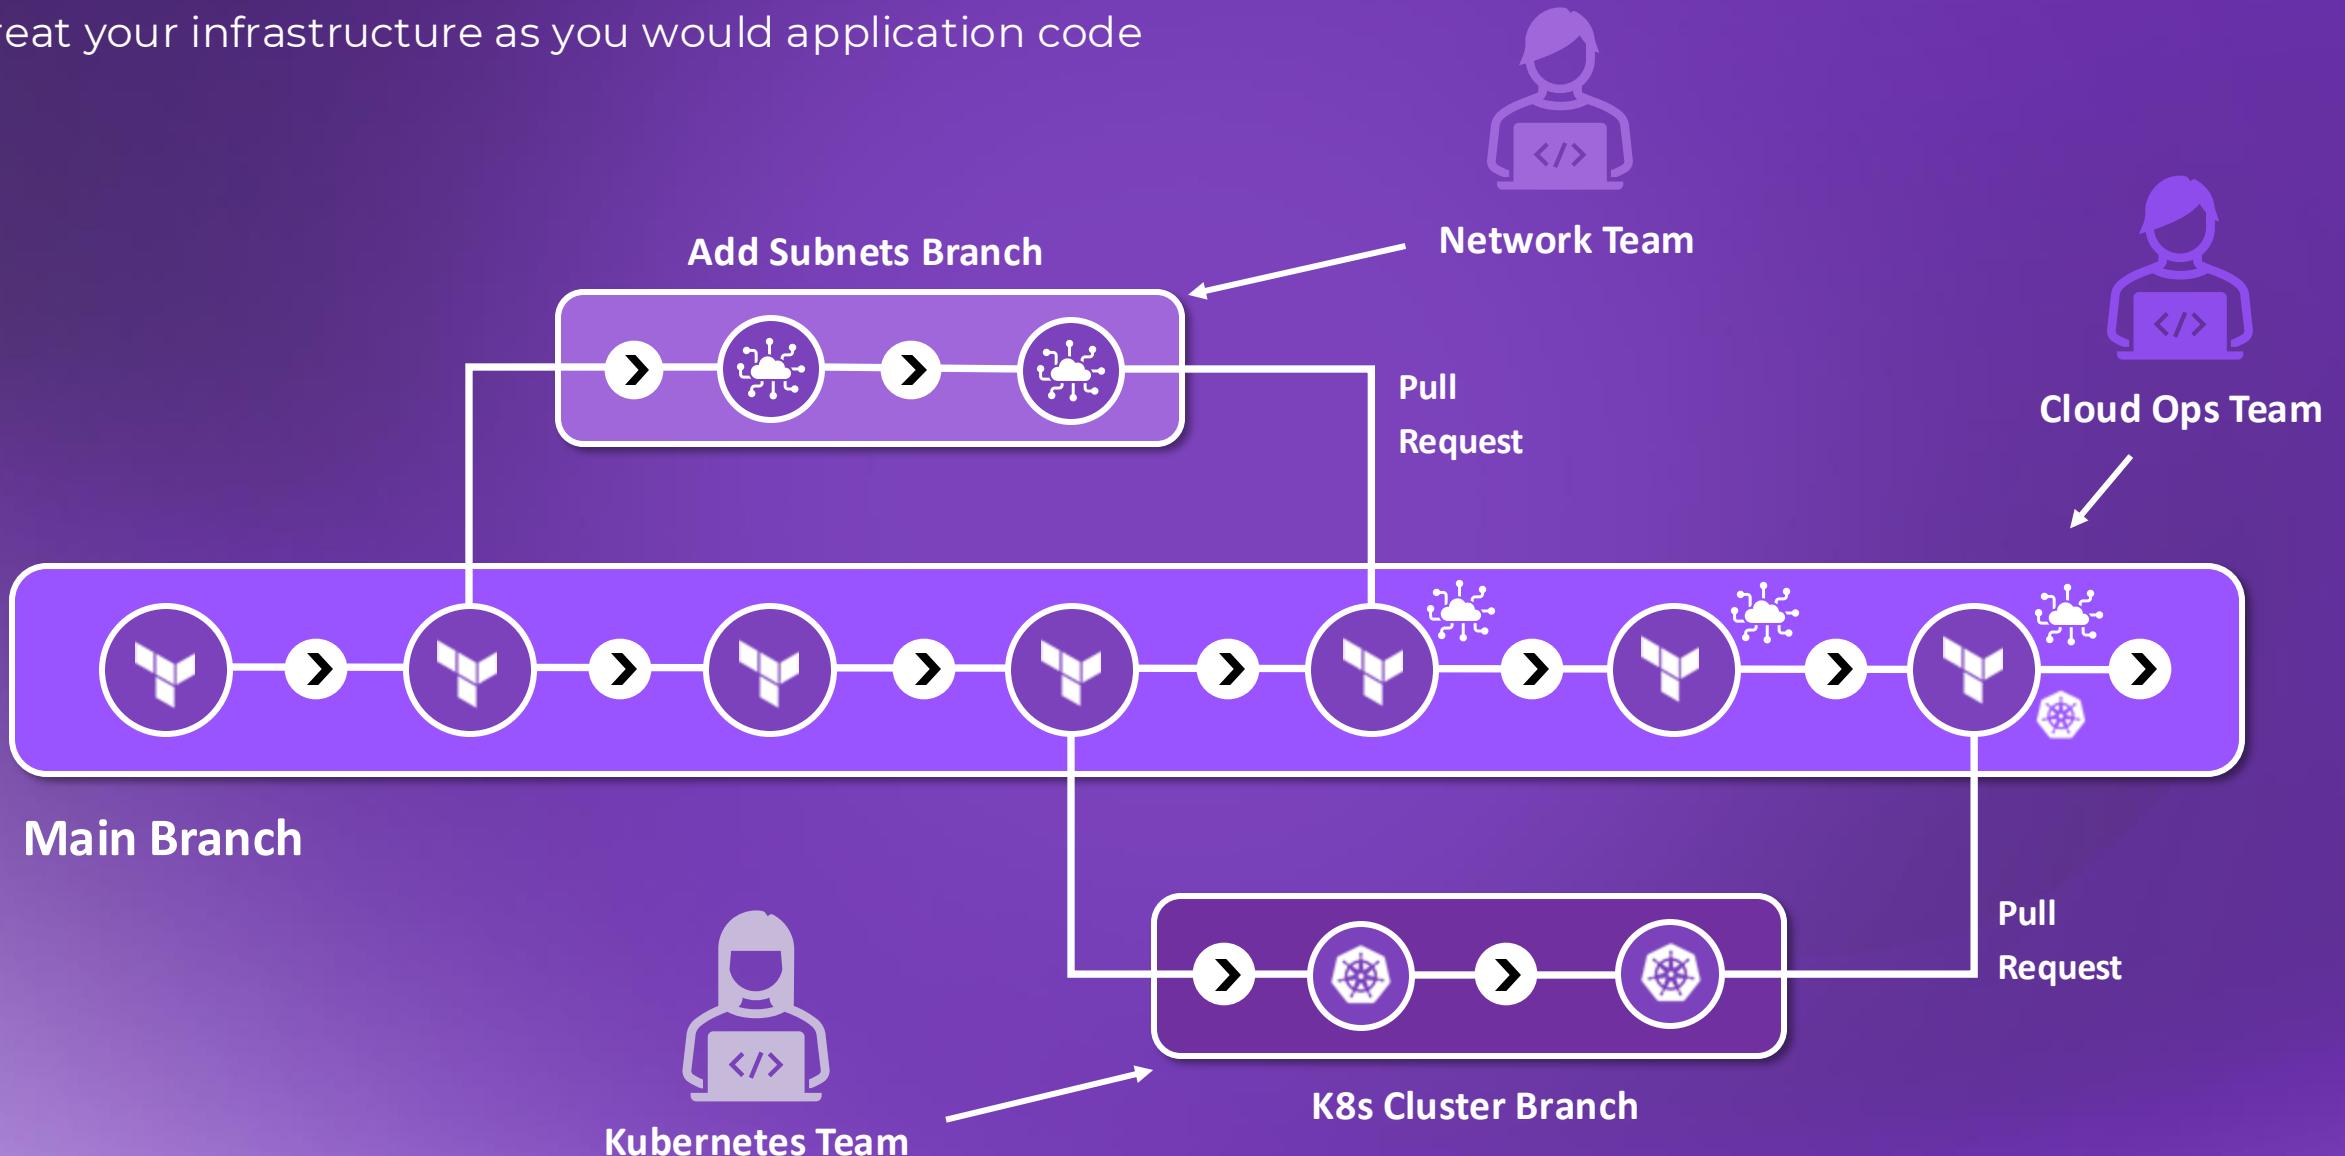


Data Application



Infrastructure as Code

Treat your infrastructure as you would application code





Speed & Consistency

Terraform can help you quickly spin up and decommission infrastructure for development, testing, QA, and production.

By defining infrastructure as code, you can reduce the risk of mistakes. You can create standardized components that can be integrated across projects consistently.

You can save infrastructure resources and their dependencies as code in a specified environment.



Core Components



Terraform Core

CLI tool that provisions and manages infrastructure resources as defined in Terraform configuration files



Providers

Extends the functionality of Terraform for specific platforms, such as public cloud providers, SaaS offerings, etc.



Resources

Infrastructure components or services that are managed by Terraform (virtual machines, networks, DNS records, etc.)



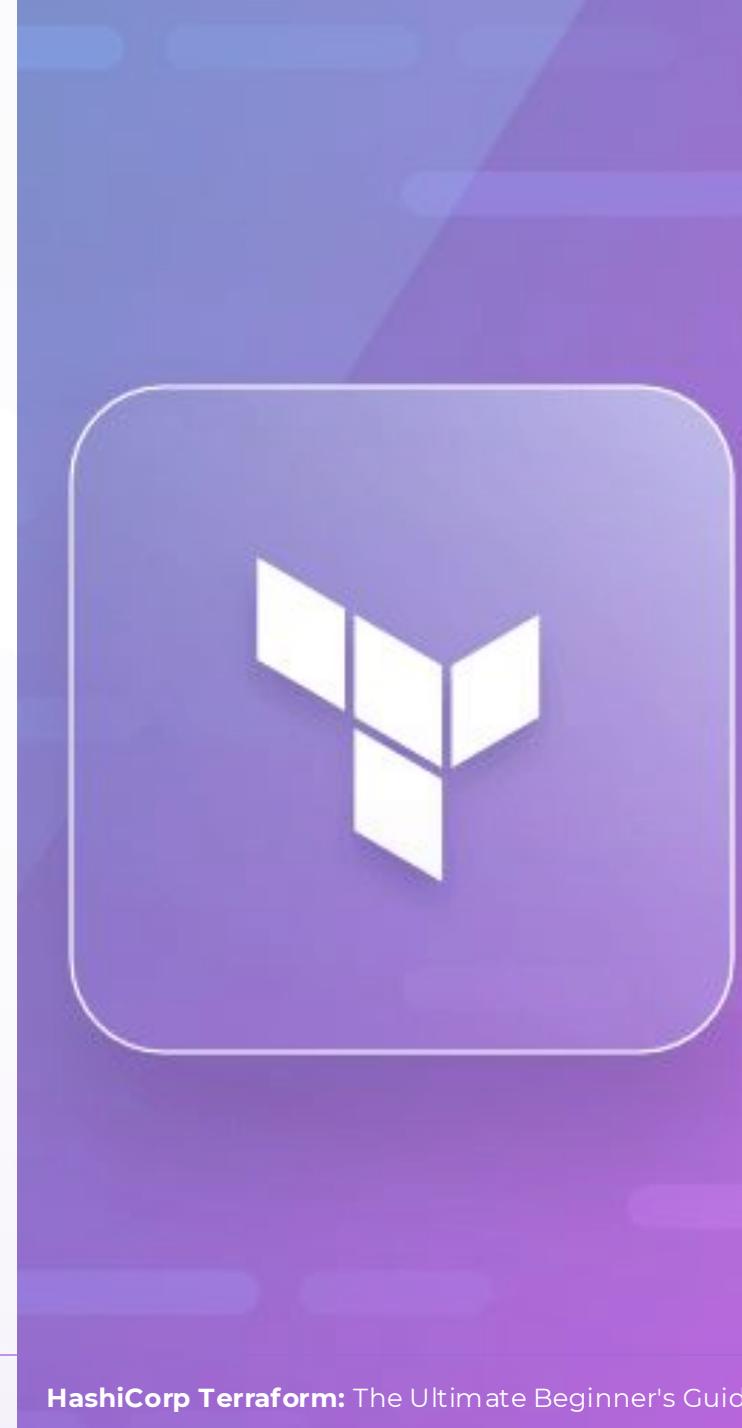
State

How Terraform maps the desired configuration with real-world resources on the target platform



Modules

Reusable and sharable blocks of code that can be called over and over again





Terraform Editions



Terraform Community



Command-Line Based



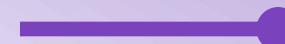
HCP Terraform



SaaS Platform



Terraform Enterprise



Self-Hosted & Managed

Compare Terraform with Other Tools

Terraform is often confused with other popular tools commonly found in many organizations.

Let's set the record straight...



Similar to Terraform



AWS
CloudFormation



Azure
Bicep



Pulumi

Different but Complementary



ANSIBLE



CHEF™



puppet

Infrastructure as Code Tools

- ❖ Infrastructure Provisioning and Management
- ❖ Declarative Language
- ❖ Immutable Infrastructure
- ❖ Modular Design

Configuration Management Tools

- ❖ Installing Packages
- ❖ Managing Files
- ❖ Manage Configuration Systems
- ❖ Automating Routine Tasks
- ❖ Orchestrating Multi-Step Tasks



What Does It Mean to Be Declarative?

A declarative approach focuses on defining *what* you want the outcome to be, not *how* to achieve it.



I want an EC2 instance with these properties.

DECLARATIVE

vs

Step 1: Build VM
Step 2: Configure VM
Step 3: Attach to Network



IMPERATIVE

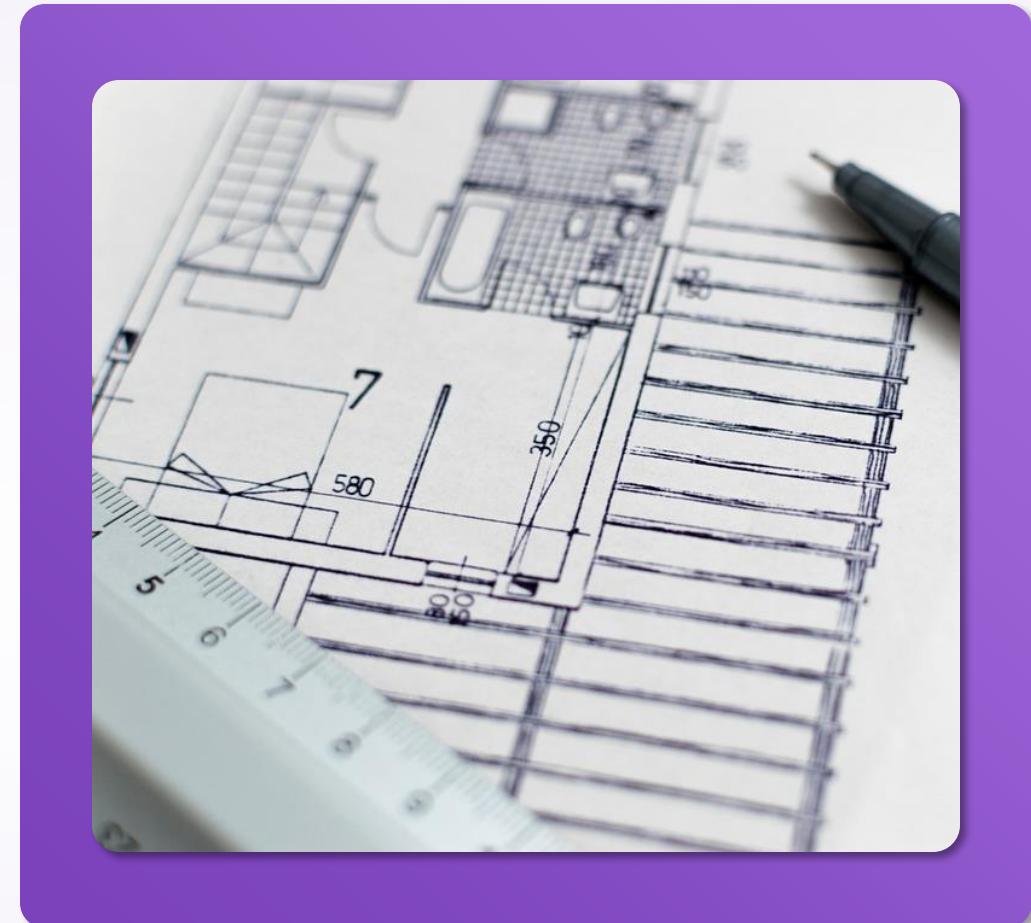


How Terraform Defines Desired State

The desired state is the blueprint of your infrastructure: resources, properties, and relationships.

What you write in your Terraform configuration is known as the desired state. You are telling Terraform "Create resources exactly how I've defined it within my Terraform configuration."

Terraform compares the current state of 'real-world' resources with the desired state and makes the appropriate changes.



How Do I Get Started with Terraform?

1

Download and Install Terraform

Get the proper binary for your machine.

2

Prepare Your Workspace

Install your favorite IDE, create directories, etc.

3

Obtain Credentials for Target Platform

Terraform needs a way to communicate with your cloud

4

Write Your First Terraform Config

Start with basic resources and iterate from there.

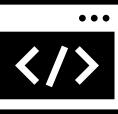
SUMMARY



Terraform is a tool created by HashiCorp that is used to manage infrastructure as code.



Terraform takes a declarative approach, which means that users describe what they want and Terraform will figure out



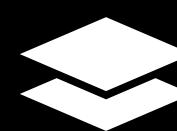
Terraform uses the HashiCorp Configuration Language (HCL) to define infrastructure



Terraform is platform-agnostic. It can be used to manage infrastructure on any cloud platform or data center.



Benefits of Terraform include version control, scalability, increased collaboration, and repeatability



Terraform's core components include Terraform Core, Providers, Resources, State, and Modules.

Lecture

Learn the Basics of HCL





HashiCorp Configuration Language (HCL)

Terraform configuration language to describe the infrastructure that Terraform manages.

I O I O
I O I O

Declarative Language

Used in Terraform (and other HashiCorp products) to define infrastructure configurations



Easy to Read and Write

HCL provides a clean and human-readable syntax for expressing resources and dependencies



Terraform's Primary Interface

Configuration files that you write tell Terraform what infrastructure you want to create or what data to retrieve

HashiCorp Configuration Language (HCL)



```
1
2 block_type "block_label" "block_label" {
3   first_argument  = expression or value
4   second_argument = expression or value
5   third          = expression or value
6 }
7
8 attribute_abc = "value_1"
9 attribute_2   = "value_2"
```





HCL Example

data
blocks {

resource
block {

```
main.tf
```

```
1 # Retrieve the list of AZs in the current AWS region
2 data "aws_availability_zones" "available" {}
3 data "aws_region" "current" {}
4
5 # Define the VPC
6 resource "aws_vpc" "vpc" {
7   cidr_block = var.vpc_cidr
8
9   tags = {
10     Name      = var.vpc_name
11     Environment = "demo_environment"
12     Terraform  = "true"
13   }
14 }
```

- single-line comment
- comment

} arguments



Block Example

Type of block Resource Type Name
aws_vpc.vpc

```
1 resource "aws_vpc" "vpc" {  
2   cidr_block = var.vpc_cidr  
3  
4   tags = {  
5     Name      = var.vpc_name  
6     Environment = "demo_environment"  
7     Terraform   = "true"  
8   }  
9 }
```



Blocks are delineated by {}

All contents within the block are arguments some type of value



Blocks start with a keyword

Defines what type of block it is, such as resource, data, variable, data, etc.



Blocks have named values

Terraform can access information created by or obtained by a block



HCL Style Guide

Writing Terraform code in a consistent style makes it easier to read and maintain.

- Ү Use `#` for single and multi-line comments
- Ү Use underscores to separate multiple words in names
- Ү Indent two spaces for each nesting level
- Ү Align equal signs for values on consecutive lines
- Ү Use empty lines to separate groups of arguments in a block

```
1 # single-line comment
2 block_type "block_label" {
3   first_argument = expression or value
4   second_argument = expression or value
5   third         = expression or value
6 }
7
8 attribute_abc = "value_1"
9 attribute_2   = "value_2"
```



Resource Referencing

Resource referencing allows you to connect resources within your configuration, making it easy to share data between resources and build dependencies.



Create More Dynamic Configurations

Terraform can use properties from one resource in another and avoid hardcoding values, creating a dynamic configuration



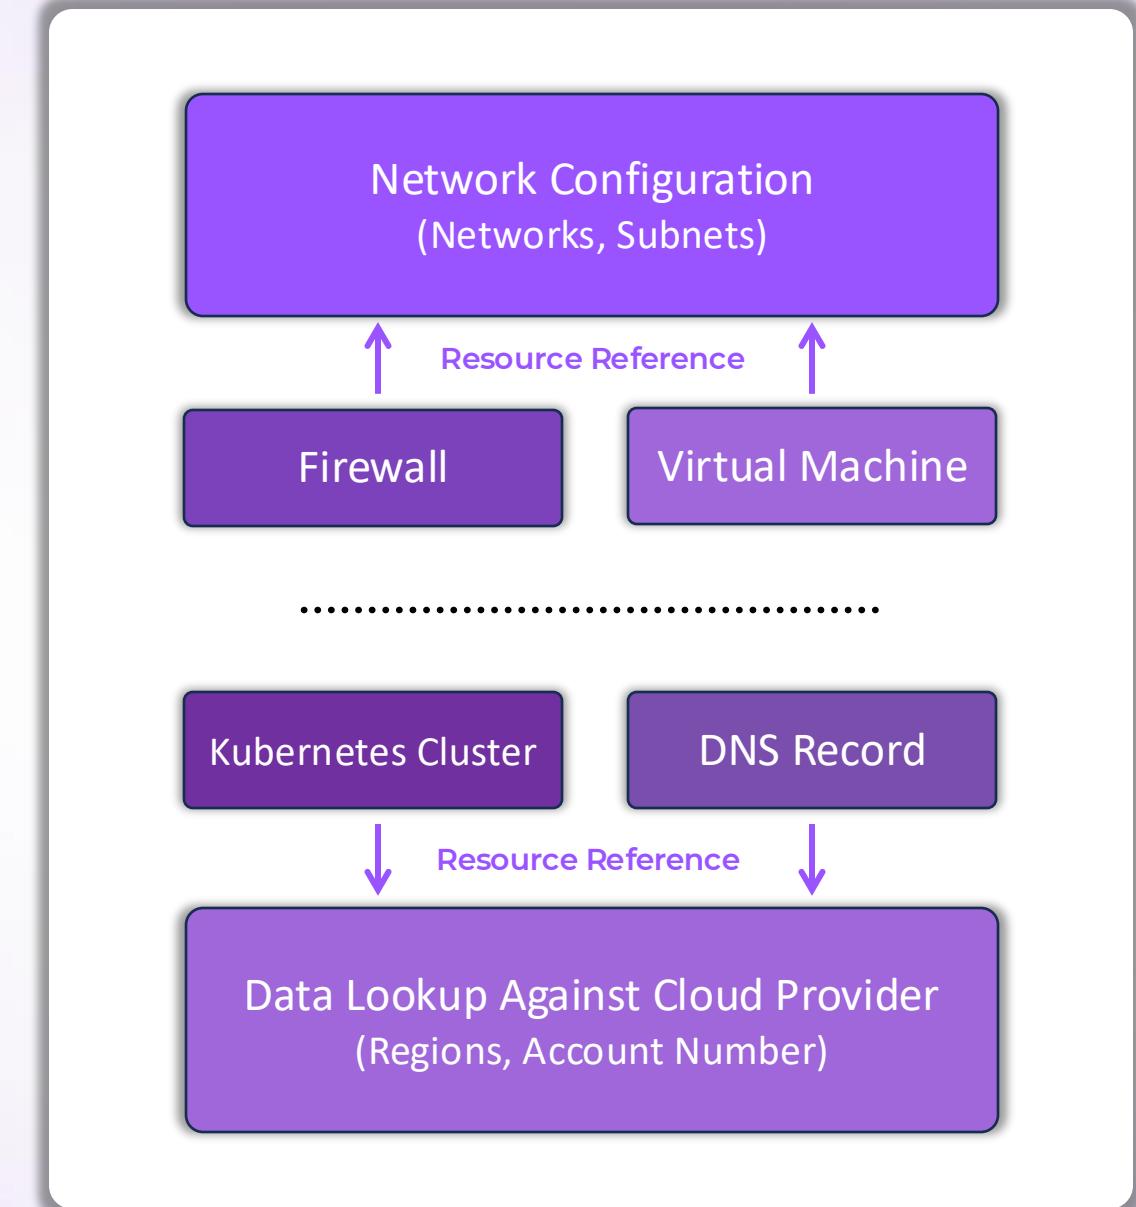
Automatic Dependency Mapping

Terraform automatically determines the order of resource creation based on references.



Resource Identifiers

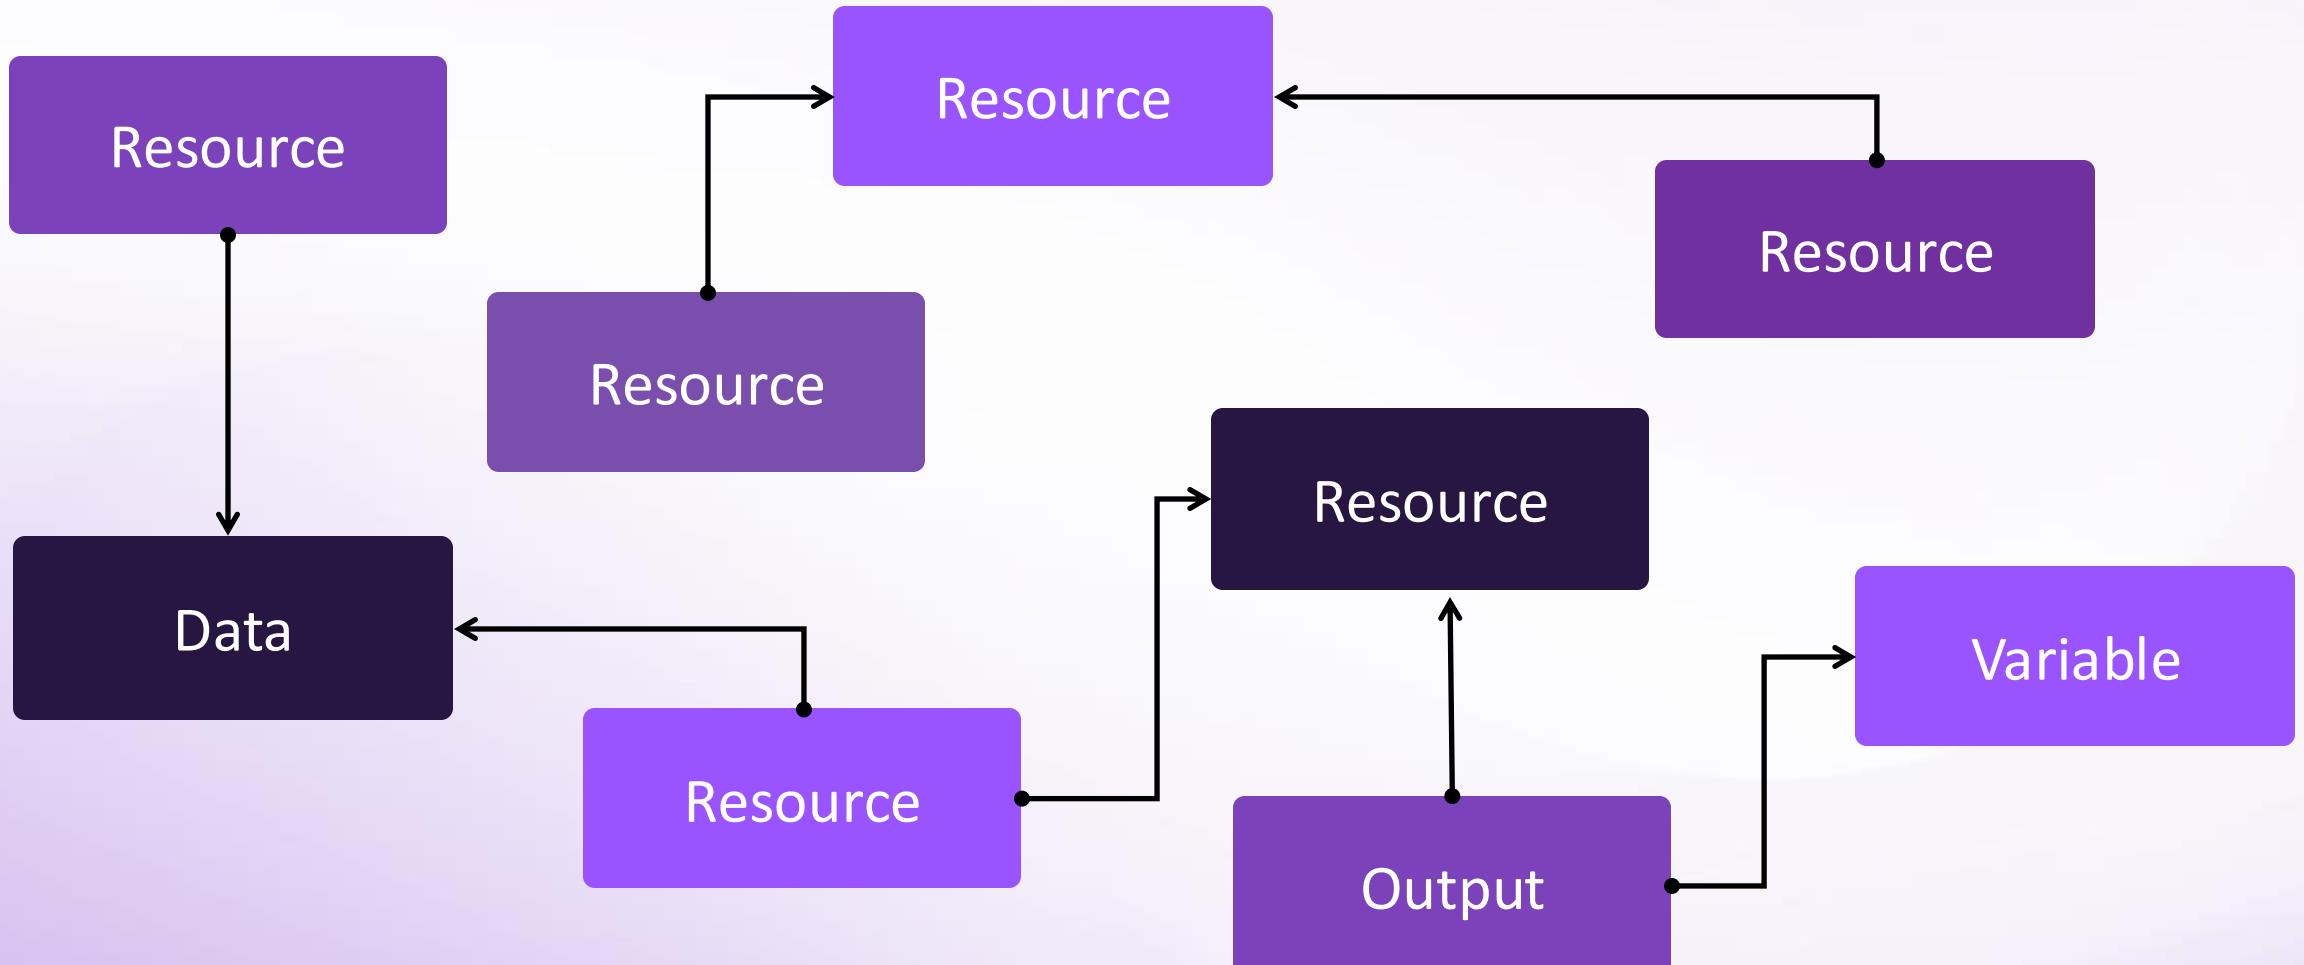
Each Terraform block is identified by a **type** and a **name**, which makes it easy to reference in other parts of your configuration





Resource Referencing

The Reality of Terraform



DFMO

Learn the Basics of HCL



Best Practices for HCL



File Extensions

Terraform HCL files use the .tf .tfvars file extension. This allows Terraform and related tools to recognize and apply configurations properly



Formatting

Correctly formatting files improves readability and maintainability. Follow the style guide or use built-in tools to format your files



Organization

Organize functionality within files (variables in a variables file, outputs in an outputs file, etc.)



Commenting and Documentation

Include comments and README.md files to provide context and instructions, making it easier for others to understand and use the code



Avoid Hardcoding Values

Use variables for values that may change which makes configurations flexible and reusable.



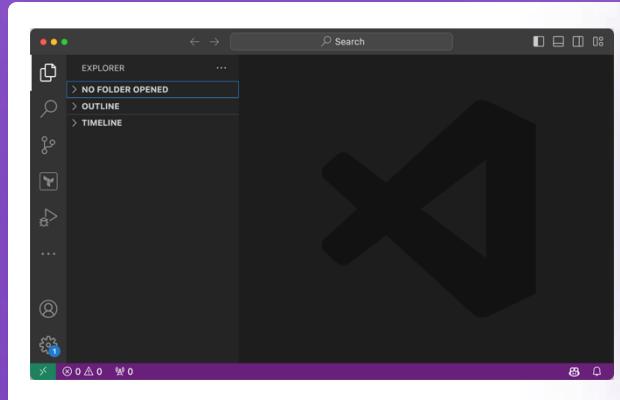


Enhancing Code Quality

Using Linters, Extensions, and Built-In Tools for Code

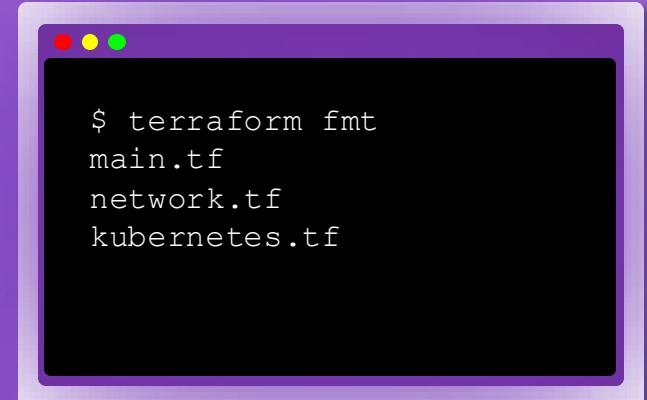
Properly formatted code improves readability and reduces errors. Use a VSCode extension or linter for real-time syntax checks and best-practice suggestions.

Or run `terraform fmt` to automatically format HCL files with a Terraform-related extension (`.tf`)



Linters and Extensions

Linters provide immediate feedback on code issues, while extensions offer additional features like autocompletion and integrated commands



Built-In Tools

`terraform fmt` automatically formats Terraform code to ensure consistency and readability, making collaboration easier across teams.

SUMMARY



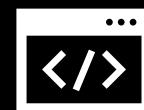
HCL is a declarative language that is easy to read and write, and is Terraform's primary interface.



A single Terraform file can have one or many blocks to define your infrastructure



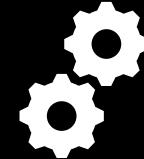
HCL uses blocks for data, resources, etc., and these blocks can be nested.



Use the '.tf' file extension so Terraform and related tools can recognize and apply configurations properly.



Single-line and block comments are used in HCL.



Use linters, extensions, and built-in tools like 'terraform fmt' to enhance code quality.

Lecture

Introduction to OpenTofu



HISTORY

2014

First Commit of Terraform Open-Source by Mitchell Hashimoto

2015

Terraform adds state locking for backend

2017

HashiCorp announces Terraform Enterprise.



2018

Terraform 0.12 released, a major upgrade introducing new language syntax improvements

2023

HashiCorp announces a license change from Mozilla to BSL. Terraform no longer "open-source". This leads to the fork of Terraform 1.5 to create OpenTofu

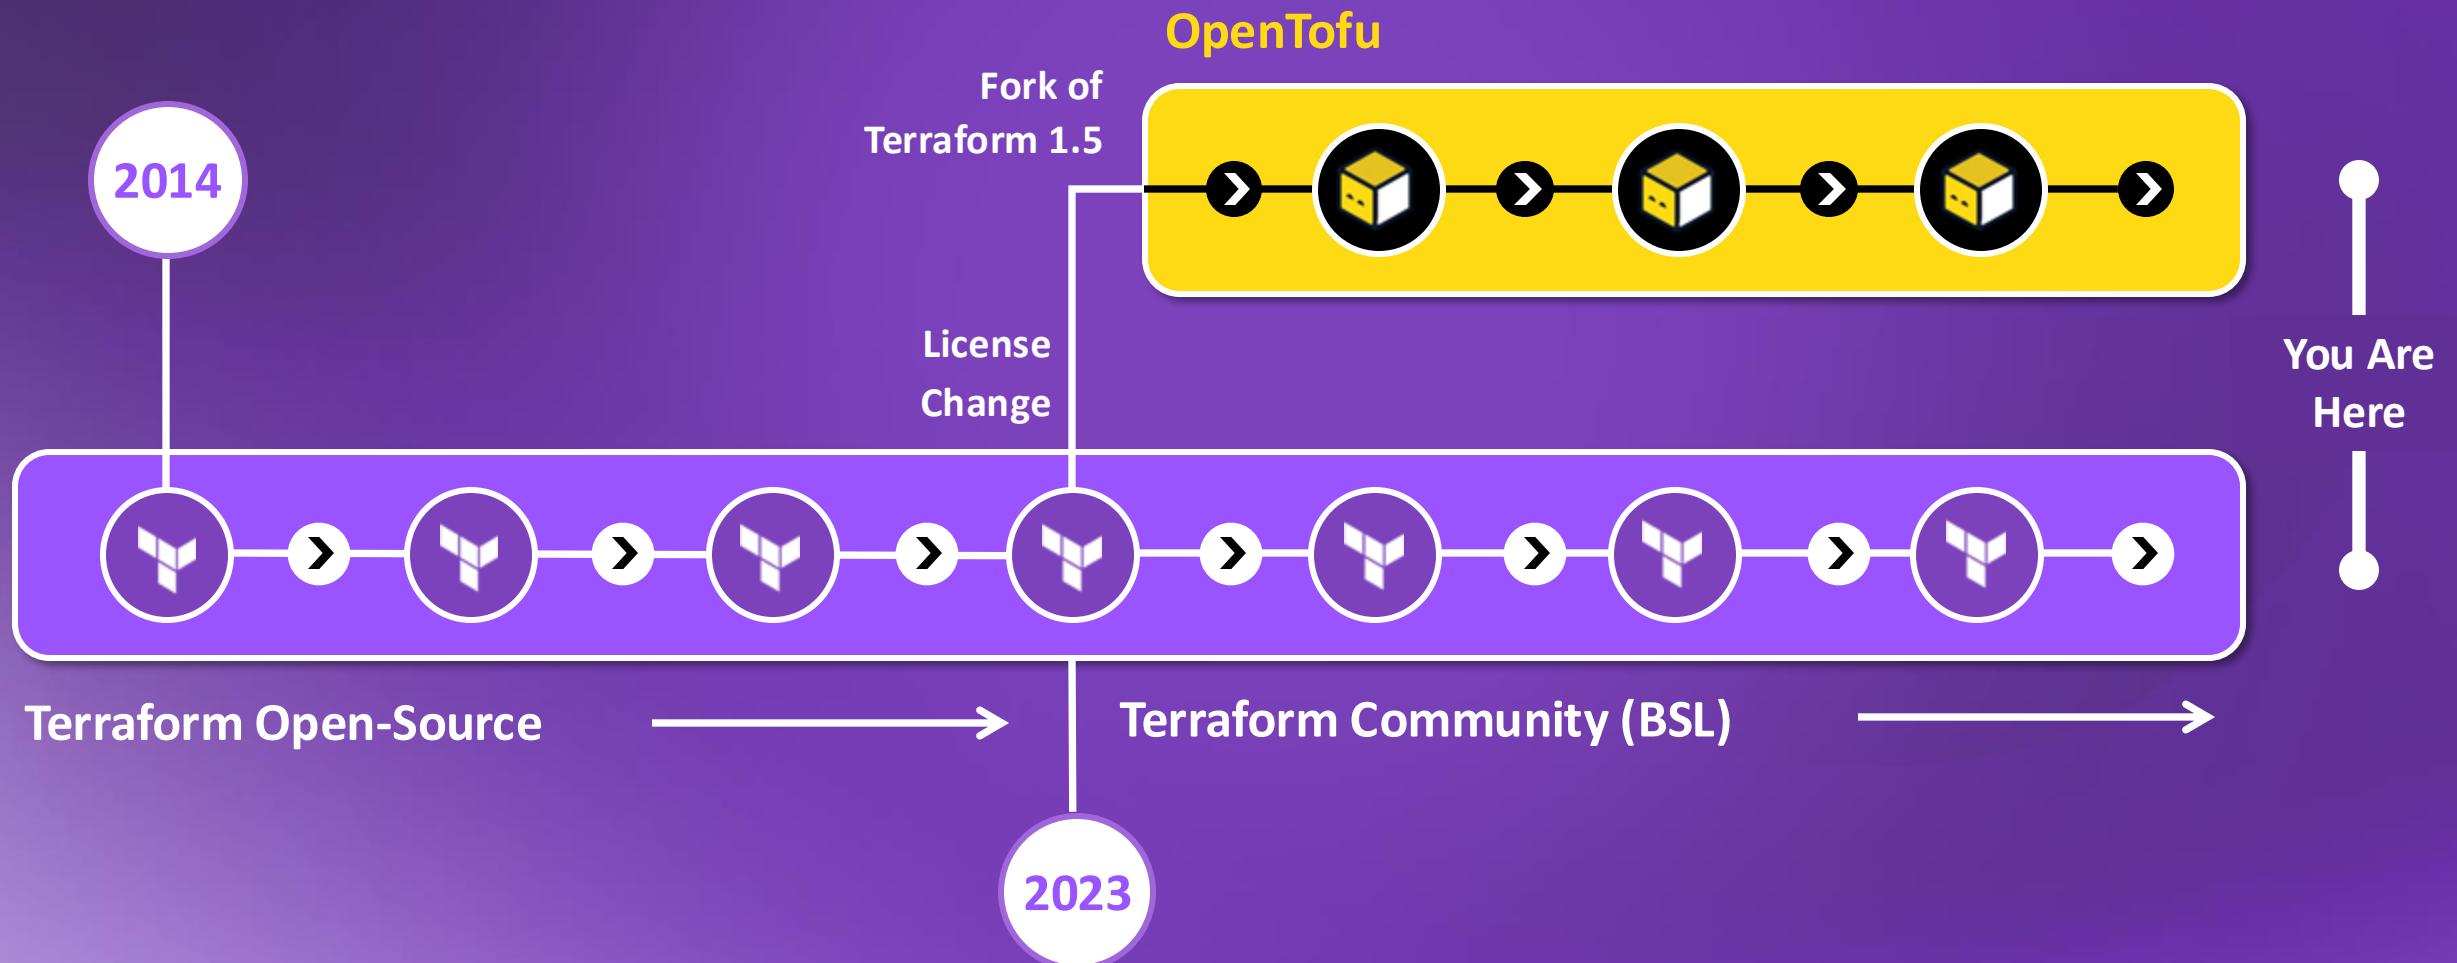
Current State

Terraform and OpenTofu are developed as separate tools.



OpenTofu and Terraform

Understanding the Relationship



Why OpenTofu?

Built from Terraform 1.5, the last MPL version, OpenTofu ensures individuals and organizations retain full access to their infrastructure code without the constraints of proprietary licensing.



Open-Source Commitment

OpenTofu maintains an open-source license (MPL 2.0) to ensure unrestricted use for individuals and organizations.



Community-Driven Development

Developers and organizations can influence its roadmap and improvements. Driven by a volunteer steering committee rather than a single company



Seamless Transition for Existing Terraform Users

Enables a smooth migration path for users impacted by the licensing changes while preserving the core Terraform experience.



Differences Between Terraform and OpenTofu



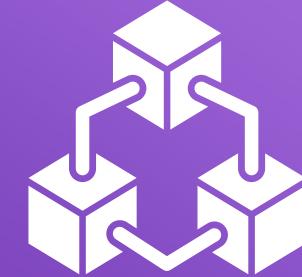
Similar but Different

Since OpenTofu is a fork of Terraform, the differences are subtle (at least for now). OpenTofu supports backward compatibility.



Licensing Concerns

OpenTofu is open-source and managed by the Linux Foundation. Terraform is "source available" but has certain restrictions of use.



New Features

OpenTofu has implemented many long-awaited features that the community has requested for years. HashiCorp is also adding features.



Are Organizations Really Moving to OpenTofu?

Absolutely. Organizations are finding that the features being added to OpenTofu are very impactful to their organization and management of Infrastructure as Code strategy moving forward.

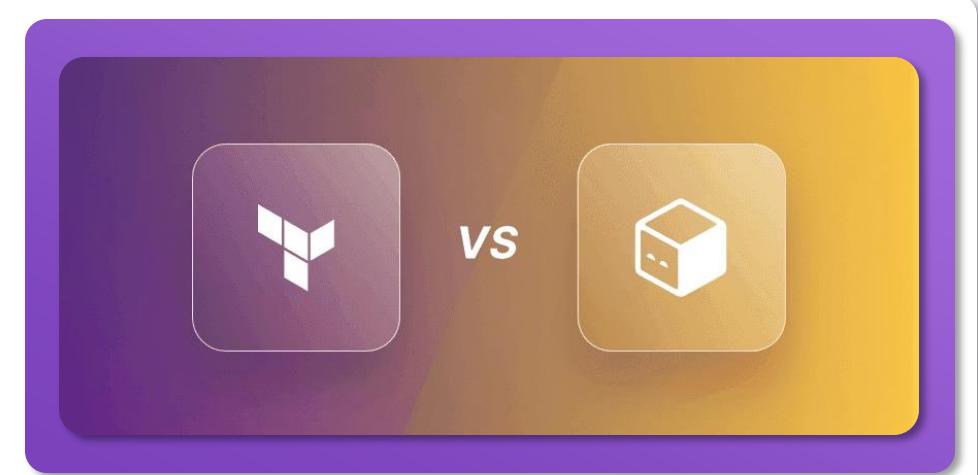
Real-world examples include:

- licensing concerns due to its chosen platform
- support of state encryption
- module source interpolation
- automation platform will not run BUSL versions of Terraform
- preference for licensing which embodies free and open-source software beliefs

Should I Use Terraform or OpenTofu?

Terraform is the de facto standard, but the market is shifting.

As Terraform and OpenTofu continue to introduce unique features at their own pace, the two products are likely to evolve in increasingly different directions. Over time, this divergence could lead to substantial differences in functionality, potentially resulting in limited or no backward compatibility between them in the future.



I Can't Choose for You

All I can do is give you the best information available, offer my own opinions, and let you make your own decision.

Read: opentofu.org/manifesto

SUMMARY



In 2023, HashiCorp changed the licensing for Terraform, leading to the creation of the open-source fork, OpenTofu.



Licensing concerns, feature advancements, and open-source support are the primary concerns for moving from Terraform



OpenTofu supports backward compatibility with Terraform



Organizations have moved to OpenTofu, but many others still use Terraform daily.

LECTURE

File Structure and Organization



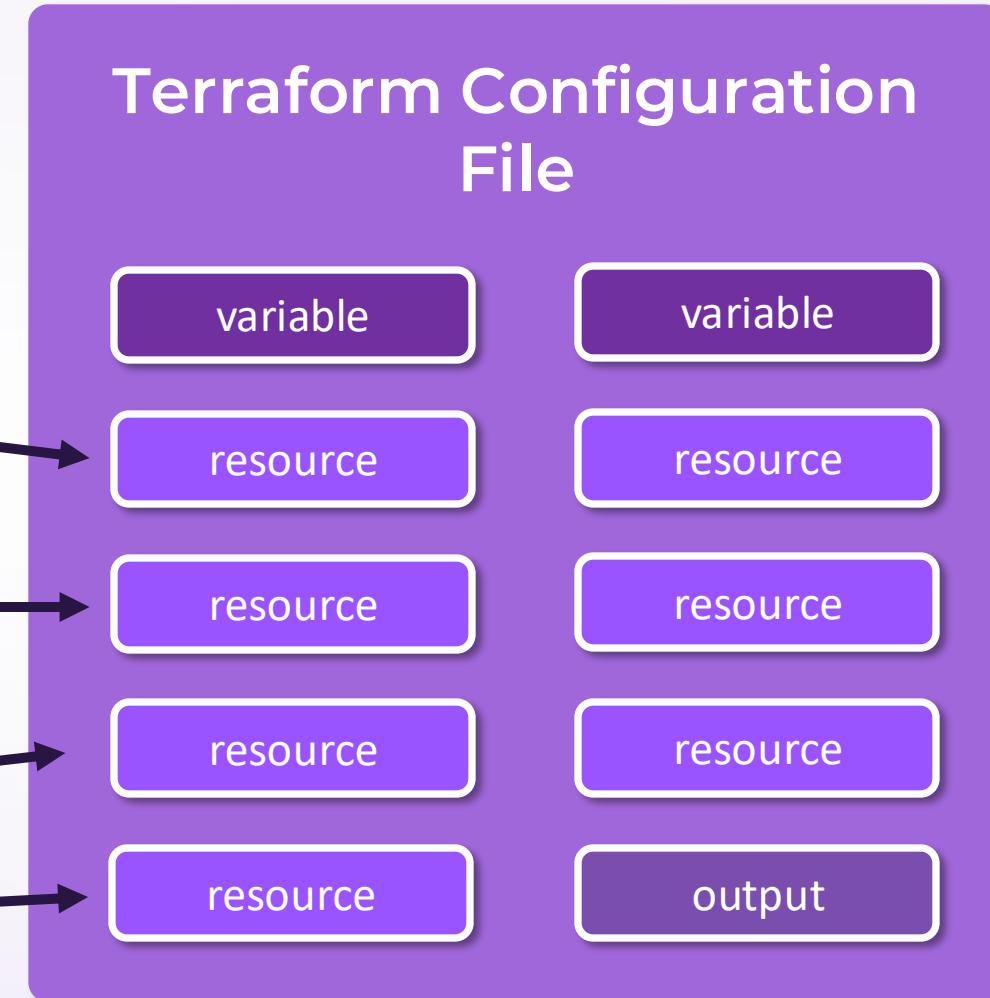


Organizing Code

customer-app.tf

Infrastructure Requirements:

- S3 Bucket
- Virtual Machine
- Subnet
- Firewall Configuration
- DNS Record
- Load Balancer
- Kubernetes Cluster





Organizing Code

Terraform Configuration File

variable

variable

resource

resource

resource

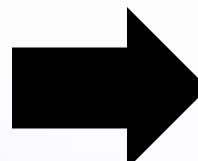
resource

resource

resource

resource

output



Working Directory

variables.tf

kubernetes.tf

network.tf

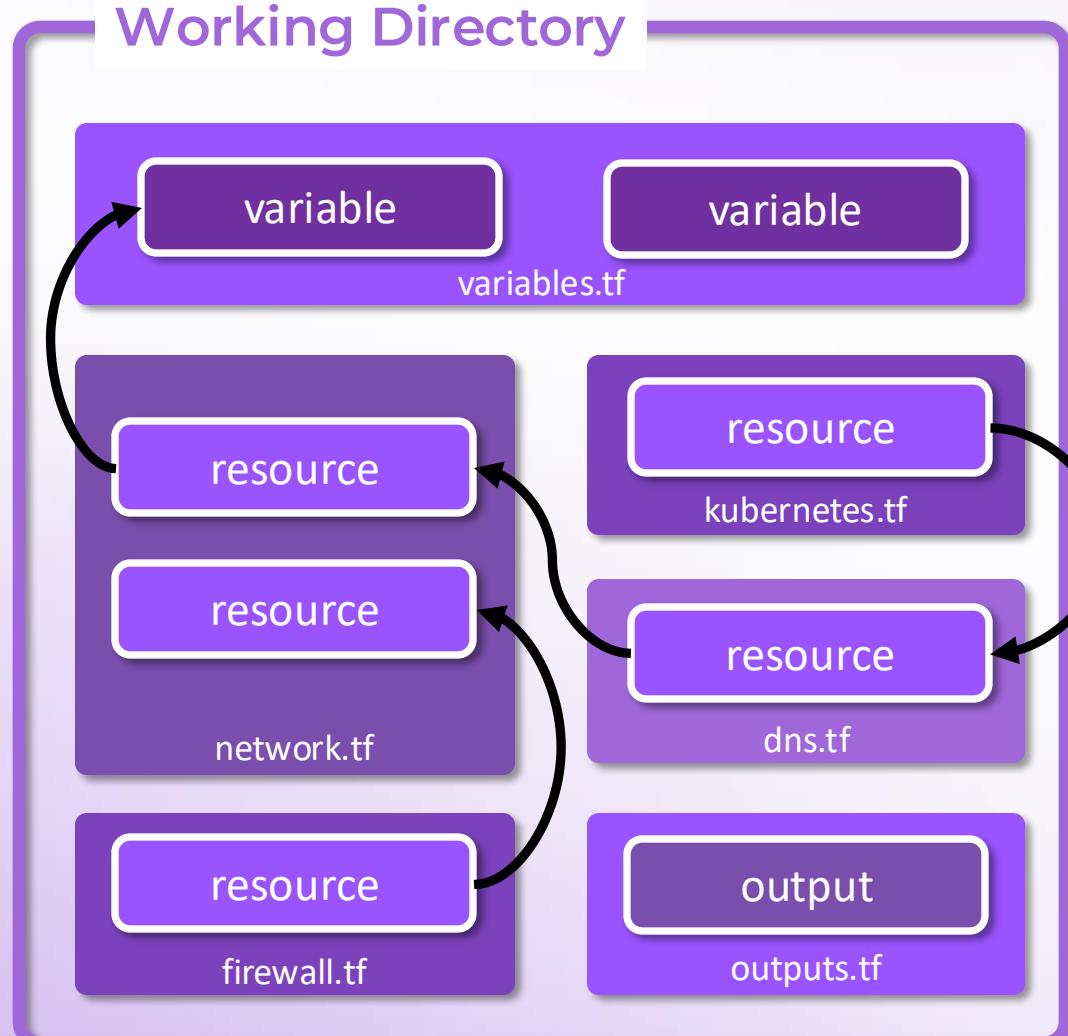
dns.tf

firewall.tf

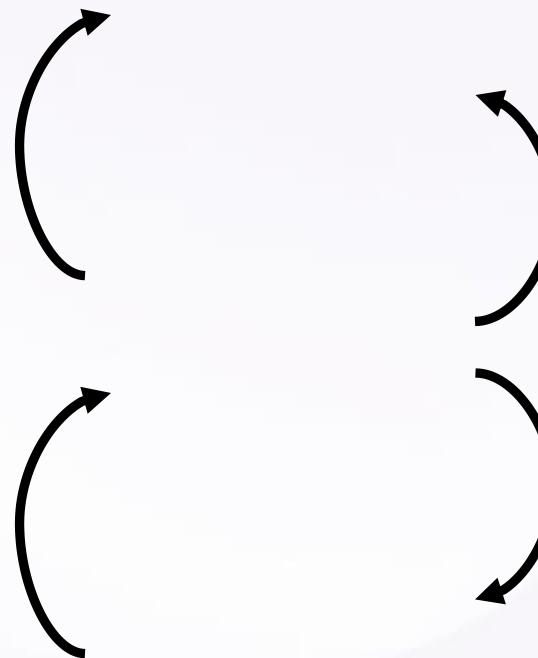
outputs.tf



Executing Code



Terraform Processes All Files Together



Common Terraform Files



main.tf

Commonly used for primary infrastructure components



variables.tf

Variable definitions used throughout your infrastructure



outputs.tf

Output blocks to output values to the terminal or to use for inputs to other modules



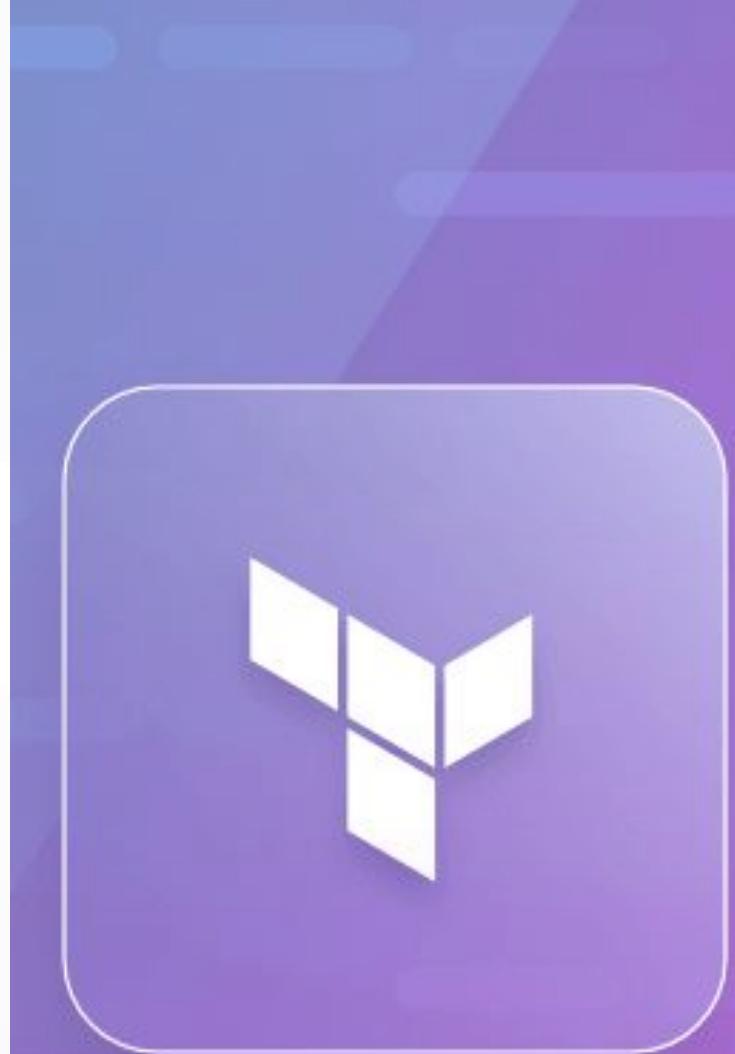
providers.tf

Provider configurations and requirements



terraform.tfvars

Define the values for your variables – usually ignored by .gitignore





Additional Files

- `terraform.tfstate`

The file that Terraform uses to store state

- `terraform.tfstate.backup`

Backup of the previous state file before Terraform writes anything new to the state file

- `.terraform.lock.hcl`

Used to track and select provider versions

- `.gitignore`

Instructs Git to ignore files included in this file

▽ `aws_deploy`

> `.terraform`

> `scripts`

 `.terraform.lock.hcl`

 `main.tf`

 `outputs.tf`

 `providers.tf`

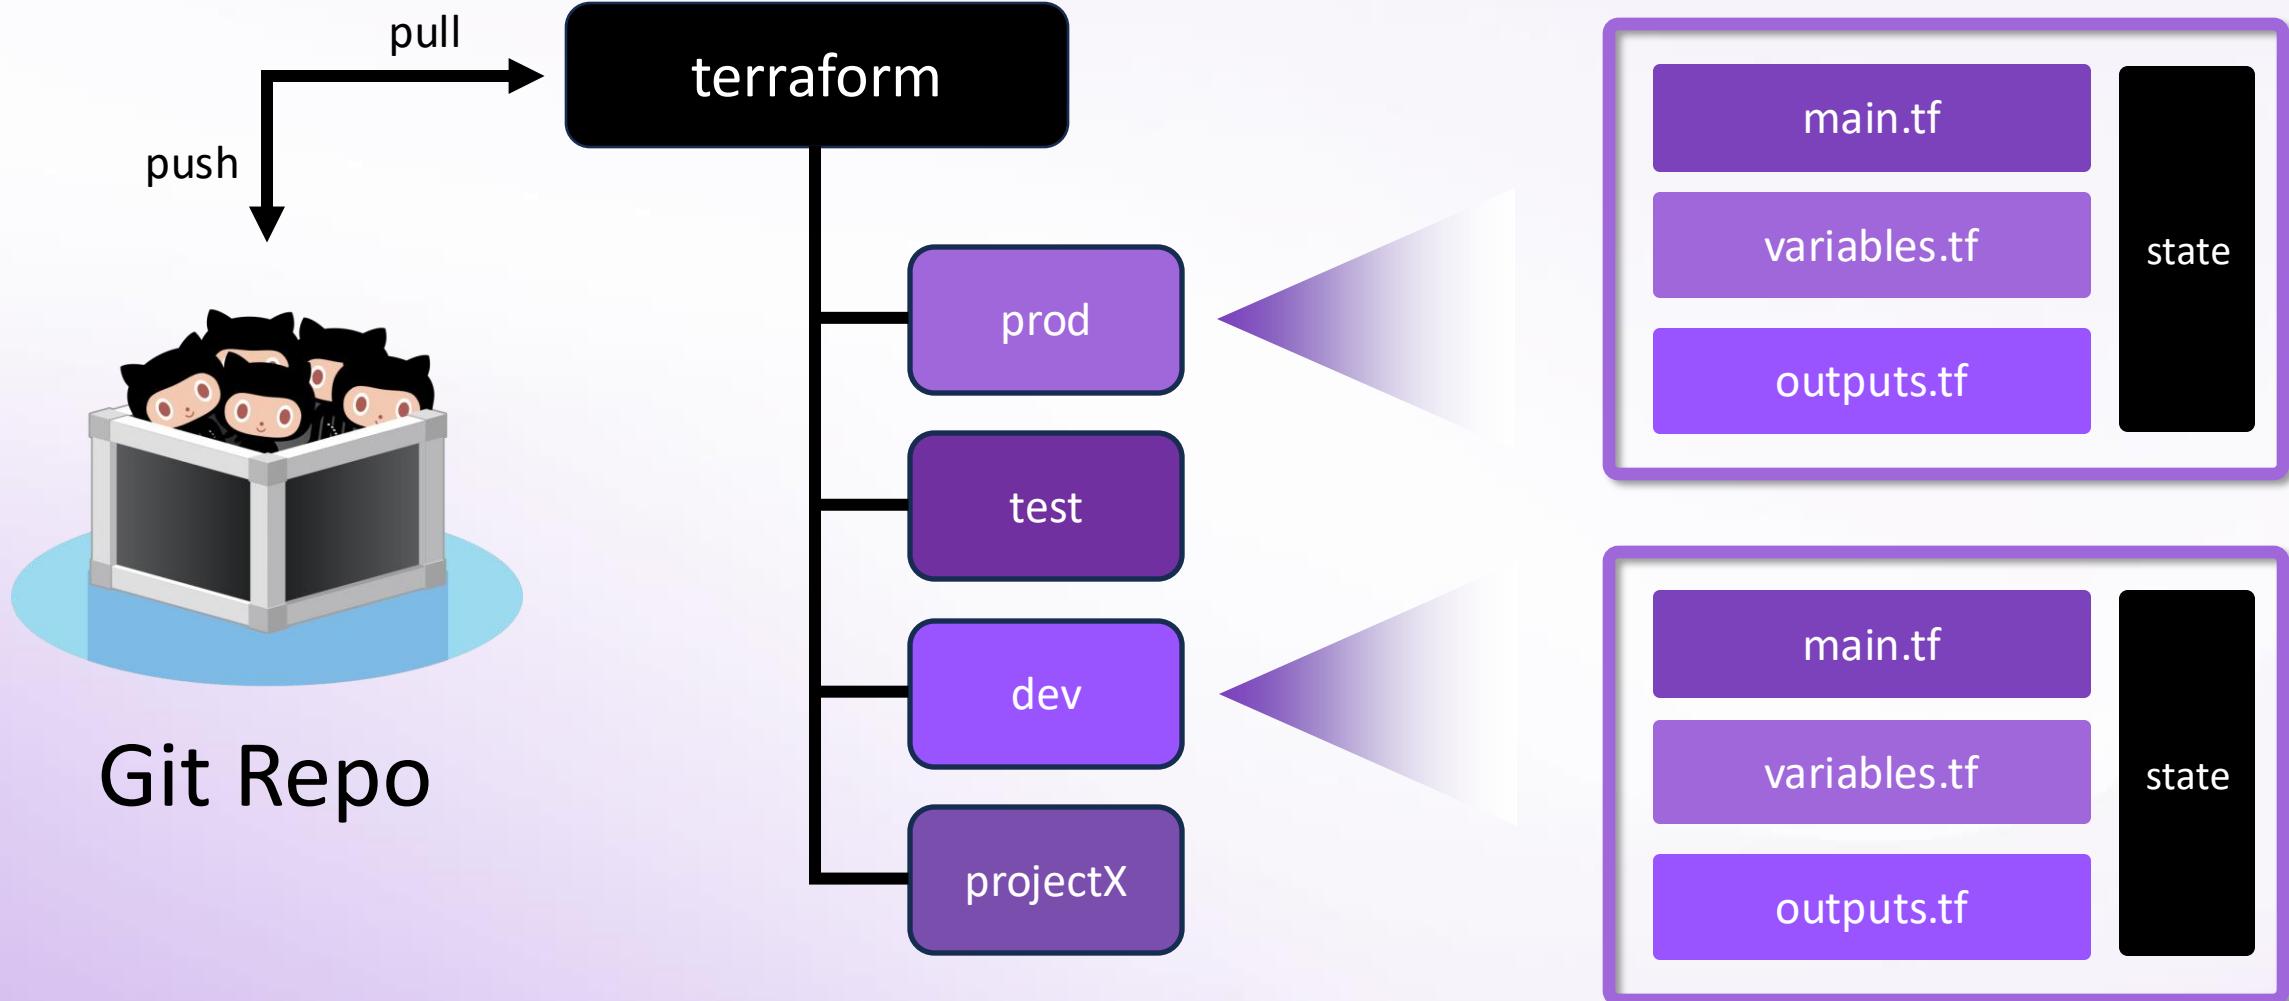
 `{ } terraform.tfstate`

 `≡ terraform.tfstate.backup`

 `variables.tf`



Working Directory



SUMMARY



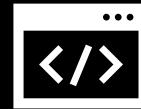
Terraform only processes the files in the current working directory unless you reference other files



While not required, there are common file names for Terraform configurations



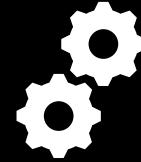
You can organize your files across multiple .tf files



Terraform creates and manages files for state and other configurations



Terraform combines all the files in the working directory when executing



Use subdirectories to break up your Terraform to better organize and reduce the blast radius of changes

INFRASTRUCTURE Understanding the Terraform Blocks



Terraform Blocks

The Building Blocks or Infrastructure



Clear Structure

Each component has a specific purpose, making it easy to organization and understand your configurations



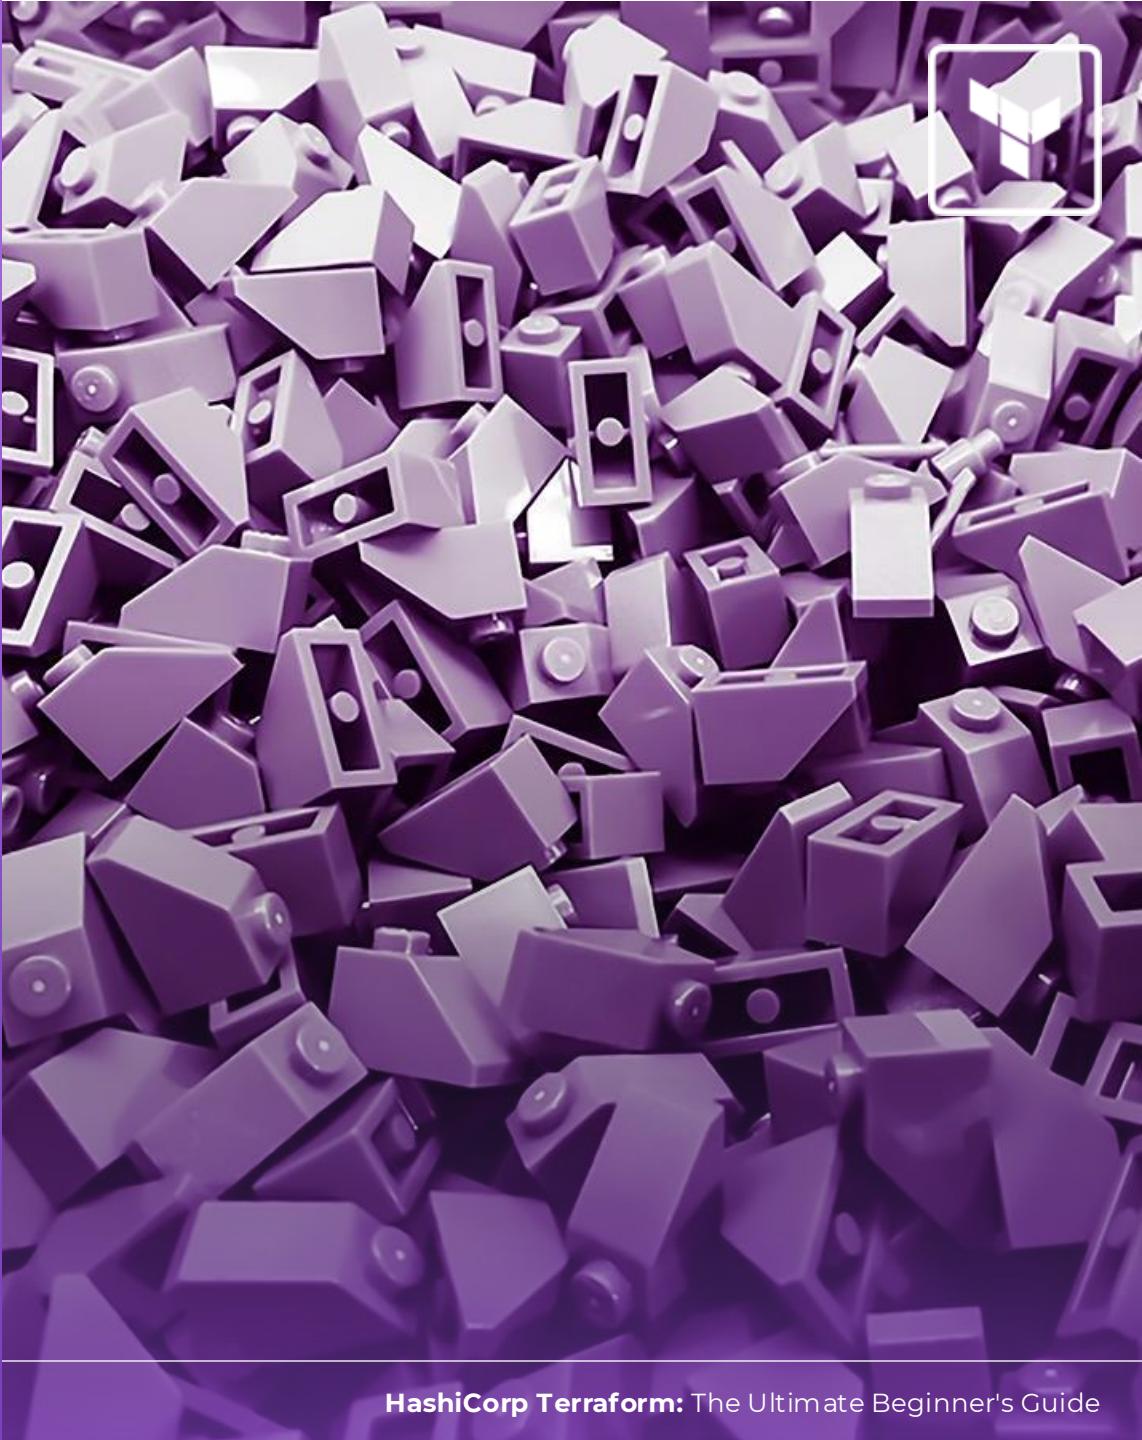
Flexible and Modular

Terraform is designed to be flexible. You don't need to use all the components to build infrastructure



Reusable Patterns

Using components thoughtfully helps create reusable, modular configurations that are easy to maintain and scale





Terraform Block Types



provider

Connect Terraform to cloud platforms or services (e.g., AWS, Azure, GCP, VMware)



resource

The infrastructure that Terraform creates, updates, or deletes (e.g., VMs, storage)



data

Retrieve information about existing resources to use in your configuration



variable

Make code reusable by defining values that can be customized for each run



output

Display or share essential information about resources after deployment



terraform

Define specific settings or resources in a configuration, like req. Terraform version



module

Reusable configurations that group related resources – makes it easy to share configs



import

Pull in existing resources into Terraform management

LECTURE

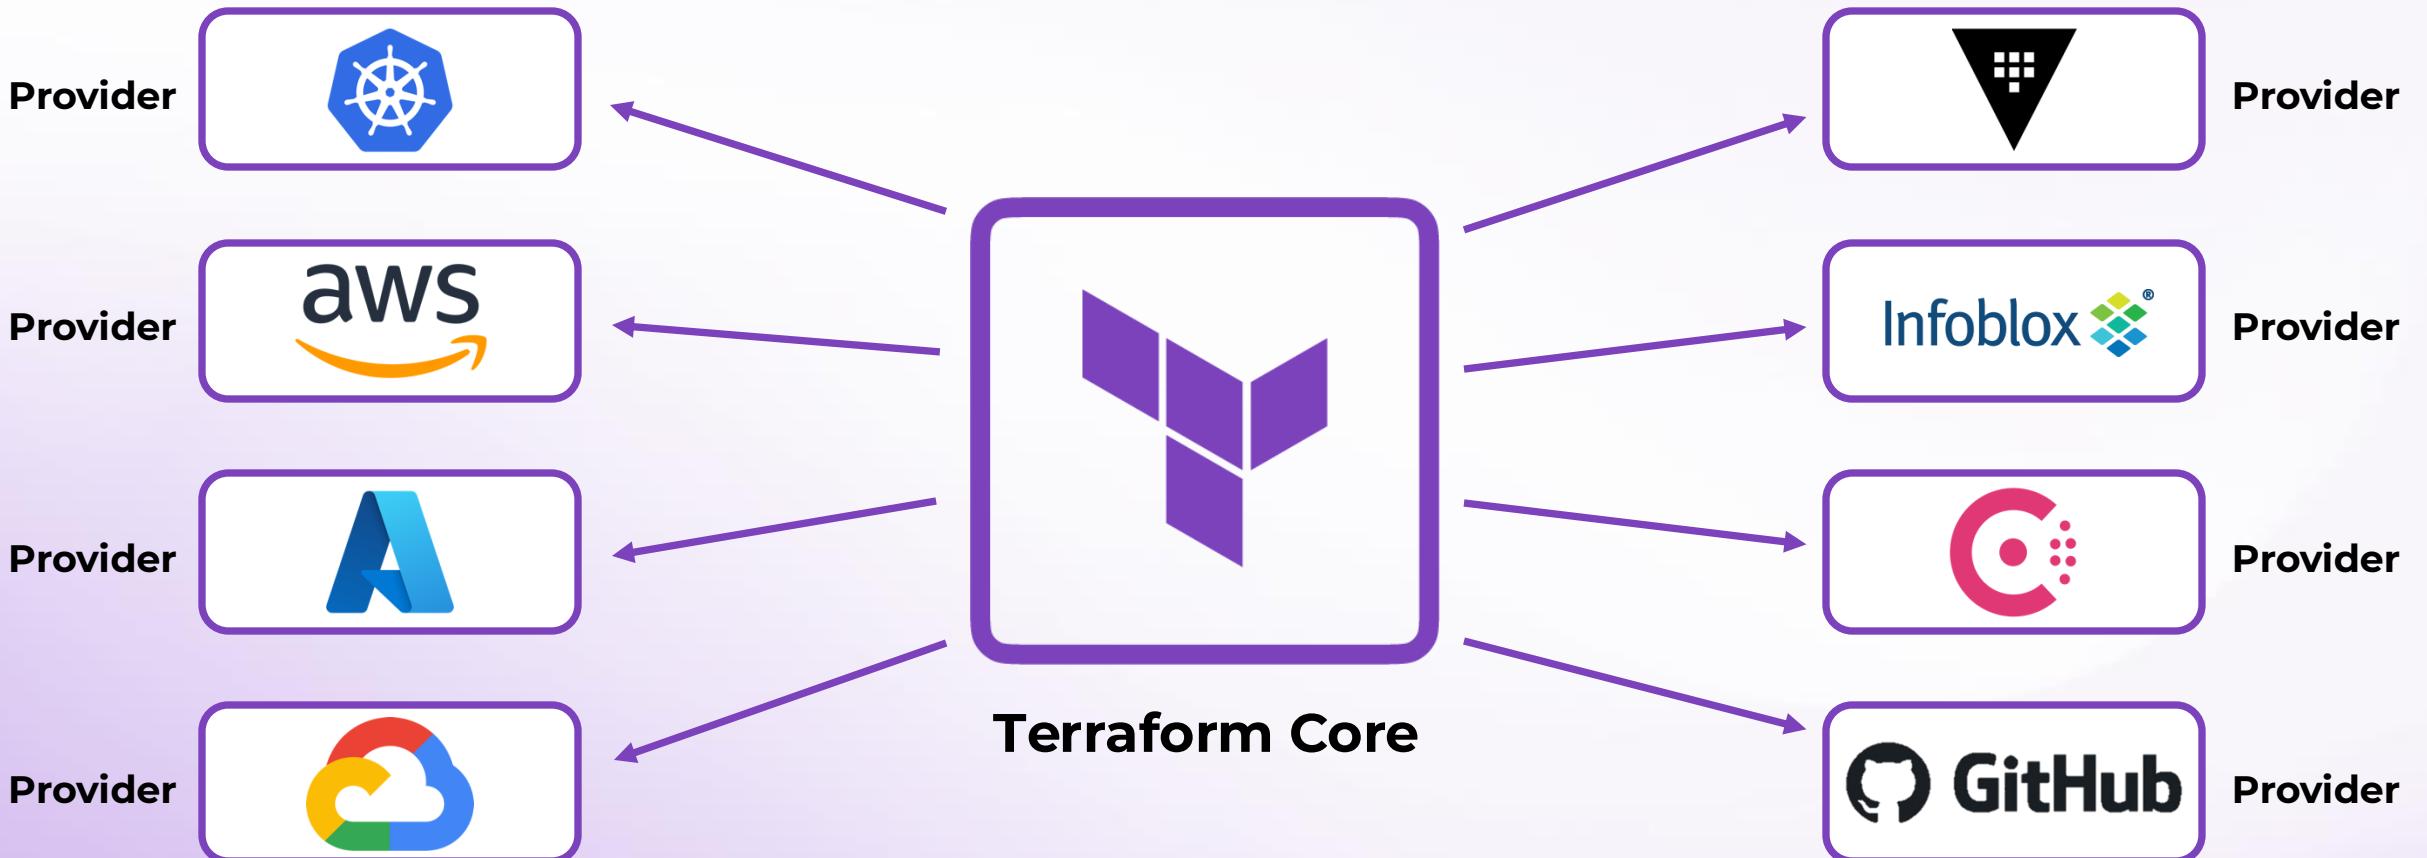
Provider Block

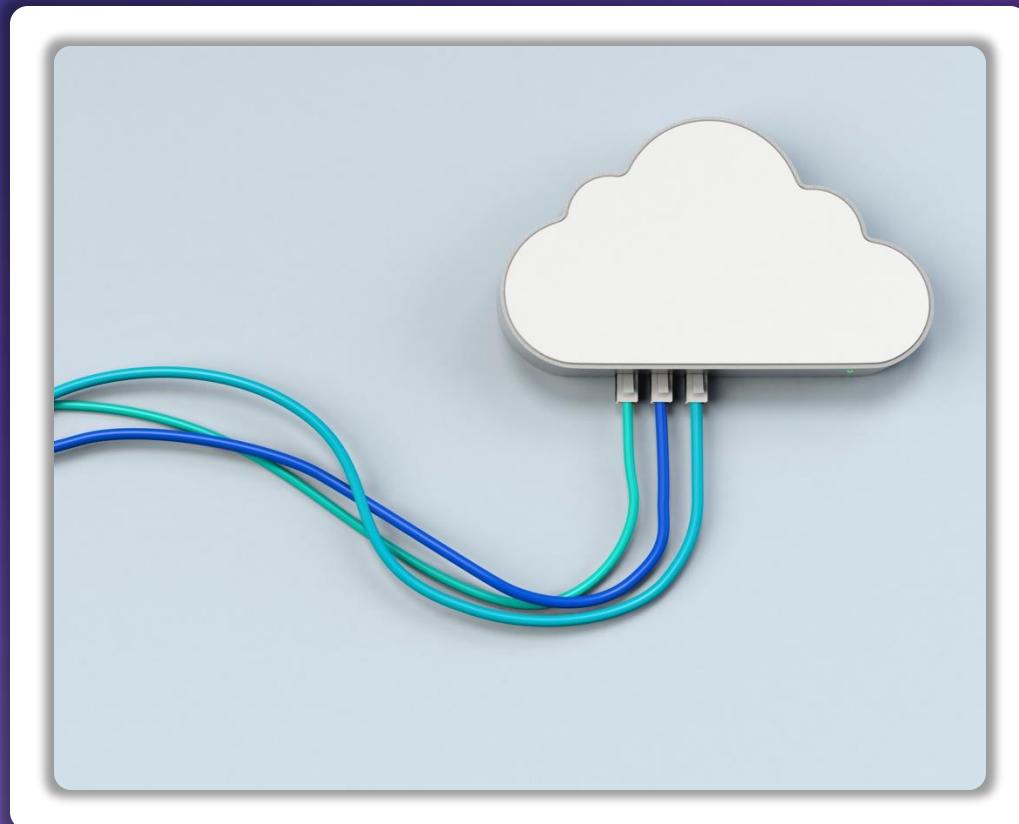


Terraform Providers



Cloud Platforms





Provider Block

The provider **block** is used to extend Terraform's functionality. It allows you to connect Terraform to any API-driven platform, such as cloud platforms, SaaS applications, on-premises hardware solutions, and more.

Provider blocks define where and what Terraform can manage, such as AWS for cloud resources or GitHub for repositories.



Provider Block

Plugins that extend the functionality of Terraform and enable it to interact with cloud providers, SaaS providers, and other API-supported apps.

- **Resources are Implemented by a Provider**

Without providers, Terraform would not be able to provision or manage any infrastructure.

- **Who Publishes Providers?**

HashiCorp maintains some providers, but some are managed by partners and the Terraform community.

- **Developed Separately From Terraform Core**

This means functionality can be added/changed per version.

A screenshot of a code editor window titled "providers.tf". The window shows a single provider block for AWS:

```
provider "aws" {  
  region  = "us-east-2"  
  profile = "prd-workload"  
}  
  
provider "azurerm" {  
  features {}  
  
  tenant_id      = "tenant-id"  
  subscription_id = "sub-id"  
  client_id       = "client-id"  
  client_secret   = "client-secret"  
}
```



Breaking Down the Provider Block

```
providers.tf
provider "aws" {
    region  = "us-east-2"
    profile = "prd-workload"
}

provider "azurerm" {
    features {}

    tenant_id          = "tenant-id"
    subscription_id   = "sub-id"
    client_id          = "client-id"
    client_secret      = "secret"
}
```

The name of the provider defined by the developer or maintainer

Arguments specific to the provider – found in documentation

Authentication arguments – should be provided via Environment Variables instead due to security concerns

®Copyright Bryan Krausen – DO NOT DISTRIBUTE



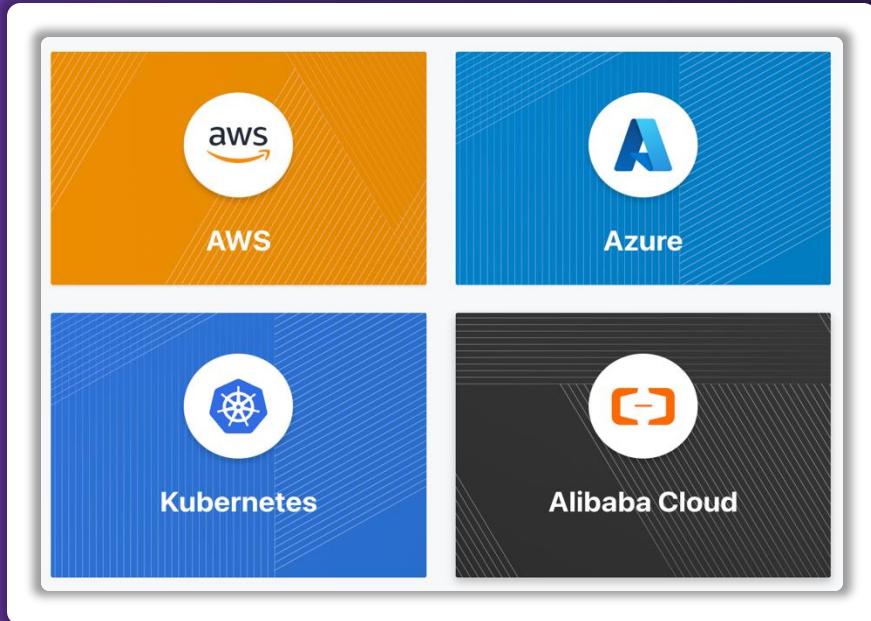
Provider Block

- Some providers need credentials or access tokens to authenticate and manage resources on the target platform (e.g., AWS access keys, GitHub API token)
- A single provider block enables Terraform to manage multiple resources on that platform (e.g., a single AWS provider enables EC2 instances, S3 buckets, EKS clusters)





Provider Block



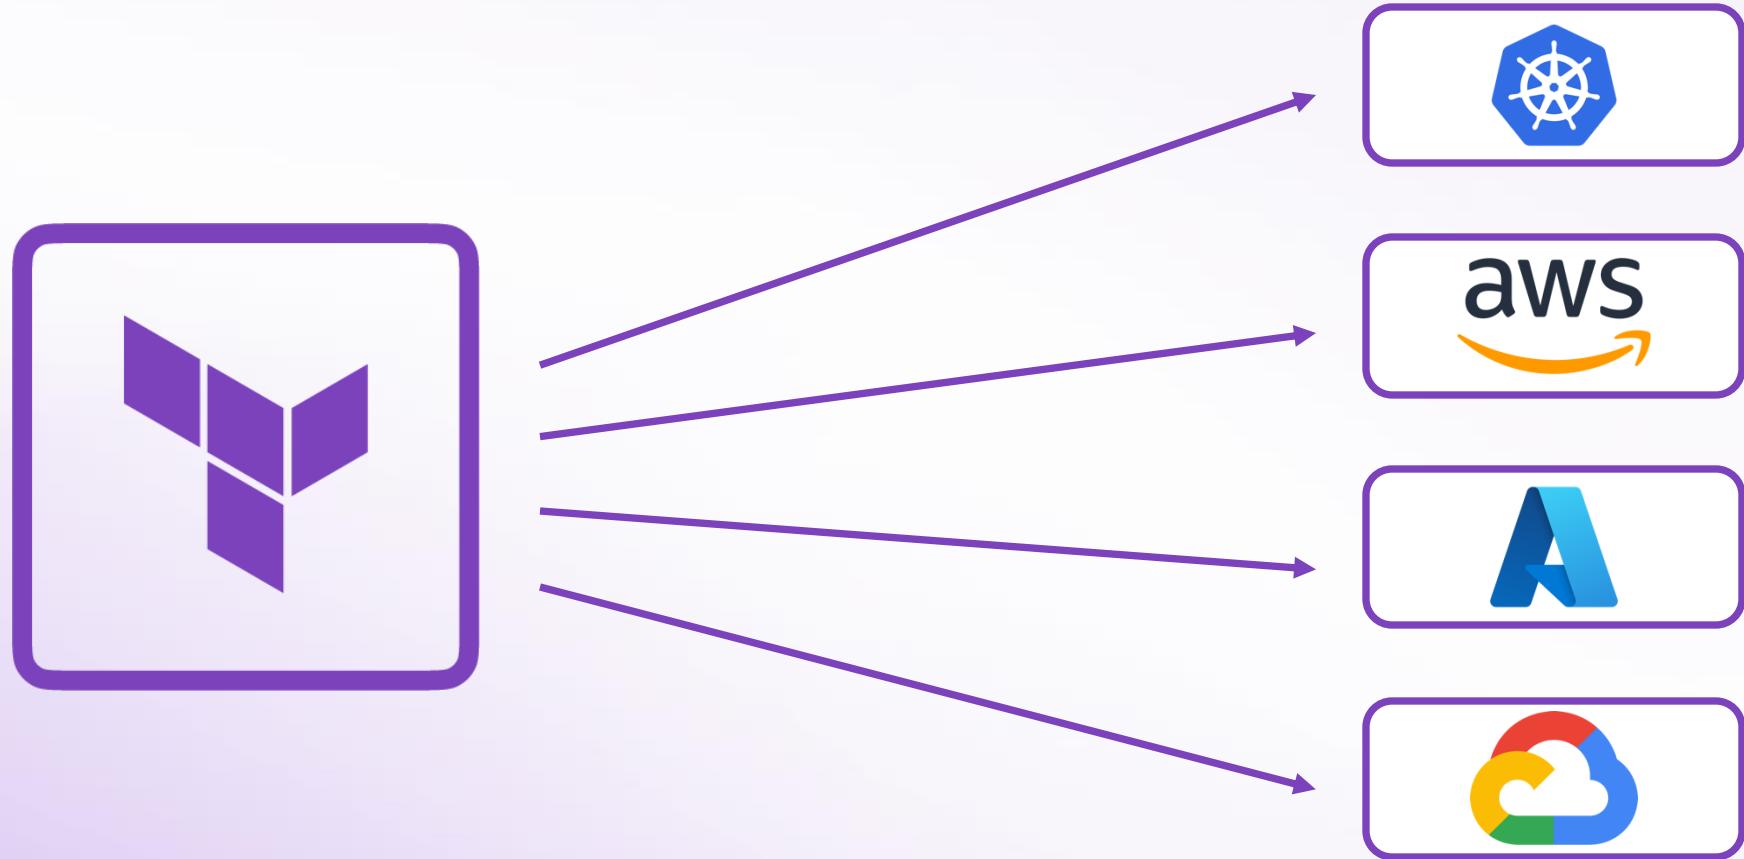
- While providers are stored in GitHub, the documentation can be found on the Terraform Registry → registry.terraform.io
- When declared in your configuration, the provider is downloaded to the machine executing Terraform

LECTURE

Resource Block



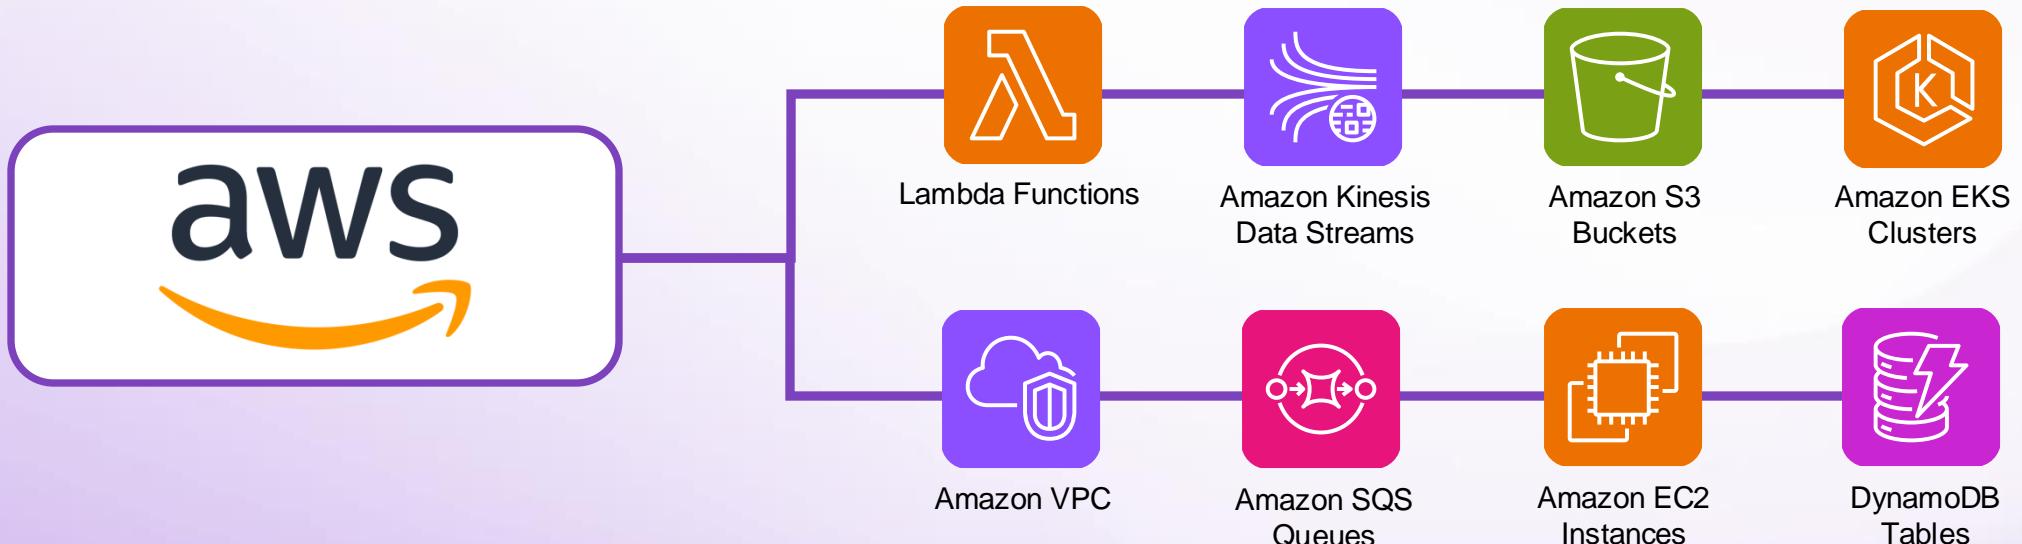
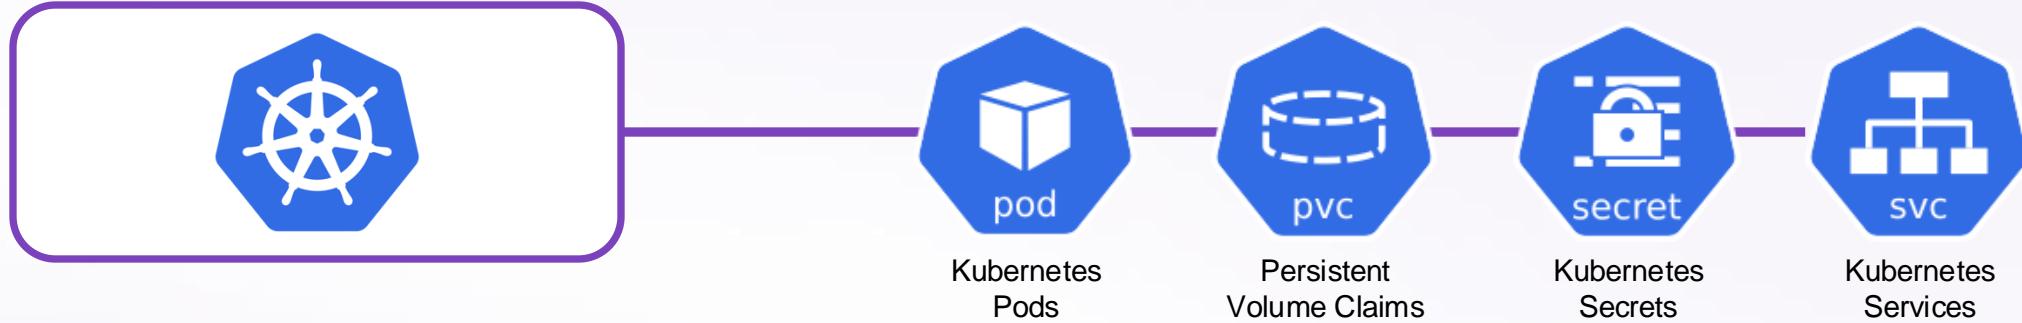
Terraform Resources



®Copyright Bryan Krausen – DO NOT DISTRIBUTE



Terraform Resources





Resource Block

The **resource block** is the core element in Terraform. It defines the specific infrastructure pieces you want to create, manage, or modify, such as virtual machines, databases, or networks.

Resource blocks use a resource type (e.g., `aws_instance`) and a unique name. It acts as a blueprint with the settings you need to configure each resource.



Resource Block

A resource block tells Terraform what to build and how it should be configured.

What Type of Resource is Needed?



The resource type identifies the type of resource that you need. Value is determined by the provider depending on the desired resource.

How Should the Resource be Configured?



Each resource block includes settings (or arguments) that configure its behavior.

Resource Identifiers



Each resource block is identified by a **type** and a **name**, which makes it easy to reference in other parts of your configuration

```
resource "aws_instance" "web" {  
    ami           = "ami-012345"  
    instance_type = "t2.micro"  
  
    key_name     = "prd-web-key"  
    subnet_id   = "subnet_12345abc"  
  
    tags = {  
        Name = "prd-web-svr_01"  
    }  
}  
  
resource "aws_instance" "db" {  
    instance_type = "m5.4xlarge"  
    ami           = "ami-0c55b159"  
}
```



Breaking Down the Resource Block

```
resource "aws_instance" "web" {
    ami             = "ami-012345"
    instance_type  = "t2.micro"

    key_name       = "prd-web-key"
    subnet_id      = "subnet_12345abc"

    tags = {
        Name = "prd-web-svr_01"
    }
}

resource "aws_instance" "db" {
    instance_type = "m5.4xlarge"
    ami           = "ami-0c55b159"
}
```

The type of resource to be created/managed - defined within the provider

The resource name – user-provided string - must be unique per configuration.

Arguments specific to the resource type
Can be found in provider documentation

A second resource block within the same Terraform configuration file

Resource Block



Resource Arguments

```
resource "aws_instance" "web" {  
    ami           = "ami-012345"  
    instance_type = "t2.micro"  
  
    key_name     = "prd-web-key"  
    subnet_id   = "subnet_12345abc"  
  
    tags = {  
        Name = "prd-web-svr_01"  
    }  
}  
  
resource "aws_instance" "db" {  
    instance_type = "m5.4xlarge"  
    ami           = "ami-0c55b159"  
}
```

The Image to Use

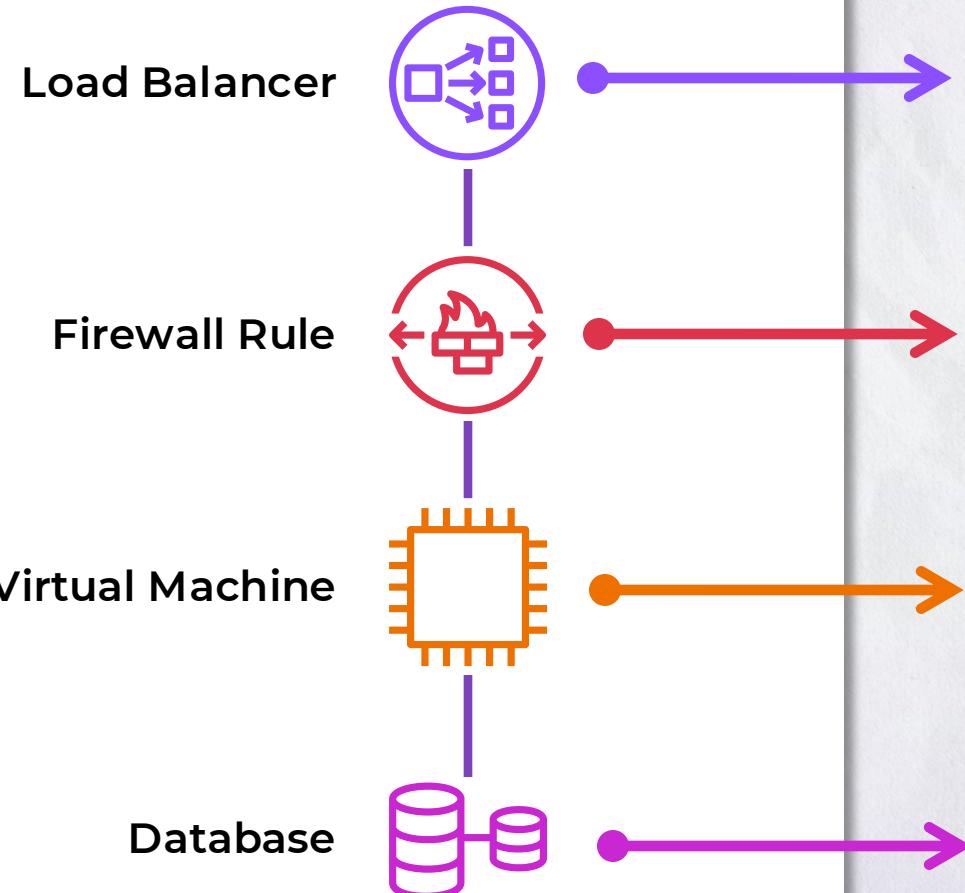
The Size of the Instance

Key Name for SSH Access

The Subnet used to Provision the Instance

Add Tags for the Resource

Resource Block



```
web_app.tf
```

```
resource "aws_lb" "public" {
  name        = "prd-web-lb"
  load_balancer_type = "network"
}

resource "fortios_firewall_policy" "db" {
  action      = "accept"
  name        = "allow_web_443"
}

resource "aws_instance" "web" {
  instance_type = "t3.large"
  ami          = "ami-0c55b159"
}

resource "aws_db_instance" "prd-db" {
  engine      = "mysql"
  instance_class = "db.t3.large"
}
```

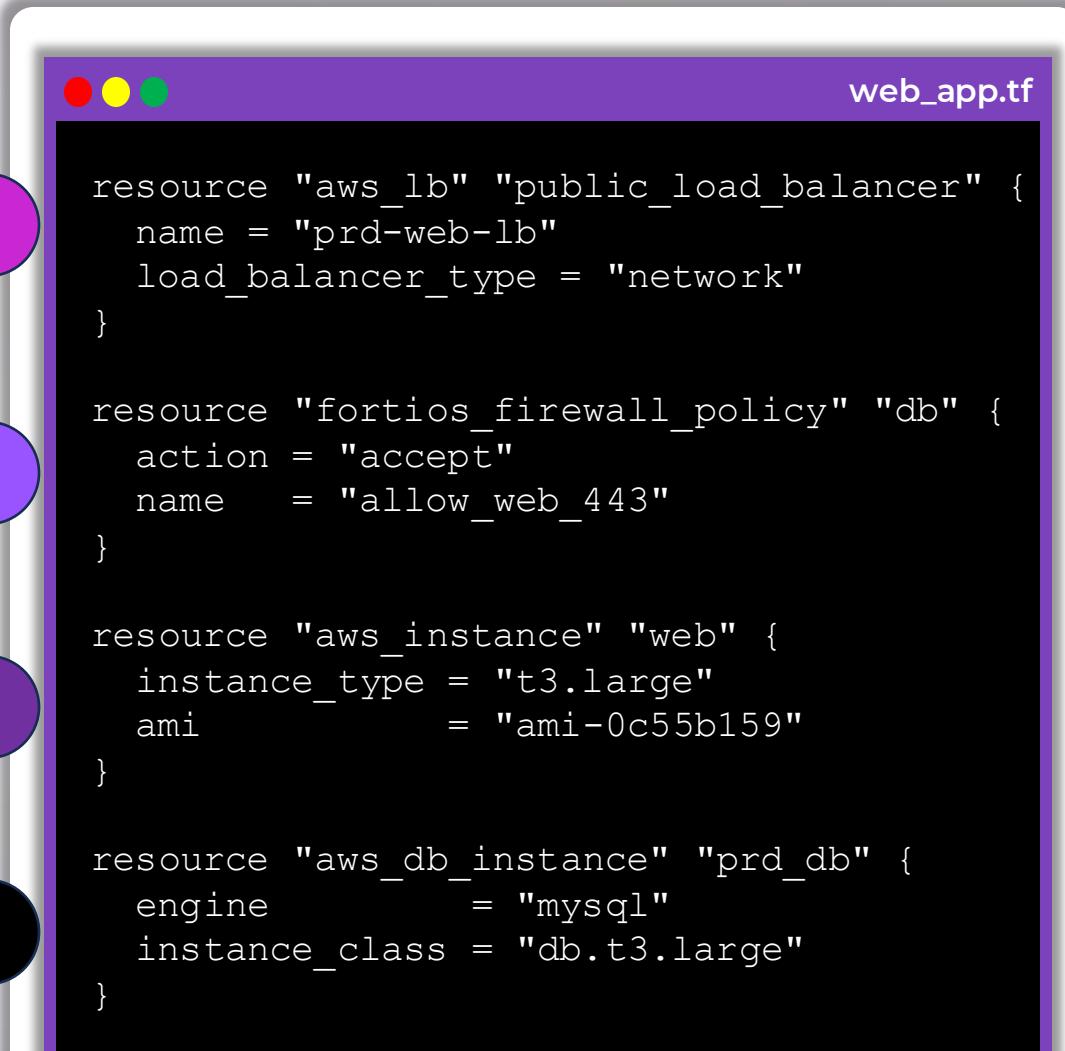


Resource Referencing

Dynamically pull information from one resource block to another, creating a connected configuration

Reference a Resource block by using the resource **type** and the **name** of the resource.

1. **aws_lb.public_load_balancer**
2. **fortios_firewall_policy.db**
3. **aws_instance.web**
4. **aws_db_instance.prd_db**



The screenshot shows a Mac OS X window titled "web_app.tf" containing Terraform configuration code. The code defines four resources: a public load balancer, a firewall policy, a web instance, and a database instance. The resources are numbered 1 through 4, corresponding to the items listed on the left.

```
resource "aws_lb" "public_load_balancer" {
  name = "prd-web-lb"
  load_balancer_type = "network"
}

resource "fortios_firewall_policy" "db" {
  action = "accept"
  name   = "allow_web_443"
}

resource "aws_instance" "web" {
  instance_type = "t3.large"
  ami           = "ami-0c55b159"
}

resource "aws_db_instance" "prd_db" {
  engine        = "mysql"
  instance_class = "db.t3.large"
}
```

Resource Block



```
main.tf
```

```
resource "github_repository" "prod_repo" {
  name      = "prod-app-xyz-repo"
  visibility = "private"
}

resource "github_branch" "default" {
  repository = ${github_repository.prod_repo.name}
  branch     = "main"
}

resource "github_branch_default" "default" {
  repository = github_repository.prod_repo.name
  branch     = ${github_branch.default.branch}
```

A diagram illustrating the structure of the Terraform code. Three resource blocks are shown as black rectangles with white text. A red box highlights the first block. Blue arrows point from the 'repository' field of the second block to the 'name' field of the first, and from the 'branch' field of the third block to the 'branch' field of the second. The file name 'main.tf' is in the top right corner of the code area.

Resource Block



- Resource names must start with a letter or underscore, and may contain only letters, digits, underscores, and dashes.

LECTURE

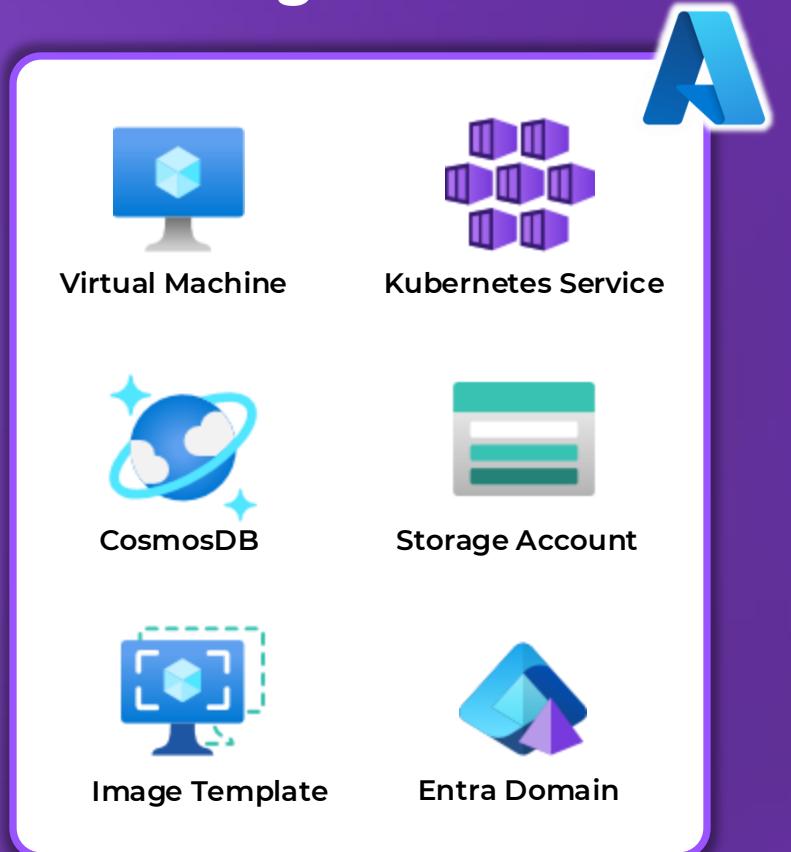
Data Block



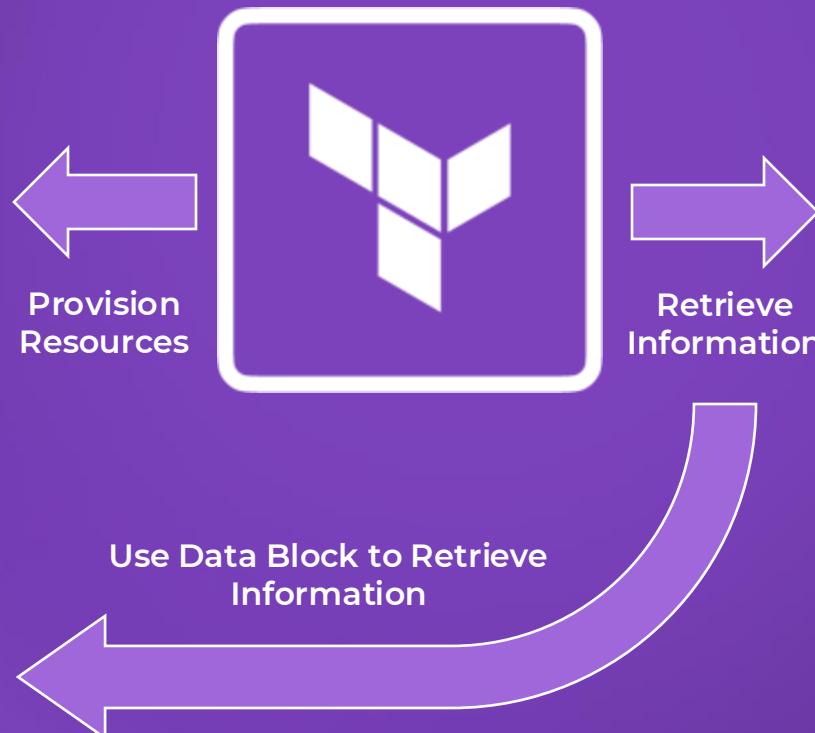


Accessing Existing Resources

Existing Resources

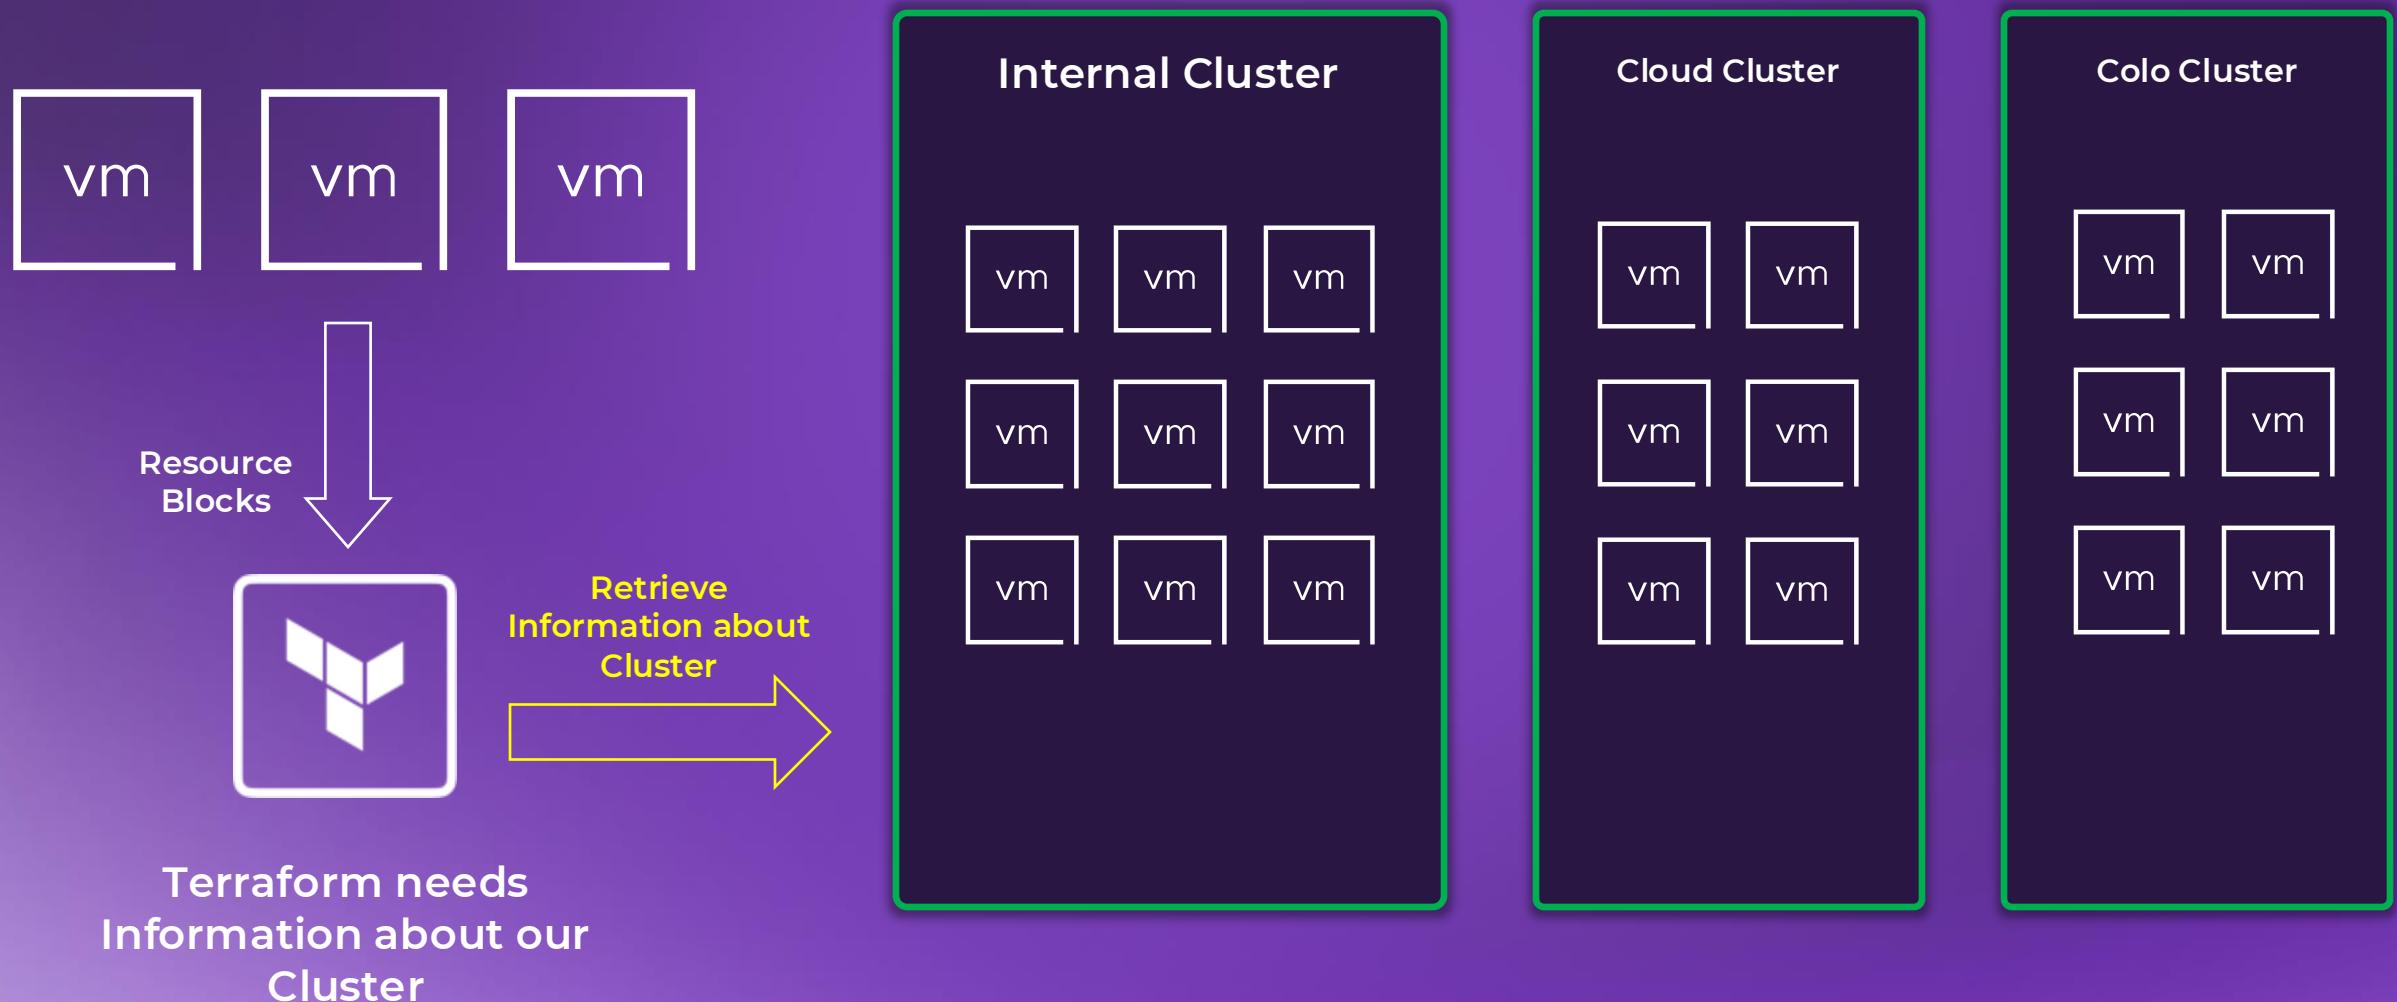


```
resource "azurerm_resource_group" "prd" {  
  name      = "example-resources"  
  location   = "West Europe"  
}  
  
resource "azurerm_virtual_network" "dv" {  
  name          = "example-network"  
  resource_group_name = "azure"  
  location       = "west"  
  address_space    = ["10.0.0.0/16"]  
}  
  
resource "azurerm_mssql_database" "db1" {  
  name      = "example-db"  
  server_id = "server_db"  
  collation = "sql_general"  
}  
  
data "azurerm_kubernetes_cluster" "aks" {  
  name      = "prod-cluster"  
  resource_group_name = "prod-group"  
}
```



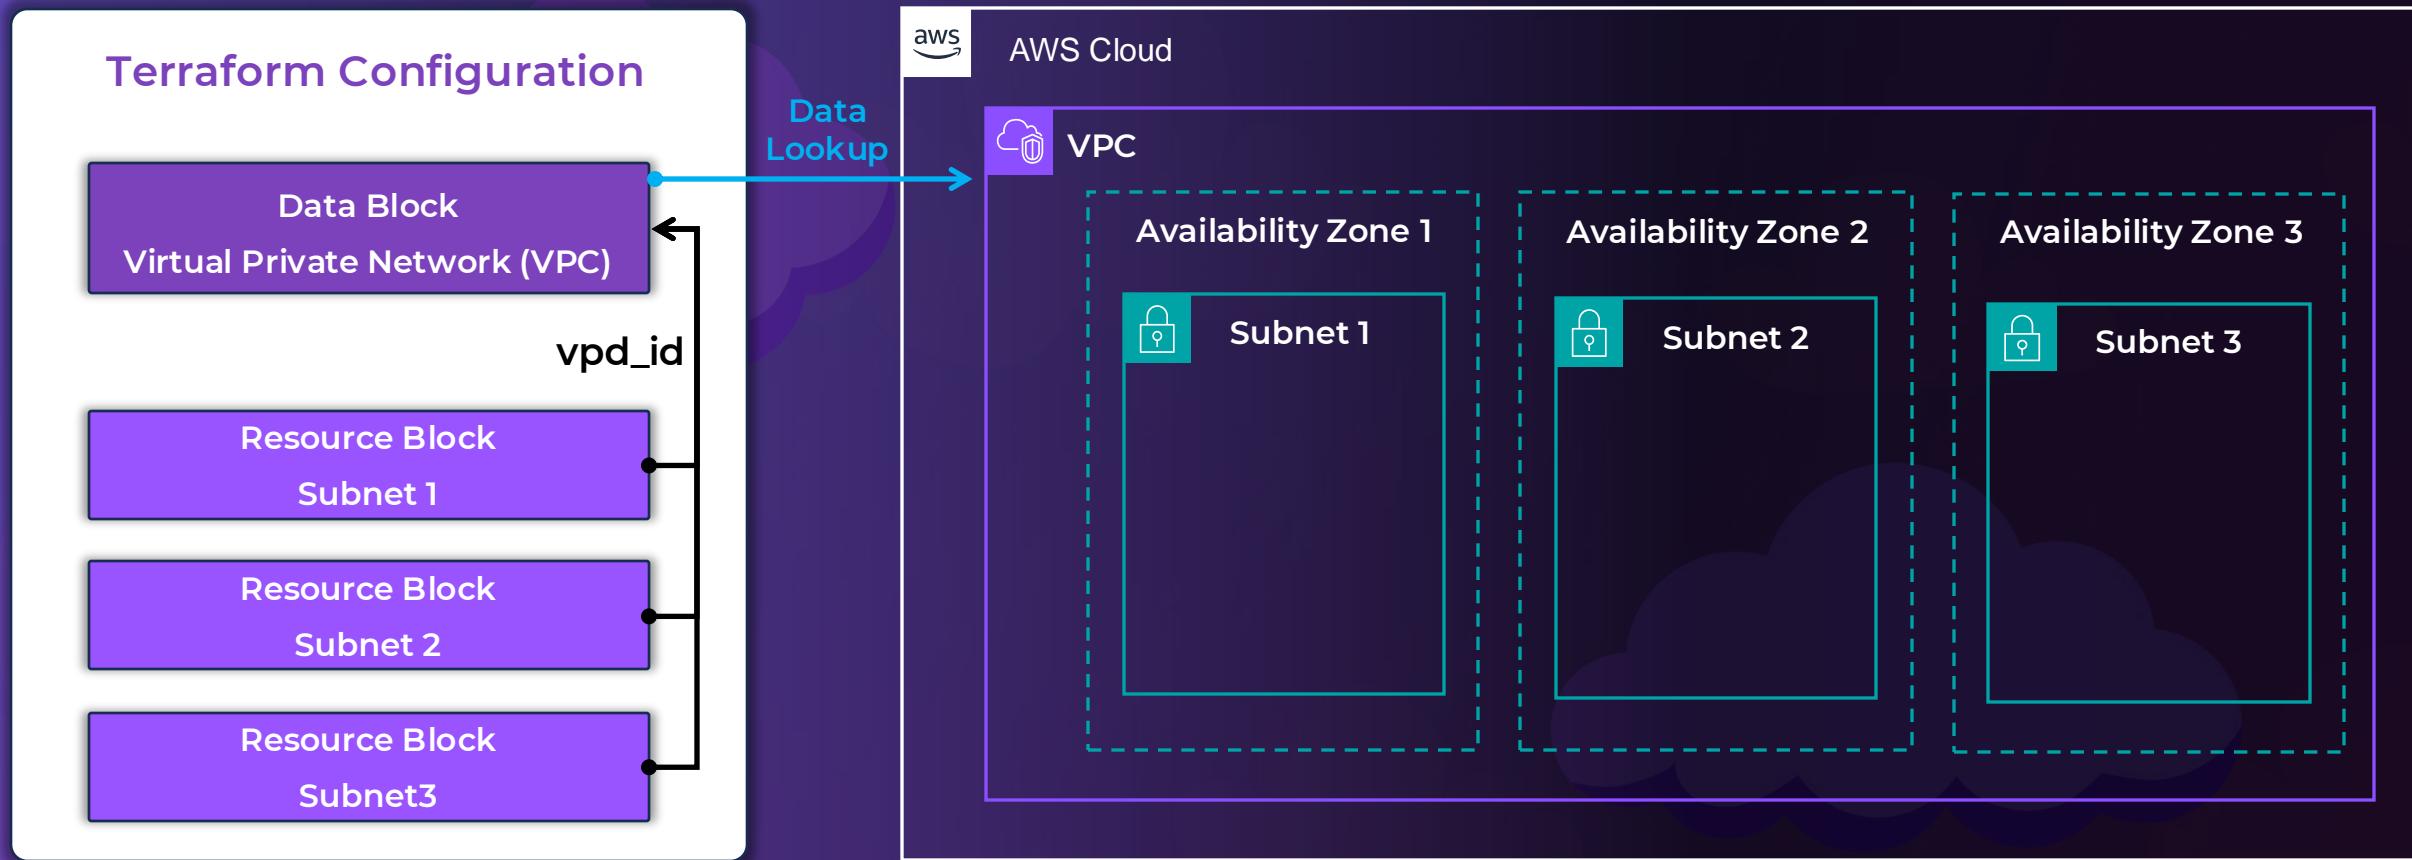


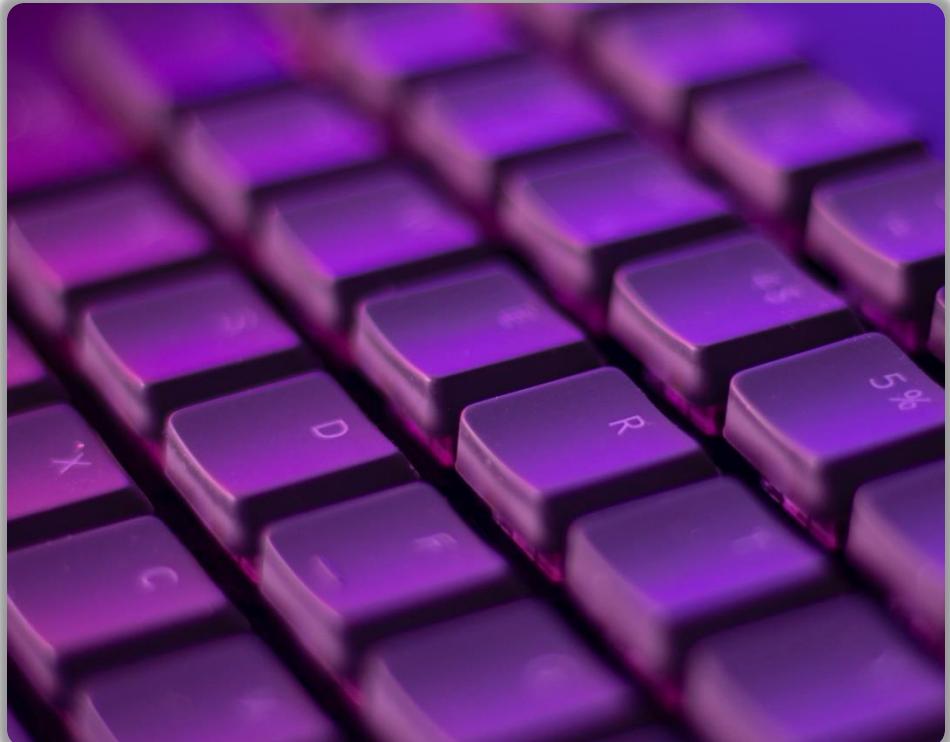
Existing Resources - Example





Existing Resources - Example





Data Block

A **data** block retrieves information about existing resources or external data without creating new infrastructure. It allows you to dynamically reference attributes, such as the ID of an existing subnet or the details of a Kubernetes cluster, making configurations flexible and reusable.

By using **data** blocks, you can integrate Terraform with resources managed outside of its configuration.



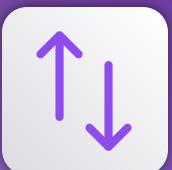
Data Block

Data blocks allow Terraform to access information about existing resources rather than creating new ones.



Access Existing Infrastructure

Retrieve the details of existing resources so they can be referenced in other Terraform blocks



Dependency Resolution

Data blocks help when one resource relies on the properties of another existing resource



Avoid Hardcoding Values

Data blocks prevent the need for hardcoded values by dynamically pulling in details, which makes configurations more adaptable

```
network.tf
```

```
data "aws_vpc" "prd" {
  filter {
    name      = "tag:Name"
    values    = ["prd-vpc"]
  }
}

resource "aws_subnet" "pub" {
  vpc_id      = data.aws_vpc.prd.id
  cidr_block = "10.0.6.0/24"
  ...
}

output "prd_vpc_subnet" {
  value = data.aws_vpc.prd.cidr_block
}
```



Breaking Down the Data Block

The given name used to reference the data block – user-defined string

The type of the resource to obtain data about - defined by provider

Query constraints defined by the data source to get the data you want

Get data about an Azure Resource Group named dev-resource-group

Get data about the default namespace for a Kubernetes cluster

```
main.tf
data "aws_vpc" "prd" {
  filter {
    name    = "tag:Name"
    values = ["prd-vpc"]
  }
}

data "azurerm_resource_group" "dev" {
  name = "dev-resource-group"
}

data "kubernetes_namespace" "app" {
  metadata {
    name = "customer-app"
  }
}
```



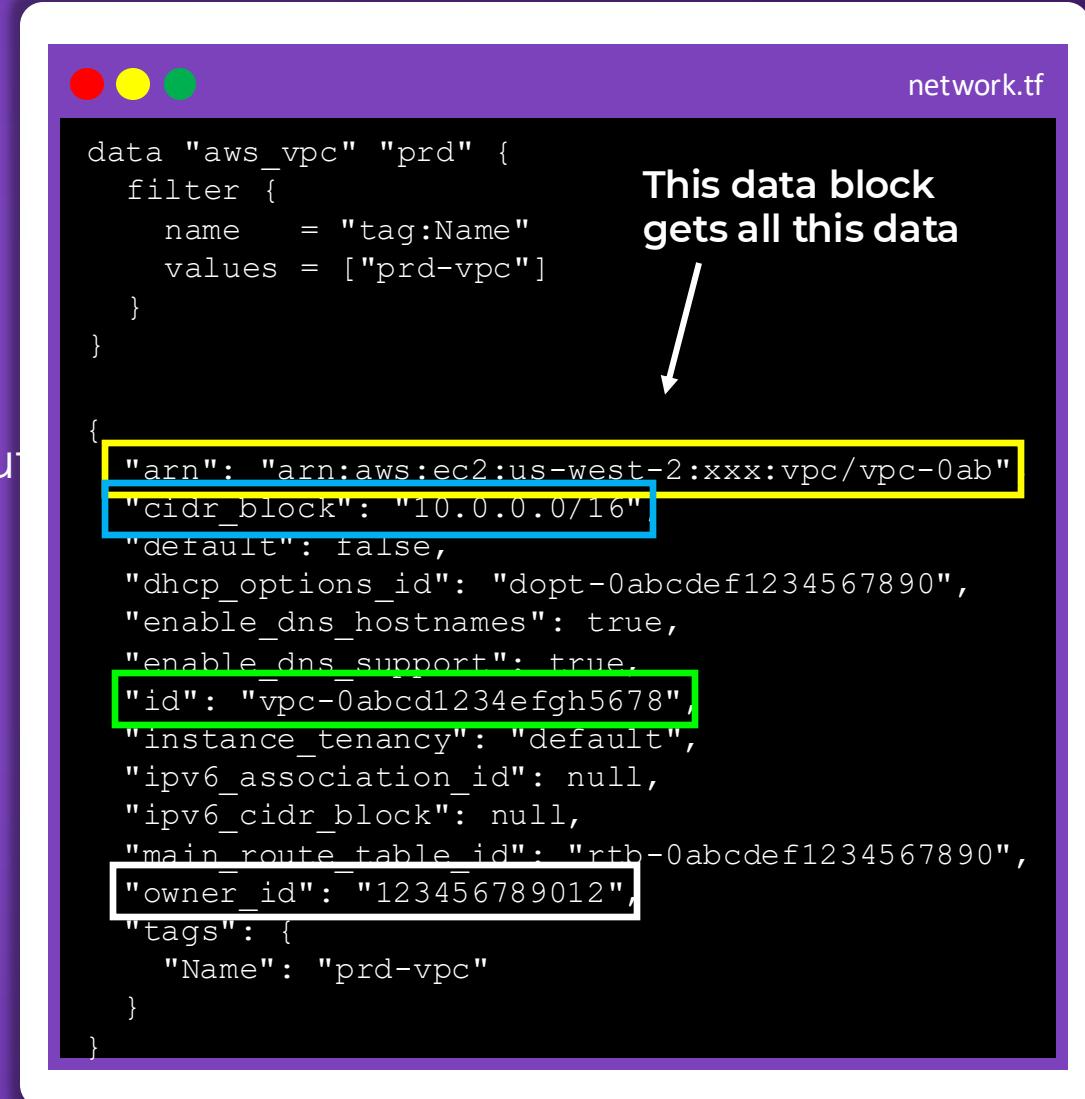
Referencing Data Blocks

Get information from a data block to use elsewhere within the configuration.

data blocks retrieve a wide range of information about the specified resource – called **attributes**.

Data block attributes are referenced just like resource attributes: `data.<type>.<name>.<attribute>`

- `data.aws_vpc.prd.arn`
- `data.aws_vpc.prd.cidr_block`
- `data.aws_vpc.prd.id`
- `data.aws_vpc.prd.owner_id`



```
network.tf
```

```
data "aws_vpc" "prd" {
  filter {
    name   = "tag:Name"
    values = ["prd-vpc"]
  }
}

{
  "arn": "arn:aws:ec2:us-west-2:xxx:vpc/vpc-0ab",
  "cidr_block": "10.0.0.0/16",
  "default": false,
  "dhcp_options_id": "dopt-0abcdef1234567890",
  "enable_dns_hostnames": true,
  "enable_dns_support": true,
  "id": "vpc-0abcd1234efgh5678",
  "instance_tenancy": "default",
  "ipv6_association_id": null,
  "ipv6_cidr_block": null,
  "main_route_table_id": "rtb-0abcdef1234567890",
  "owner_id": "123456789012",
  "tags": {
    "Name": "prd-vpc"
  }
}
```

This data block gets all this data

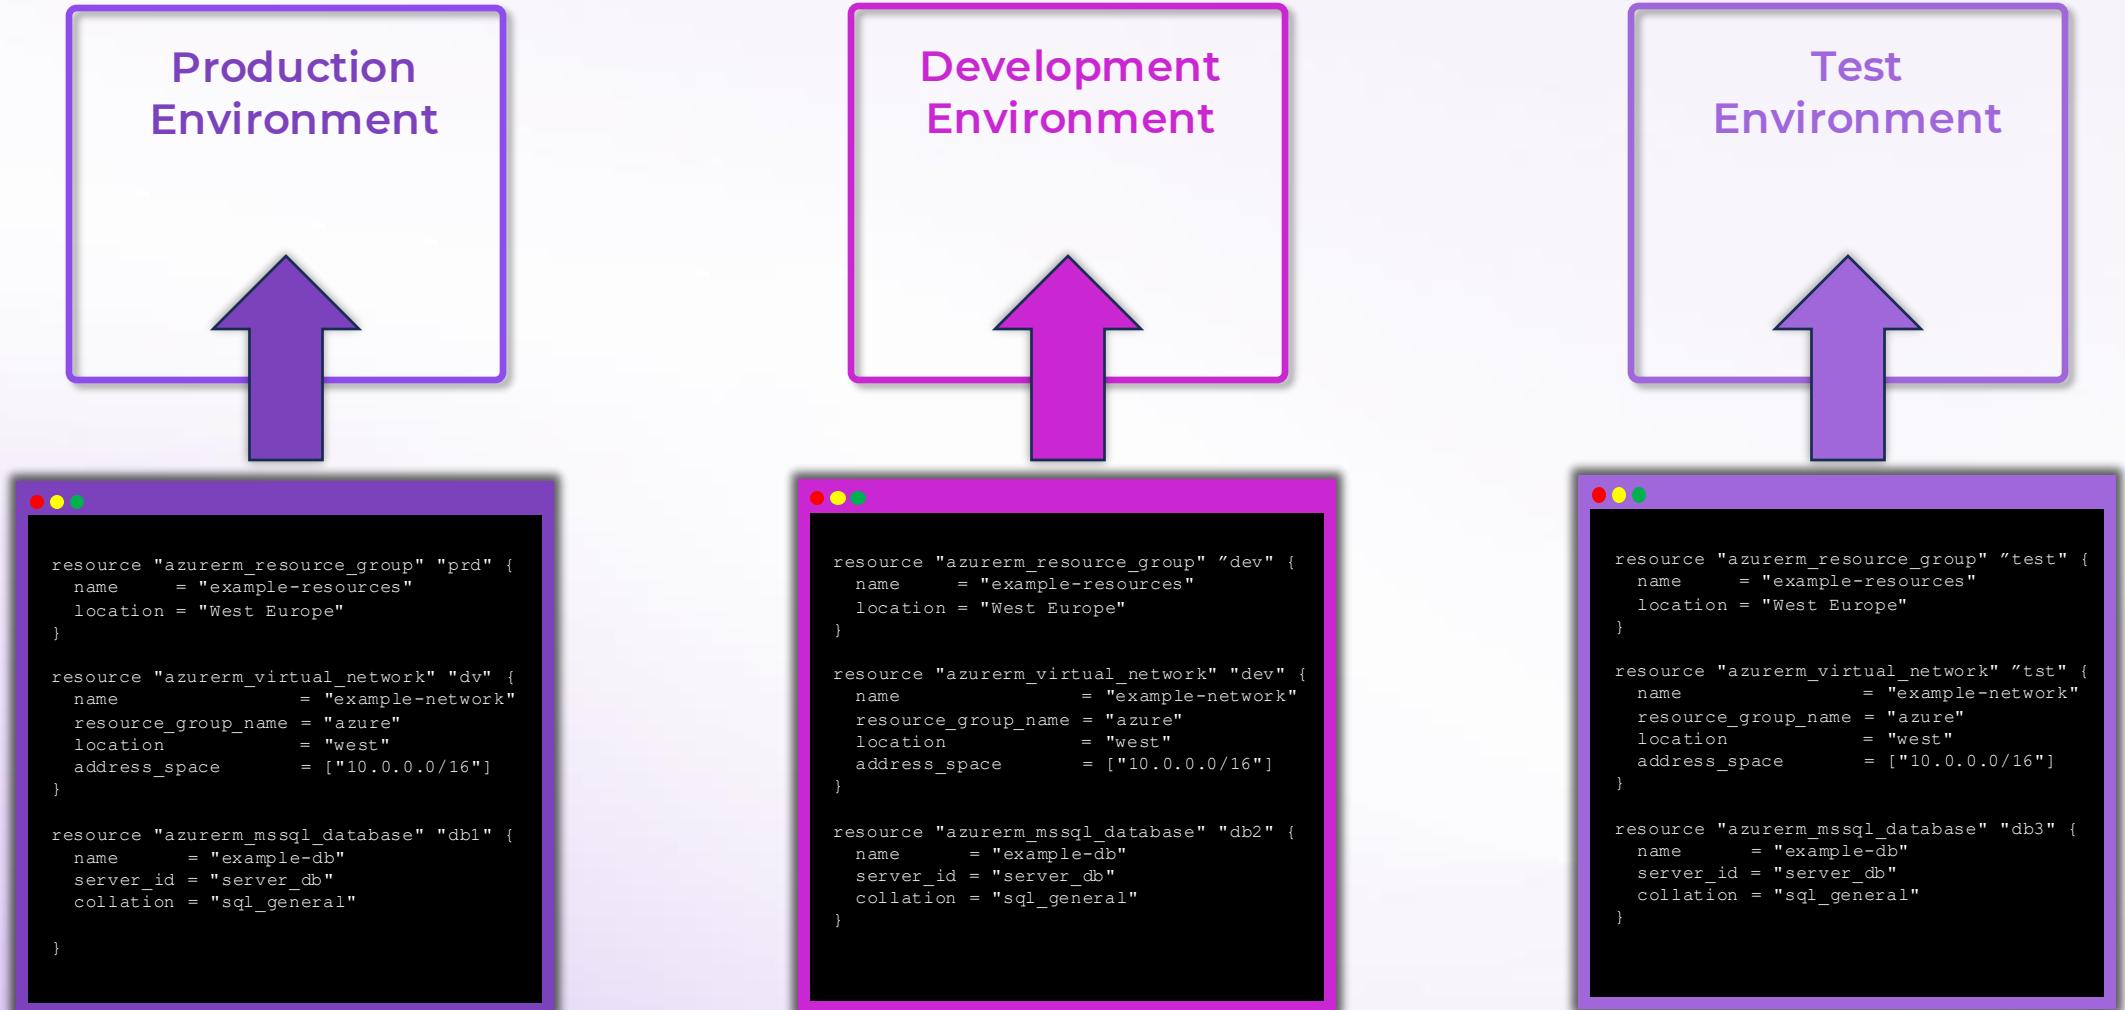
LECTURE

Variable Block



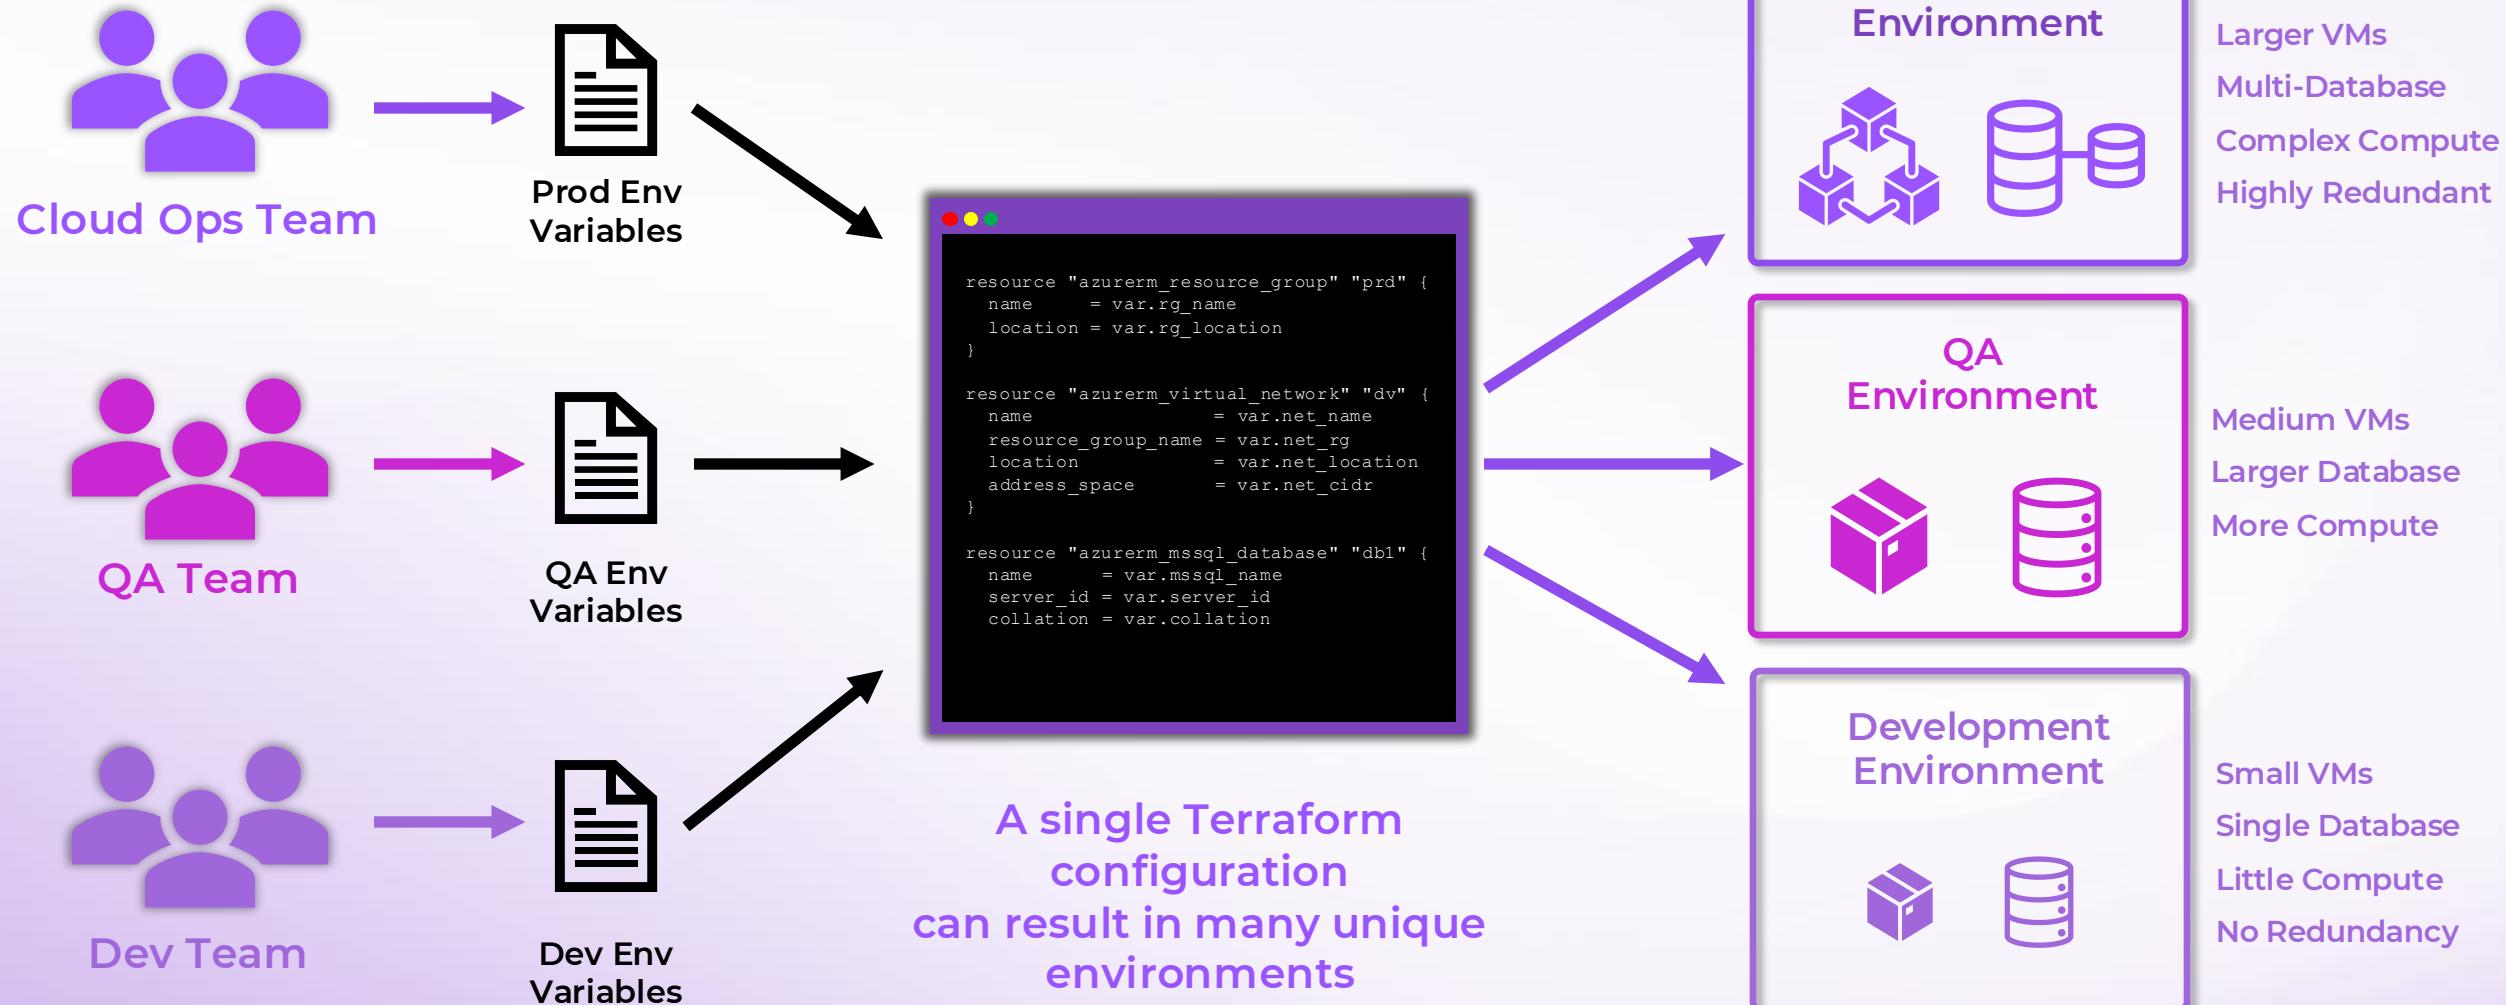


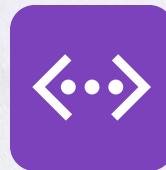
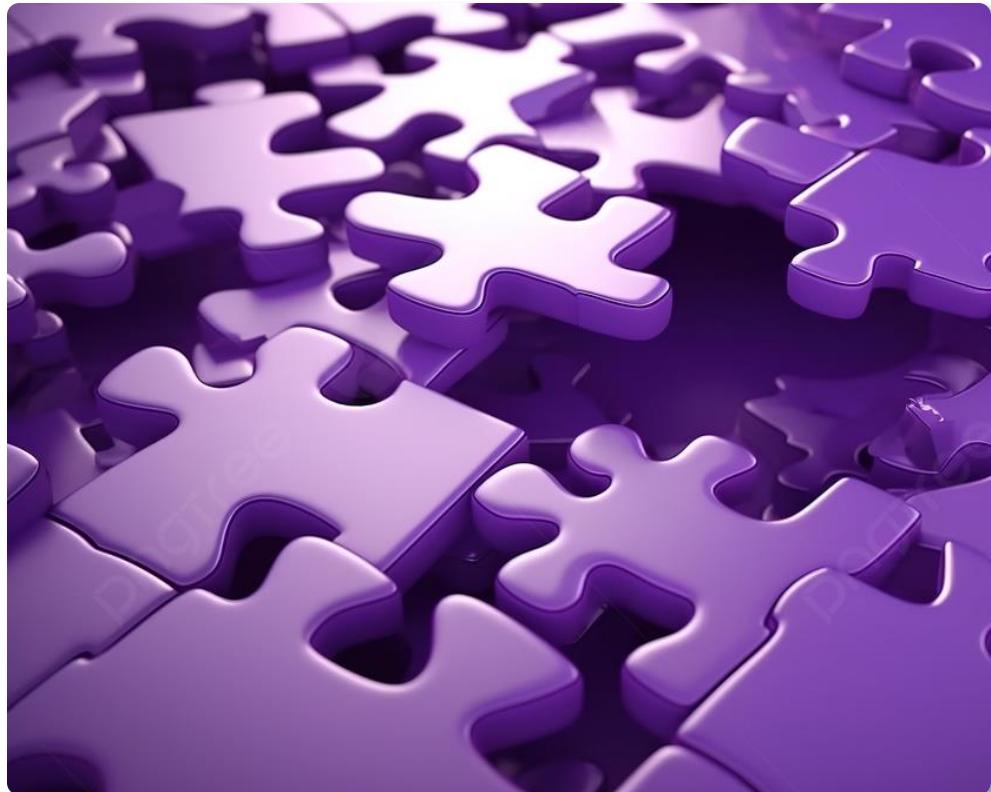
The Problem: Hardcoding Values in Terraform





The Solution: Using Variables in Terraform





Variable Block

A **variable** block is used to define inputs that can be reused and customized, making your configurations more flexible and dynamic. Variables allow you to parameterize values instead of hardcoding them directly into resource blocks.

This approach improves reusability, simplifies managing multiple environments, and helps maintain cleaner, more maintainable configurations.



Variable Block

Define an input that can be reused and customized to make configurations more flexible and dynamic



Dynamic Inputs for Dynamic Code

Variables allow configurations to adapt by using different values for different environments or scenarios.



Centralized Management of Values

By defining variables in one place, you can update values easily without modifying multiple resource blocks.



Improved Reusability of Code

Variable blocks make Terraform configurations modular and reusable across projects and teams.

```
variable "vsphere_datacenter" {  
    description = "Name of datacenter"  
    type        = string  
    default     = "prd-workload-dc"  
}  
  
variable "vsphere_networks" {  
    description = "List of networks"  
    type        = list(string)  
  
    default     = [  
        "VM Network",  
        "Management Network"  
    ]  
}
```



Breaking Down the Variable Block

```
variable 'vsphere_datacenter' {
  description = "Name of datacenter"
  type        = string
  default     = "prd-workload-dc"
}

variable "vsphere_networks" {
  description = "List of networks"
  type        = list(string)

  default = [
    "VM Network",
    "Management Network"
  ]
}
```

- The name for the variable – a user-provided string

- Friendly description of the variable and what it's used for...

- Define what value types are accepted for the variable

- The default value of the variable if not set anywhere else in Terraform



Types of Variables

Primitive Variable Types

ABC

String

a sequence of characters
representing some text



```
1 variable "region" {  
2   type    = string  
3   default = "us-east-2"  
4 }
```

9

Number

a numeric value



```
1 variable "num_of_vms" {  
2   type    = number  
3   default = 3  
4 }
```

T/F

Bool

a boolean value – either True
or False



```
1 variable "enable_ha" {  
2   type    = bool  
3   default = true  
4 }
```



Types of Variables

Complex Variable Types



List

A sequence of similar values

Presented by a pair of square brackets [] containing a comma-separated sequence of values

Note: lists start at 0



```
1 variable "permitted_size" {  
2   type  = list  
3   default = ["t3.small", "t4g.micro"]  
4 }
```

index = 0

index = 1



Types of Variables

Complex Variable Types

Map

A group of values identified by named labels

Presented by a pair of curly braces {} and contains a series of key=value pairs

```
● ● ●  
1 variable "course_details" {  
2   description = "Course details"  
3   type        = map(string)  
4   default     = {  
5     instructor = "bryan"  
6     course     = "terraform"  
7   }  
8 }
```

var.course_details.instructor

var.course_details.course



Types of Variables

Complex Variable Types



Set

A collection of guaranteed unique values

Terraform does not support accessing elements of a set using an index because sets are unordered collections.

Note: Convert to a list to access



```
1 variable "pub_subnet_ids" {  
2   type    = set(string)  
3   default = ["subnet-12345", "subnet-67890"]  
4 }
```



Resource Referencing

Reference a Variable within the Terraform configuration to use the value

Reference a Variable by using the **name** of the variable.

1. **var.vsphere_datacenter**
2. **var.vsphere_networks[0]**

The screenshot shows a Mac OS X window with a white title bar containing the standard red, yellow, and green buttons. The main area of the window displays a Terraform configuration file with two numbered callouts. Callout 1 points to the first variable definition:

```
variable "vsphere_datacenter" " {
```

Callout 2 points to the second variable definition:

```
variable "vsphere_networks" {
```

Both definitions include a description, type, and default value. The configuration file continues with a default list of networks.

```
description = "Name of datacenter"
type      = string
default   = "prd-workload-dc"
}

description = "List of networks"
type      = list(string)

default = [
  "VM Network",
  "Management Network"
]
```

Assigning Values to Variables

After variables are declared in your Terraform configuration, the values can be set in many different ways...

Variable Defaults



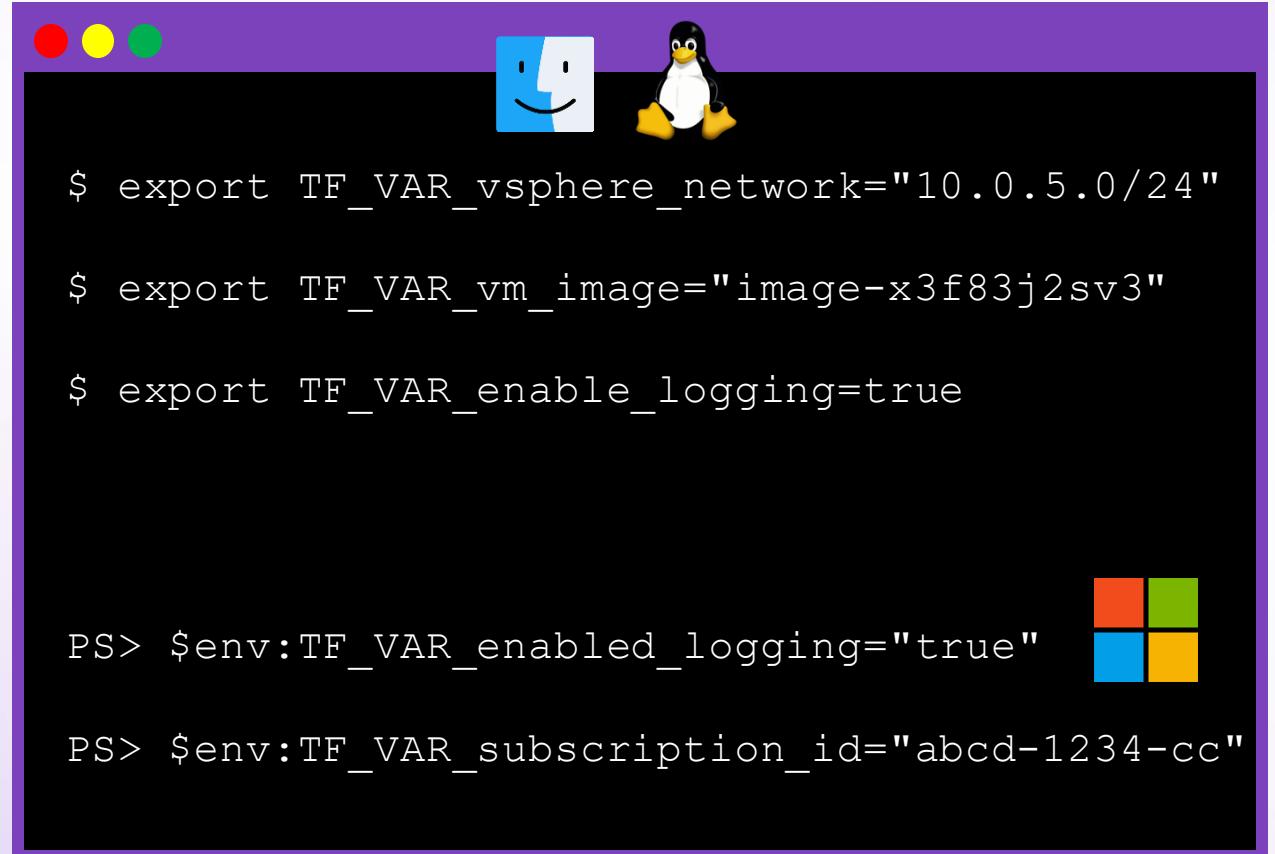
Set a value directly in
the variable block
declaration

```
● ● ●  
1 variable "pub_subnet_ids" {  
2   type    = set(string)  
3   default = ["subnet-12345", "subnet-67890"]  
4 }
```

Environment Variables



- Use an environment variable to set the value
- Must start with `TF_VAR_` followed by the variable name
- Useful for when you don't want to commit the value to a repo

A terminal window with a black background and a purple header bar. The header bar features the macOS logo, the Linux Tux logo, and the Windows logo. The terminal displays four lines of code:

```
$ export TF_VAR_vsphere_network="10.0.5.0/24"  
$ export TF_VAR_vm_image="image-x3f83j2sv3"  
$ export TF_VAR_enable_logging=true  
  
PS> $env:TF_VAR_enabled_logging="true"  
PS> $env:TF_VAR_subscription_id="abcd-1234-cc"
```

tfvars File



- Y Use a .tfvars file to set variable values
- Y Terraform automatically loads a .tfvars file if they exist
- Y Set values using a key=value format using the variable names
- Y A .tfvars file is usually excluded from version control (.gitignore)

A screenshot of a Mac OS X window titled "terraform.tfvars". The window contains the following configuration code:

```
# VM-level configurations
vsphere_network = "10.0.5.0/24"
vm_image        = "image-x3f83j2sv3"

# Application Configurations
enable_logging = true

# Cloud Account Configurations
subscription_id = "abcd-1234-cc"
```



Command Line Flags

- Pass variable values directly from the command line.
- Use `-var="key=value"` to set the value for a specific variable.
- Command-line `-var` values override both default values and values in `.tfvars` files.
- Useful for quick, temporary changes or testing without modifying files.



TERMINAL

```
$ terraform plan -var="enable_logging=true"

$ terraform apply -var="region=us-west-2" -var="vm_image=image-x3f83j2sv3"
```



Order of Precedence

What hierarchy does Terraform follow when resolving variable values from different sources?



Flexibility for Values

Understanding the order of precedence ensures the right value is applied



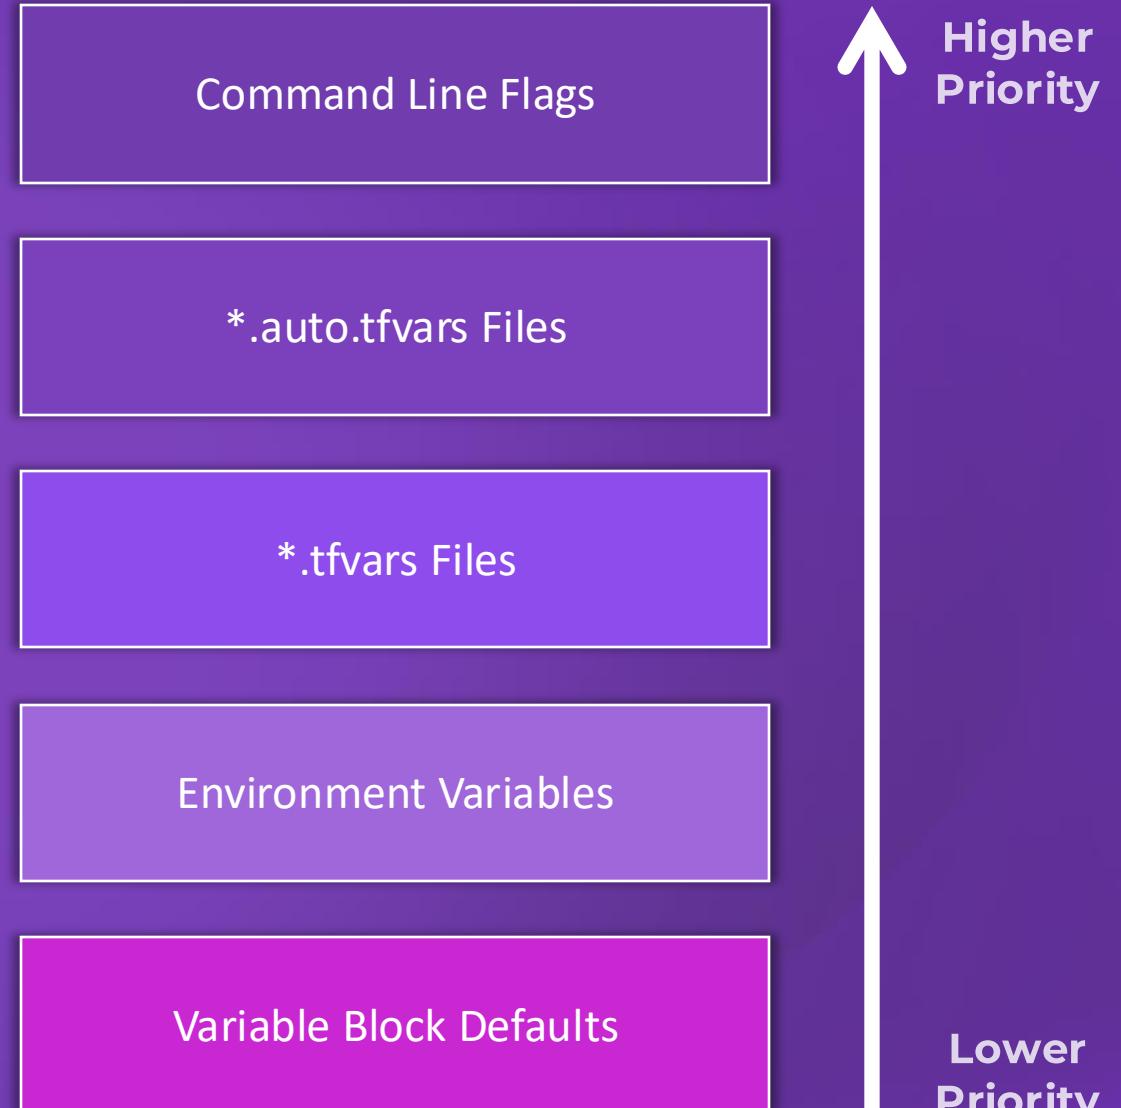
Clear Override Rules

Higher-priority sources override lower-priority defaults, preventing unexpected behavior.



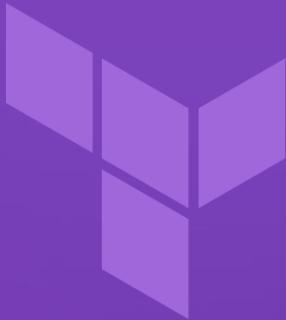
Adaptability Across Environments

Precedence allows you to manage variables for different environments without changing configurations.



LECTURE

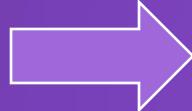
Output Block





Accessing New Resources

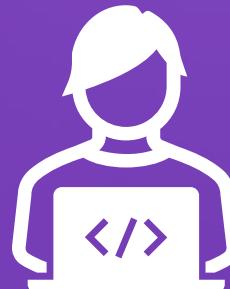
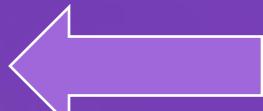
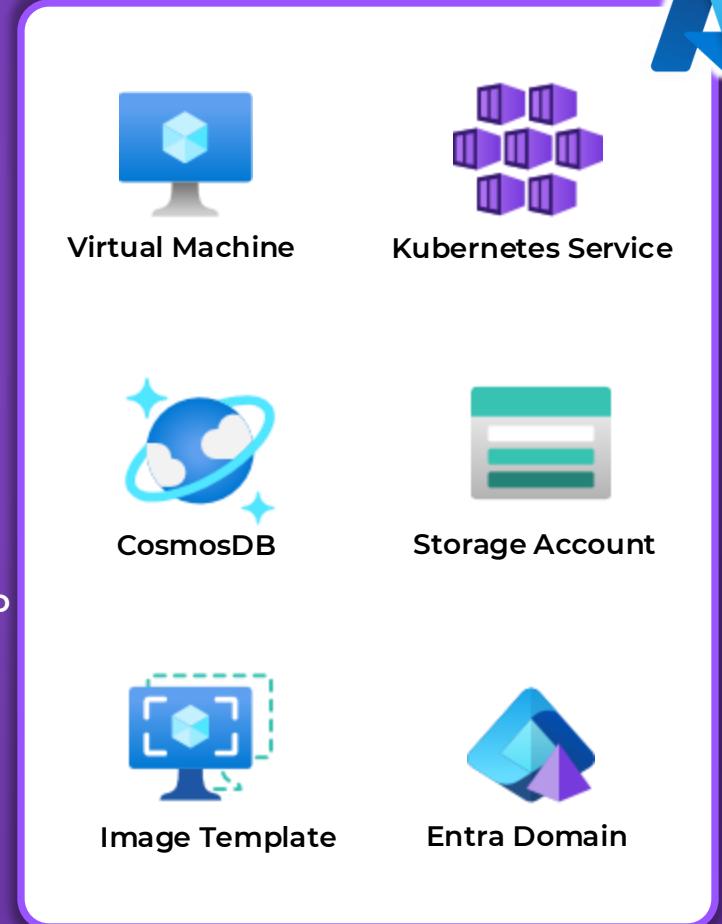
```
resource "azurerm_resource_group" "prd" {  
    name      = "example-resources"  
    location  = "West Europe"  
}  
  
resource "azurerm_virtual_network" "dv" {  
    name          = "example-network"  
    resource_group_name = "azure"  
    location      = "west"  
    address_space  = ["10.0.0.0/16"]  
}
```



Write
Code



Provision
Resources



Output Information to
Access Resources



<https://btk.me/btk>



Output Block

An **output** block in Terraform is used to display key information about your infrastructure after a deployment, such as IP addresses, DNS names, or resource IDs. Outputs make it easy to retrieve and share important details without manually looking them up in the cloud provider's console.

They're also useful for passing data between modules or automating workflows.



Output Block

An output block displays important information about your infrastructure, like resource IDs or URLs, after deployment.

Display Critical Information

Outputs make it easy to retrieve important details like IP addresses, DNS names, or resource IDs after deployment.

Enable Module Integration

Outputs can pass data between Terraform modules, simplifying workflows and making configs reusable.

Improve Workflow Efficiency

Outputs automate the retrieval of key values, reducing the need for manual lookups in provider consoles.

®Copyright Bryan Krausen – DO NOT DISTRIBUTE



outputs.tf

```
# Output Instance IP Address
output "instance_public_ip" {
  description = "Public IP of Server"
  value        = aws_instance.web.public_ip
}

# Output DNS Name for Load Balancer
output "website_dns" {
  description = "Website DNS Record"
  value        = aws_elb.web_app.dns_name
}

# Output Friendly URL of Website
output "website_url" {
  value = "https://${aws_alb.web.dns_name}"
}
```



Breaking Down the Output Block

```
outputs.tf

# Output Instance IP Address
output "instance_public_ip" {
  description = "Public IP of Web server"
  value        = aws_instance.web.public_ip
}

# Output DNS Name for Load Balancer
output "website_dns" {
  description = "Website DNS Record"
  value       = aws_elb.web_app.dns_name
}

# Output Friendly URL of Website
output "website_url" {
  value = "https://${aws_alb.web.dns_name}"
}
```

The name of the output – user-defined string

The value of the output

Another output block – a more advanced value combining the value of a resource within a string

Lecture

Terraform Block





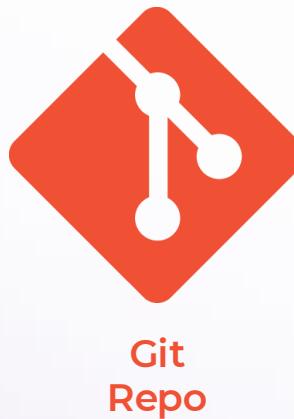
Imagine This Real-World Scenario...

Your Workstation

```
resource "azurerm_resource_group" "prd" {  
    name      = "example-resources"  
    location  = "West Europe"  
}  
  
resource "azurerm_virtual_network" "dv" {  
    name          = "example-network"  
    resource_group_name = "azure"  
    location      = "west"  
    address_space  = ["10.0.0.0/16"]  
}  
  
resource "azurerm_mssql_database" "db1" {  
    name      = "example-db"  
    server_id = "server_db"  
    collation = "sql_general"  
}
```



Commit
Code



Pull
Code

```
resource "azurerm_resource_group" "prd" {  
    name      = "example-resources"  
    location  = "West Europe"  
}  
  
resource "azurerm_virtual_network" "dv" {  
    name          = "example-network"  
    resource_group_name = "azure"  
    location      = "west"  
    address_space  = ["10.0.0.0/16"]  
}  
  
resource "azurerm_mssql_database" "db1" {  
    name      = "example-db"  
    server_id = "server_db"  
    collation = "sql_general"  
}
```



%)#@!!

A purple silhouette of a person with arms outstretched, looking confused or surprised.

Terraform Version
1.10

Colleague's Workstation

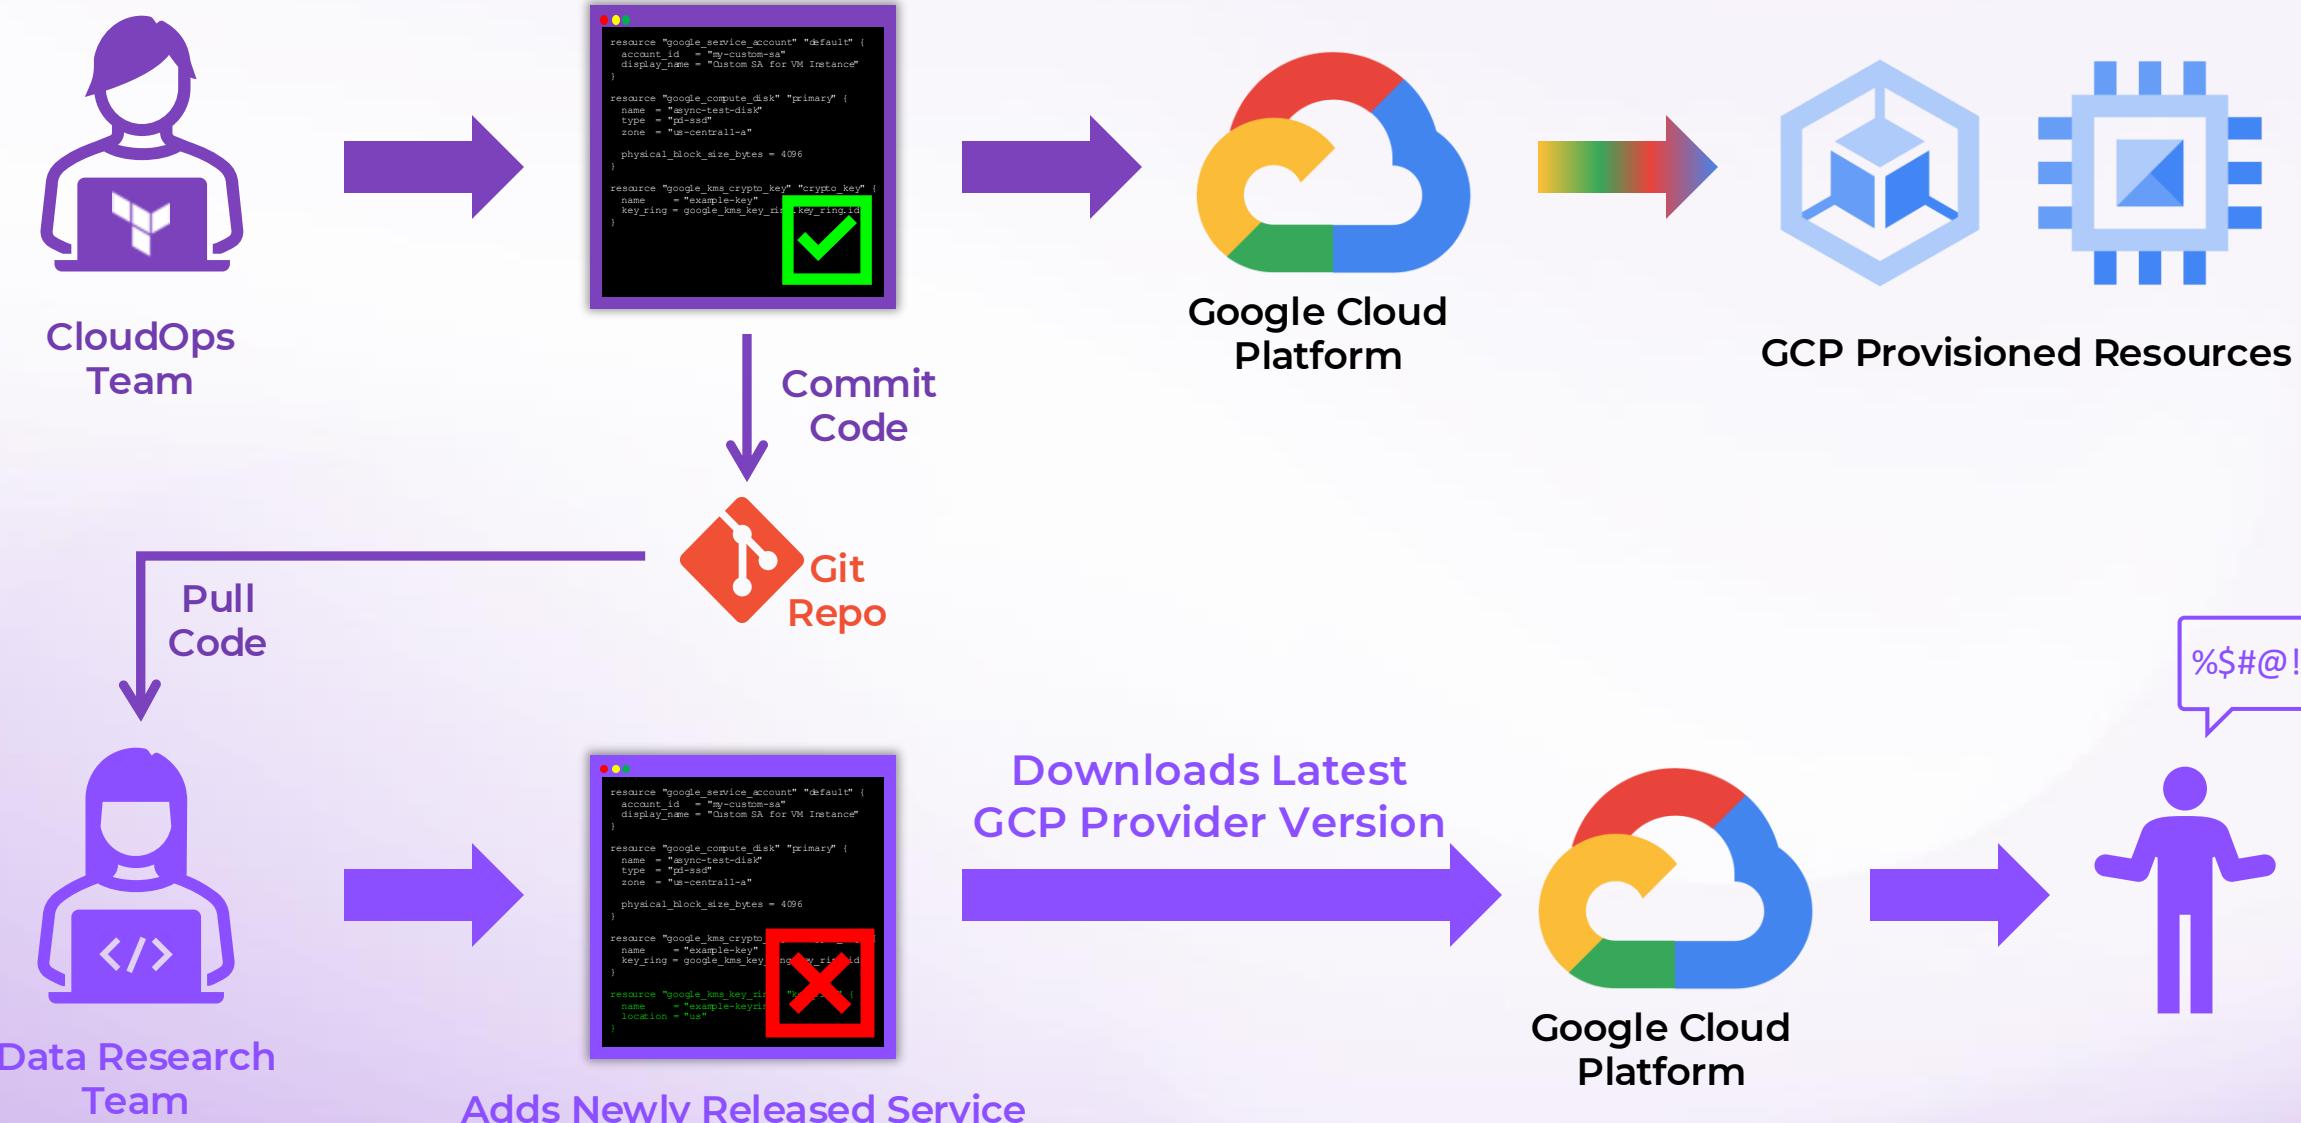
Terraform Version
1.11

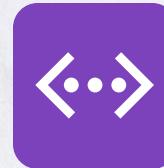
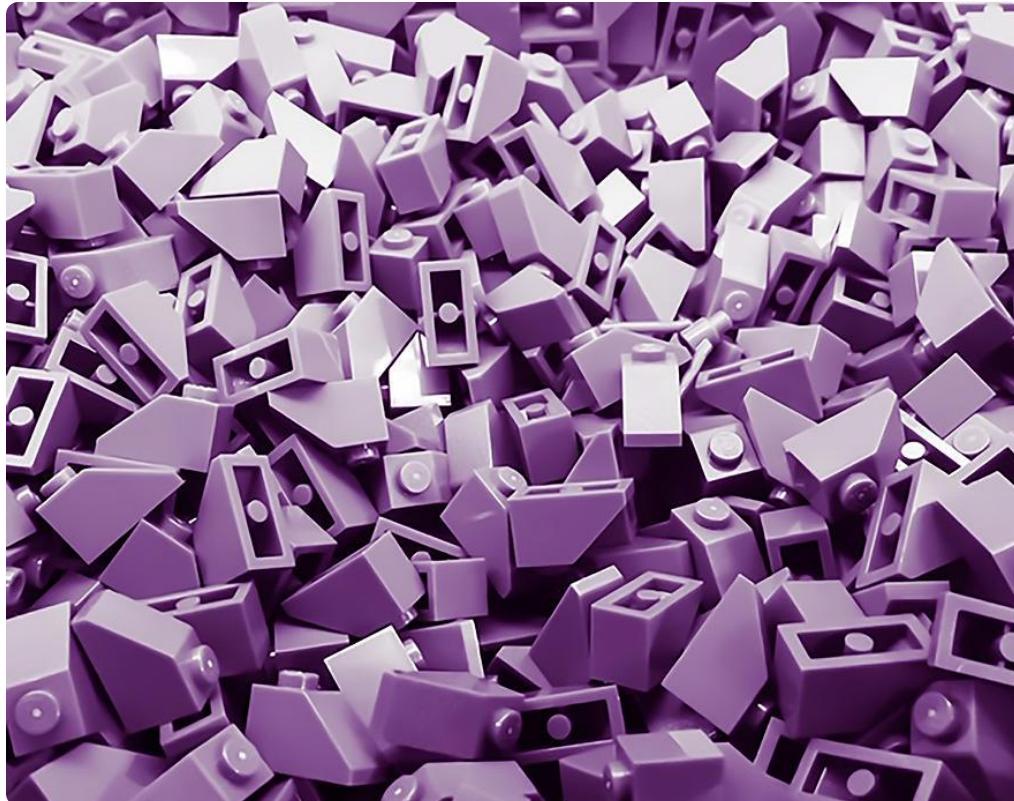
®Copyright Bryan Krausen – DO NOT DISTRIBUTE

HashiCorp Terraform: The Ultimate Beginner's Guide



The Risk of Mismatched Provider Versions





Terraform Block

The **terraform** block sets global configurations for your Terraform project, such as required provider versions, Terraform version constraints, and backend settings for state management. It ensures consistency across teams and environments by enforcing specific versions and configurations.

Using the **terraform** block helps avoid errors and provides a solid foundation for managing infrastructure effectively.



Terraform Block

The `terraform` block ensures your Terraform setup is consistent, and scalable across teams and environments.

Version Control

1.1

Specifies the required Terraform and provider versions to avoid compatibility issues

Backend Configuration



Configures where Terraform state is stored to ensure secure and reliable state management

Provider Requirements



Defines the providers and the versions needed for your infrastructure

```
terraform {  
    required_version = "~> 1.10.0"  
  
    required_providers {  
        aws = {  
            source  = "hashicorp/aws"  
            version = "~> 5.0"  
        }  
    }  
  
    backend "s3" {  
        bucket = "prd-terraform-east-2"  
        key    = "state/terraform.tfstate"  
        region = "us-east-2"  
    }  
}
```



Breaking Down the Terraform Block

```
● ● ●  
terraform {  
    required_version = "~> 1.10.0"  
  
    required_providers {  
        aws = {  
            source = "hashicorp/aws"  
            version = "5.87.0"  
        }  
    }  
  
    backend "s3" {  
        bucket = "prd-terraform-east-2"  
        key = "state/terraform.tfstate"  
        region = "us-east-2"  
    }  
}
```

Ensures Terraform version 1.10.x is used to run this configuration

Specify a required provider and require version 5.87.0

The backend configuration of where you want to store Terraform state

Version Constraints

```
terraform {  
    required_version = "1.9.8"  
  
    required_providers {  
        aws = {  
            source = "hashicorp/aws"  
            version = "5.77.0"  
        }  
    }  
}
```

- Only Terraform v1.9.8 can be used to run this configuration
- Installs/uses AWS provider version 5.77.0

```
terraform {  
    required_version = ">= 1.9.8"  
  
    required_providers {  
        aws = {  
            source = "hashicorp/aws"  
            version = ">= 5.77.0"  
        }  
    }  
}
```

- Terraform v1.9.8 or greater can be used to run this configuration
- Installs/uses AWS provider version greater or equal to 5.77.0

```
terraform {  
    required_version = "~> 1.11.0"  
  
    required_providers {  
        aws = {  
            source = "hashicorp/aws"  
            version = "~> 5.87.0"  
        }  
    }  
}
```

- Terraform v1.11.x or greater can be used to run this configuration
- Installs/uses AWS provider version greater or equal to 5.87.x

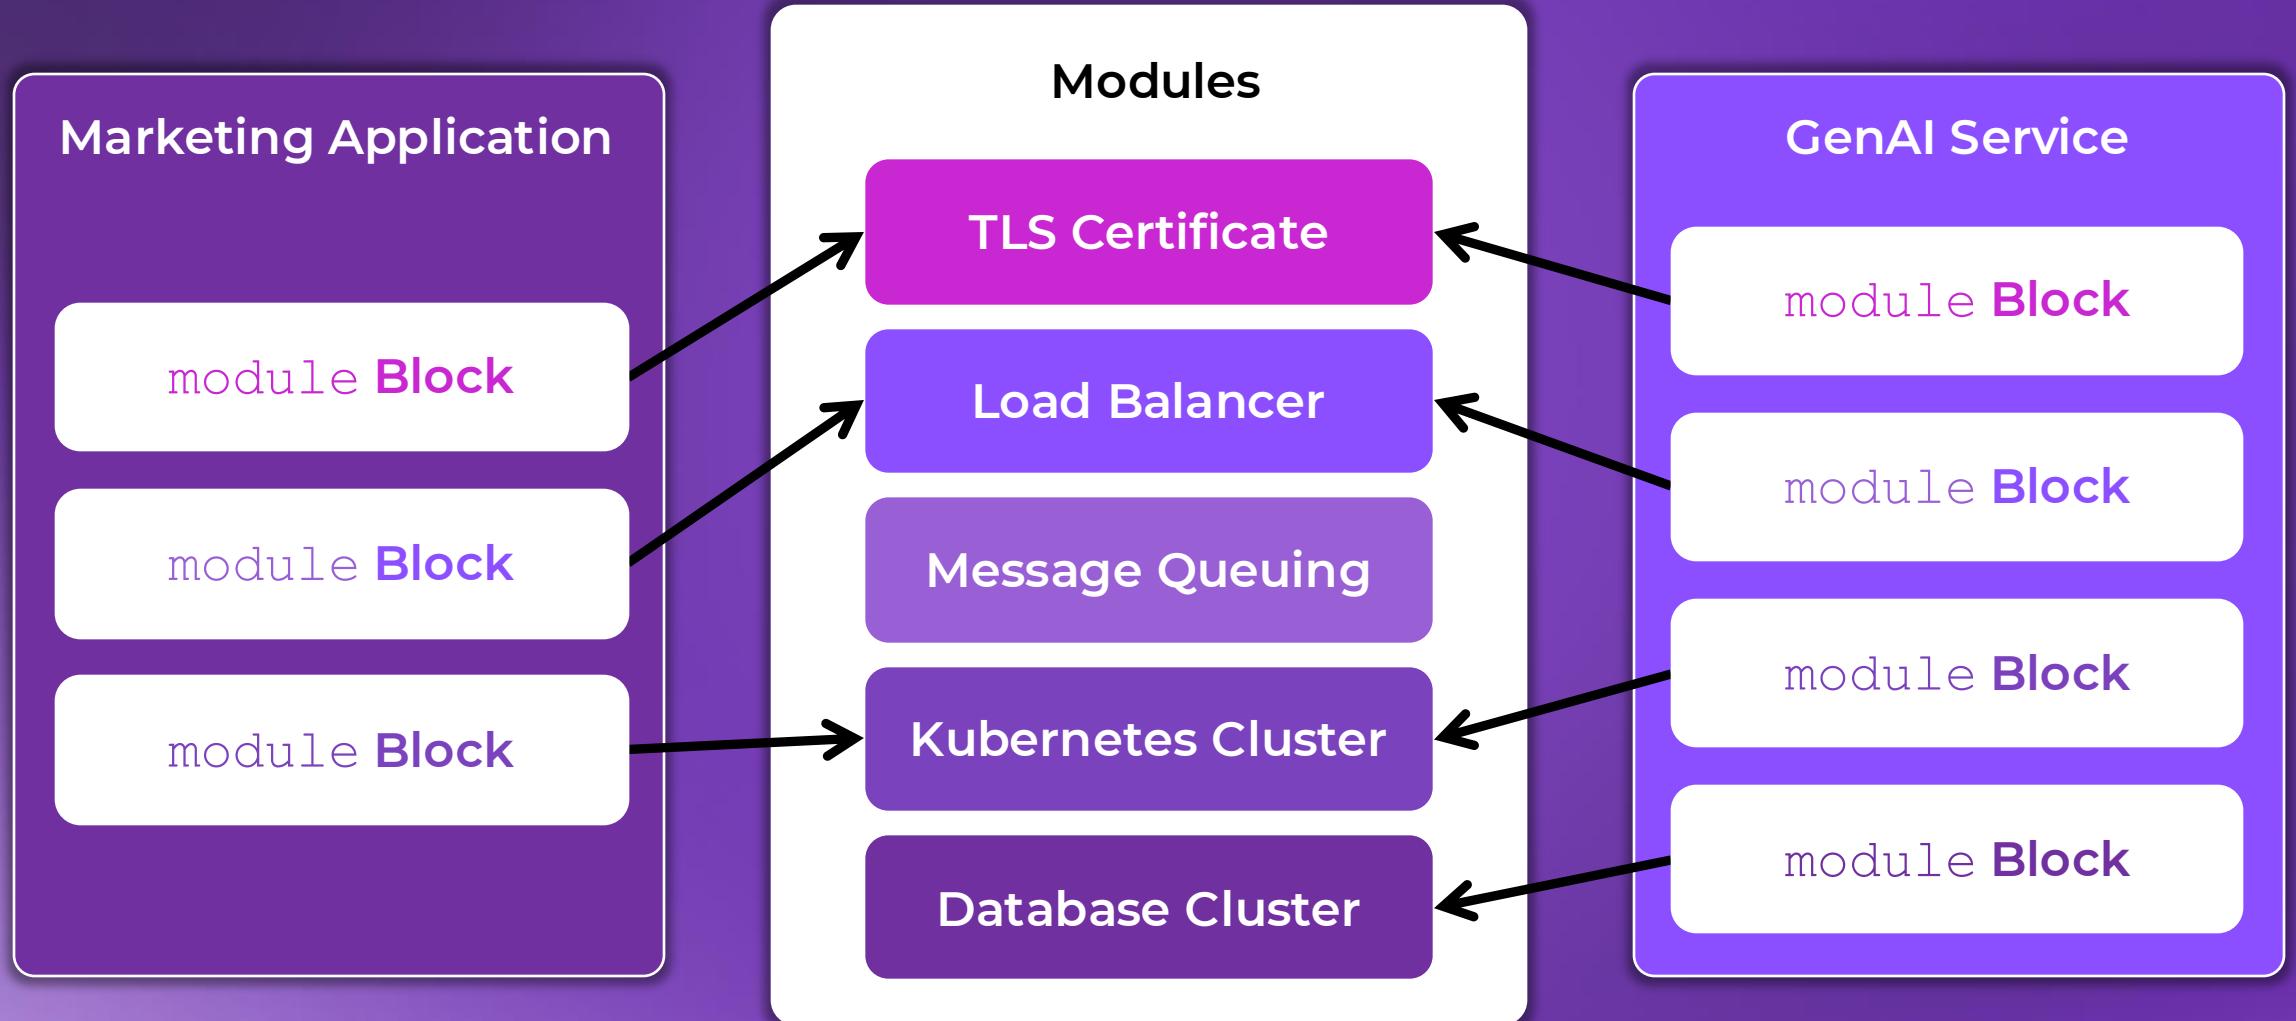
~ allows only the right-most version component to increment

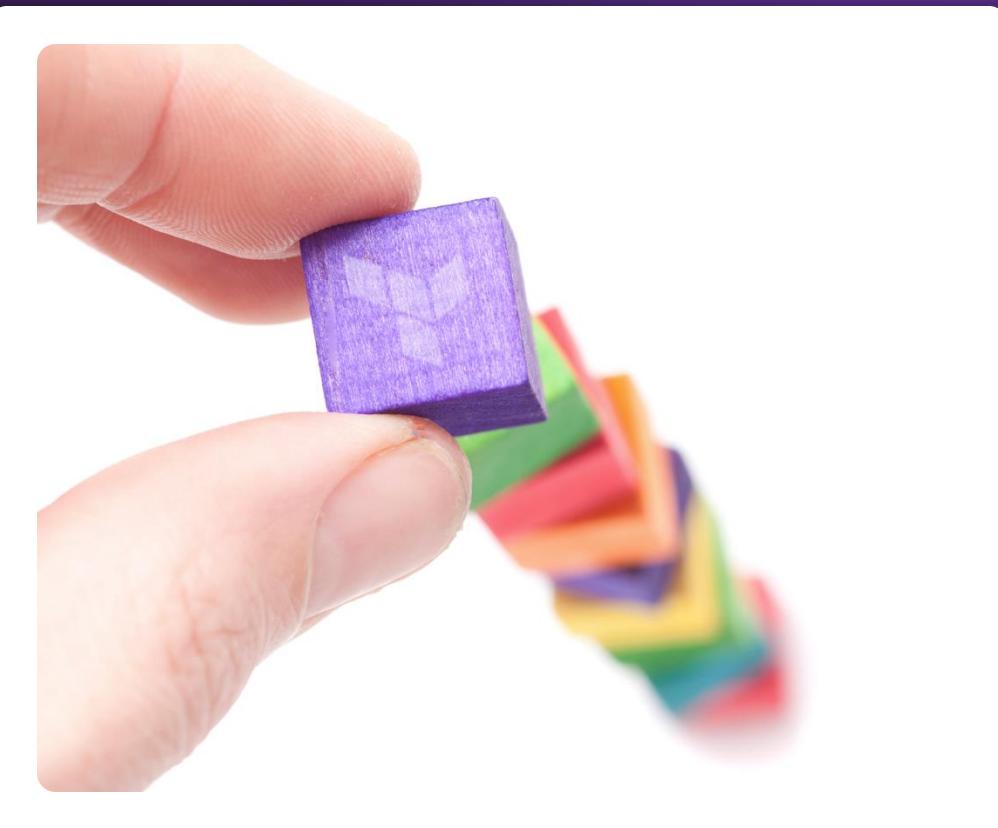
Lecture

Module Block



Terraform Building Blocks





Module Block

The **module** block allows you to reuse and organize Terraform configurations by calling pre-defined modules, which will simplify infrastructure management. Modules allow you to group resources into reusable components, such as a network or database module.

By passing variables to modules and using outputs, you can integrate them into your configuration for more dynamic and flexible deployments.



Module Block

Modules help you create maintainable and standardized infrastructure by packaging commonly used resources together.

Code Reusability

Create standardized infrastructure patterns that can be used across multiple projects and teams.

Simplified Management

Break complex infrastructure into smaller, manageable components that are easier to maintain and update.

Consistent Deployment

Ensure infrastructure consistency by using tested, version-controlled module configurations across environments.

```
main.tf

module "vpc" {
  source = "terraform-aws-modules/vpc/aws"
  version = "4.0.1"

  name = var.vpc_name
  cidr = var.vpc_cidr_block

  azs = ["us-west-2a", "us-west-2b"]
  private_subnets = ["10.0.1.0/24"]
  public_subnets = ["10.0.101.0/24"]

  enable_nat_gateway = true
  single_nat_gateway = true
}
```



Breaking Down the Module Block

```
main.tf

module "vpc" {
  source  = "terraform-aws-modules/vpc/aws"
  version = "4.0.1"

  name        = var.vpc_name
  cidr       = var.vpc_cidr_block

  azs         = ["us-west-2a", "us-west-2b"]
  private_subnets = ["10.0.1.0/24"]
  public_subnets  = ["10.0.101.0/24"]

  enable_nat_gateway = true
  single_nat_gateway = true
}
```

Where the module is located to download

The version of the module to use

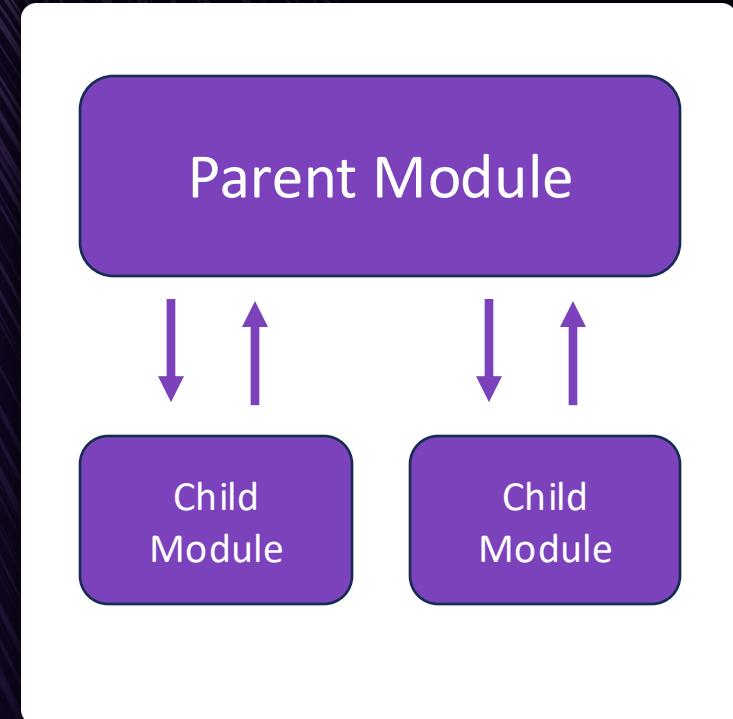
Local variables/values to pass into the module

Module Block



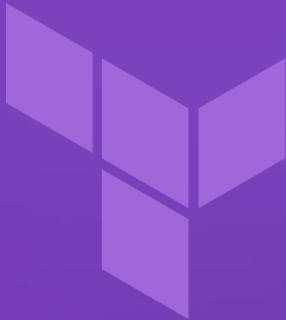
Parent Module vs. Child Module

- A parent (calling) module references and configures other modules, while a child module is the reusable infrastructure component being called.
- Modules can be retrieved from a registry, a repo, or stored locally.



LECTURE

Import Block



Manage all the Things with Terraform...



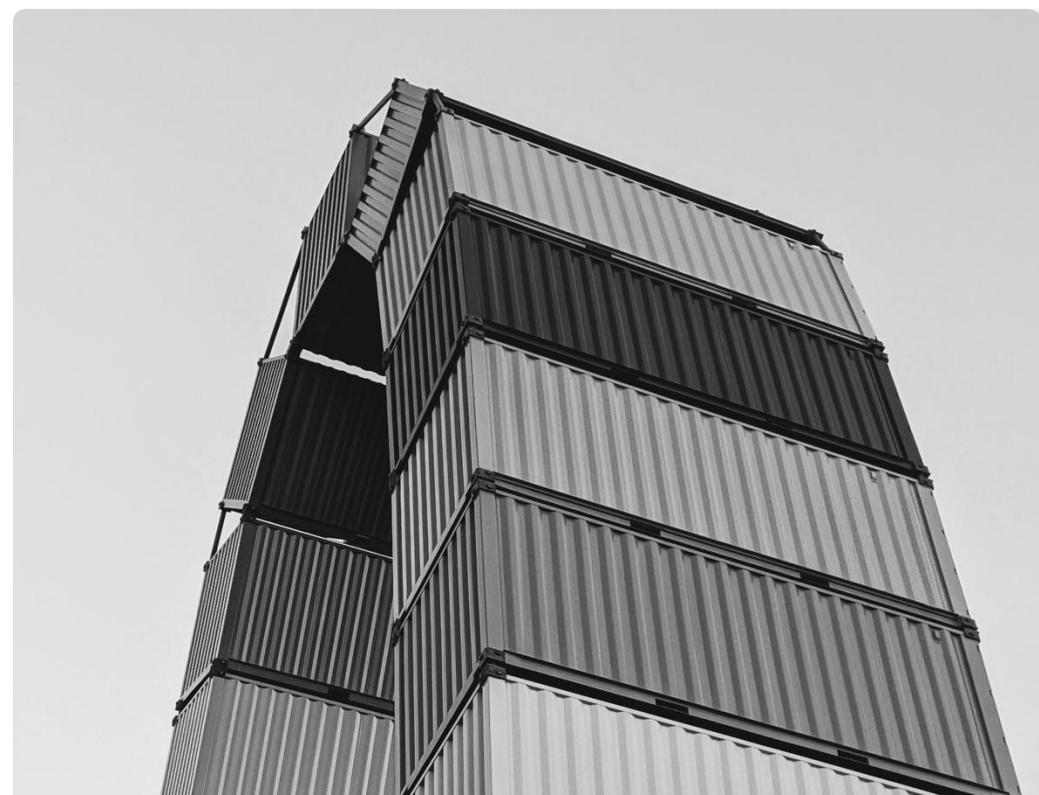
```
resource "aws_instance" "web" {  
    ami           = "ami-012345"  
    instance_type = "t2.micro"  
  
    key_name   = "prd-web-key"  
    subnet_id  = "subnet_12345abc"  
  
    tags = {  
        Name = "prd-web-svr_01"  
    }  
}  
  
resource "aws_instance" "db" {  
    instance_type = "m5.4xlarge"  
    ami           = "ami-0c55b159"  
}
```

After learning Terraform, you want to start provisioning all resources using Infrastructure as Code

Existing Resources



...but what about these resources already running in the cloud/on-premises?



Import Block

The `import` block in Terraform lets you bring existing infrastructure under Terraform management by mapping real-world resources to your configuration. This allows Terraform to track and manage them without manual setup.

It simplifies migration, reduces effort, and ensures infrastructure is versioned and controlled.



Import Block

Define an input that can be reused and customized to make configurations more flexible and dynamic



Map Existing Resources

Links real-world infrastructure to Terraform without changing it



Simplifies Migration

Reduces manual setup when adopting Terraform



Enables State Management

Brings unmanaged resources under Terraform control

```
# Import Repo
import {
  to = github_repository.prd-app-repo
  id = "my-org/production"
}

# Import Azure Resource Group
import {
  to = azurerm_resource_group.prd-rg
  id = "/subscriptions/1234/rg/my-rg"
}

# Import Vault Secret
import {
  to = vault_kv_secret_v2.my-secret
  id = "secret/data/my-secret"
}
```



Breaking Down the Import Block

```
# Import Repo
import {
  to = github repository.prd-app-repo
  id = "my-org/production"
}

# Import Azure Resource Group
import {
  to = azurerm_resource_group.prd-rg
  id = "/subscriptions/1234/rg/my-rg"
}

# Import Vault Secret
import {
  to = vault_kv_secret_v2.my-secret
  id = "secret/data/my-secret"
}
```

• The import block & curly braces

What resource block in Terraform will this be mapped to?

What is the actual resource that you want to import? A VM? A repo? A network?

Another example of importing a resource from Vault

Import Block - Requirements



- Before importing a resource, you need to create the resource block that will manage the imported resource moving forward.
- You can write it manually or use `terraform plan -generate-config-out=PATH`

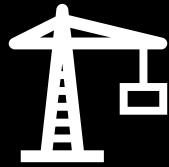
```
# Import Repo
import {
  to = github_repository.prd-app
  id = "my-org/production"
}

# Import Vault Secret
import {
  to = vault_kv_secret_v2.secret
  id = "secret/data/my-secret"
}
```

```
resource "github_repository" "prd-app" {
  name      = "production"
  description = "Production Codebase"
}

resource "vault_kv_secret_v2" "secret" {
  mount      = "secret"
  name      = my-secret
  data_json = jsonencode ({admin=pass})
}
```

SUMMARY



Blocks are the fundamental building blocks of Terraform configurations, each serving a specific purpose in IaC



Output blocks enable visibility into deployment results, crucial for monitoring and debugging



Provider and resource blocks form the core of infrastructure deployment, connecting to platforms and defining what to create



Module and import blocks promote code reuse and easier management of existing infrastructure



Data and variable blocks enhance flexibility by allowing dynamic value injection and existing resource queries



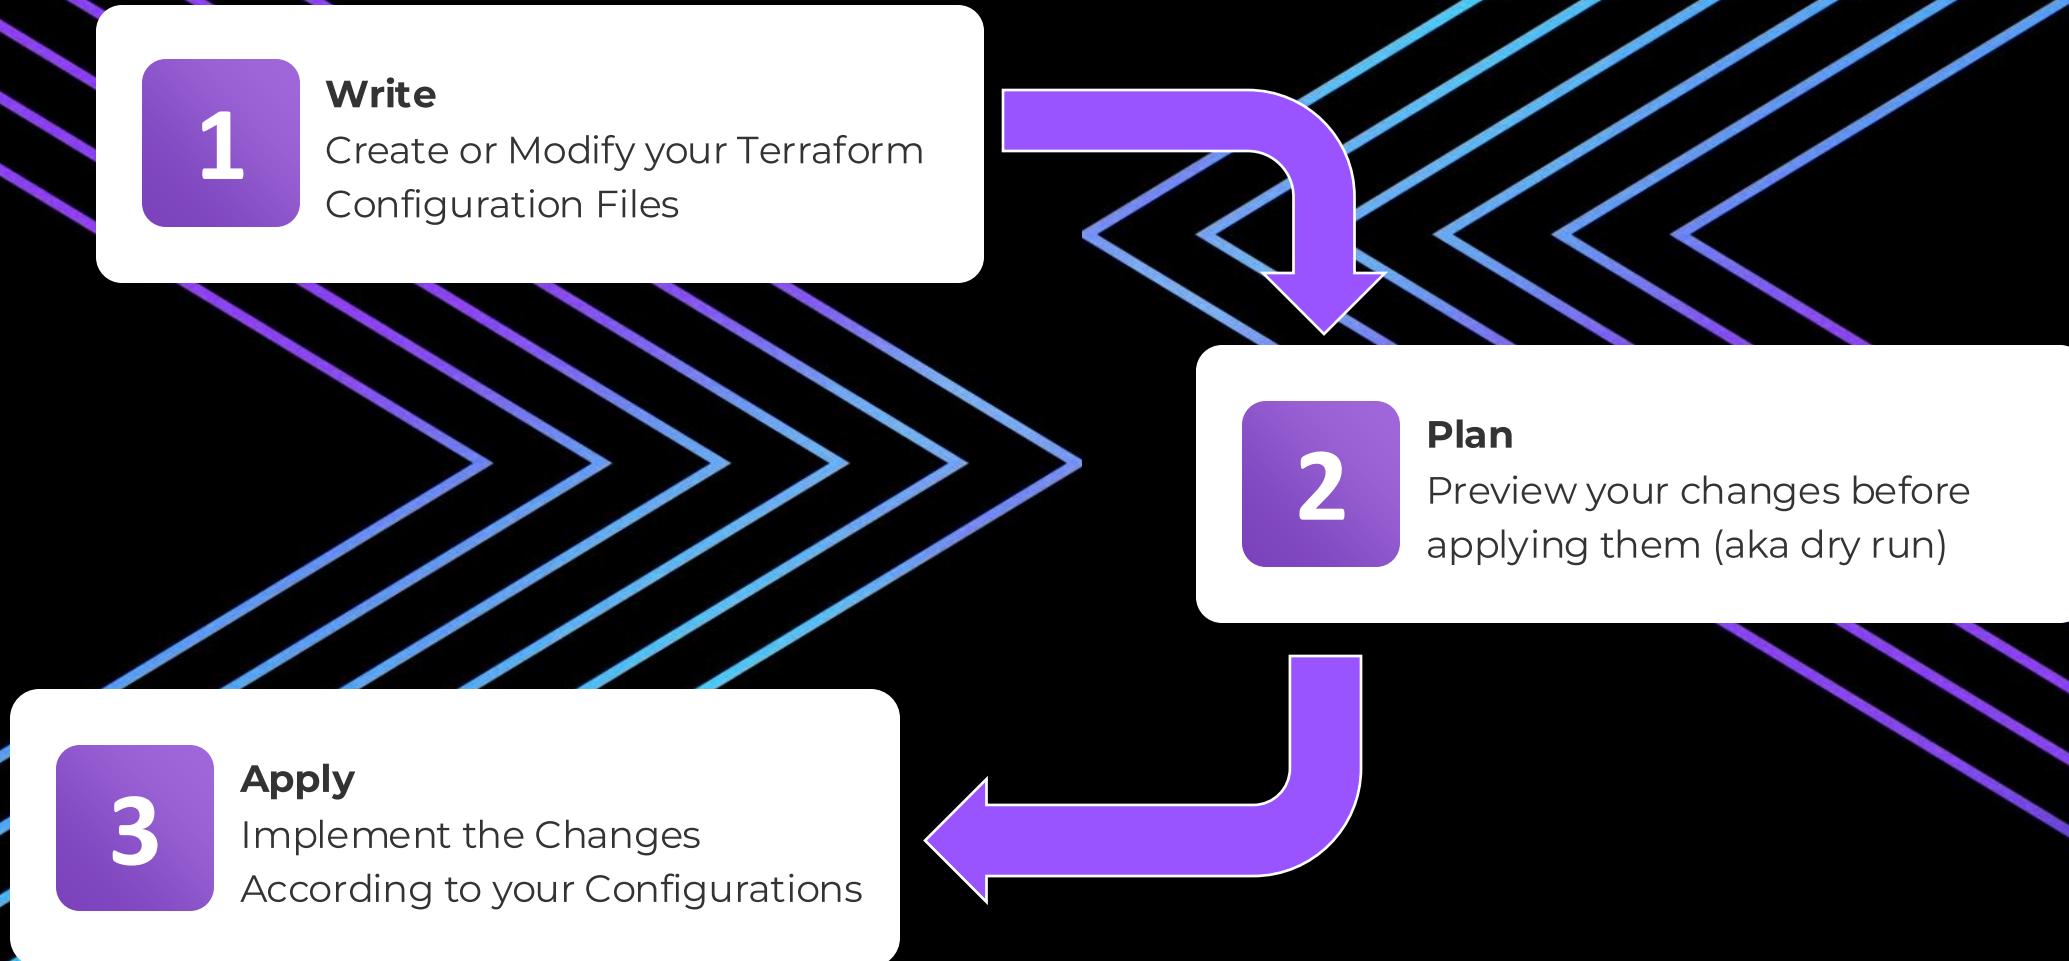
Together, these blocks create a complete framework for infrastructure lifecycle management

Lecture

Introduction to the Terraform Workflow

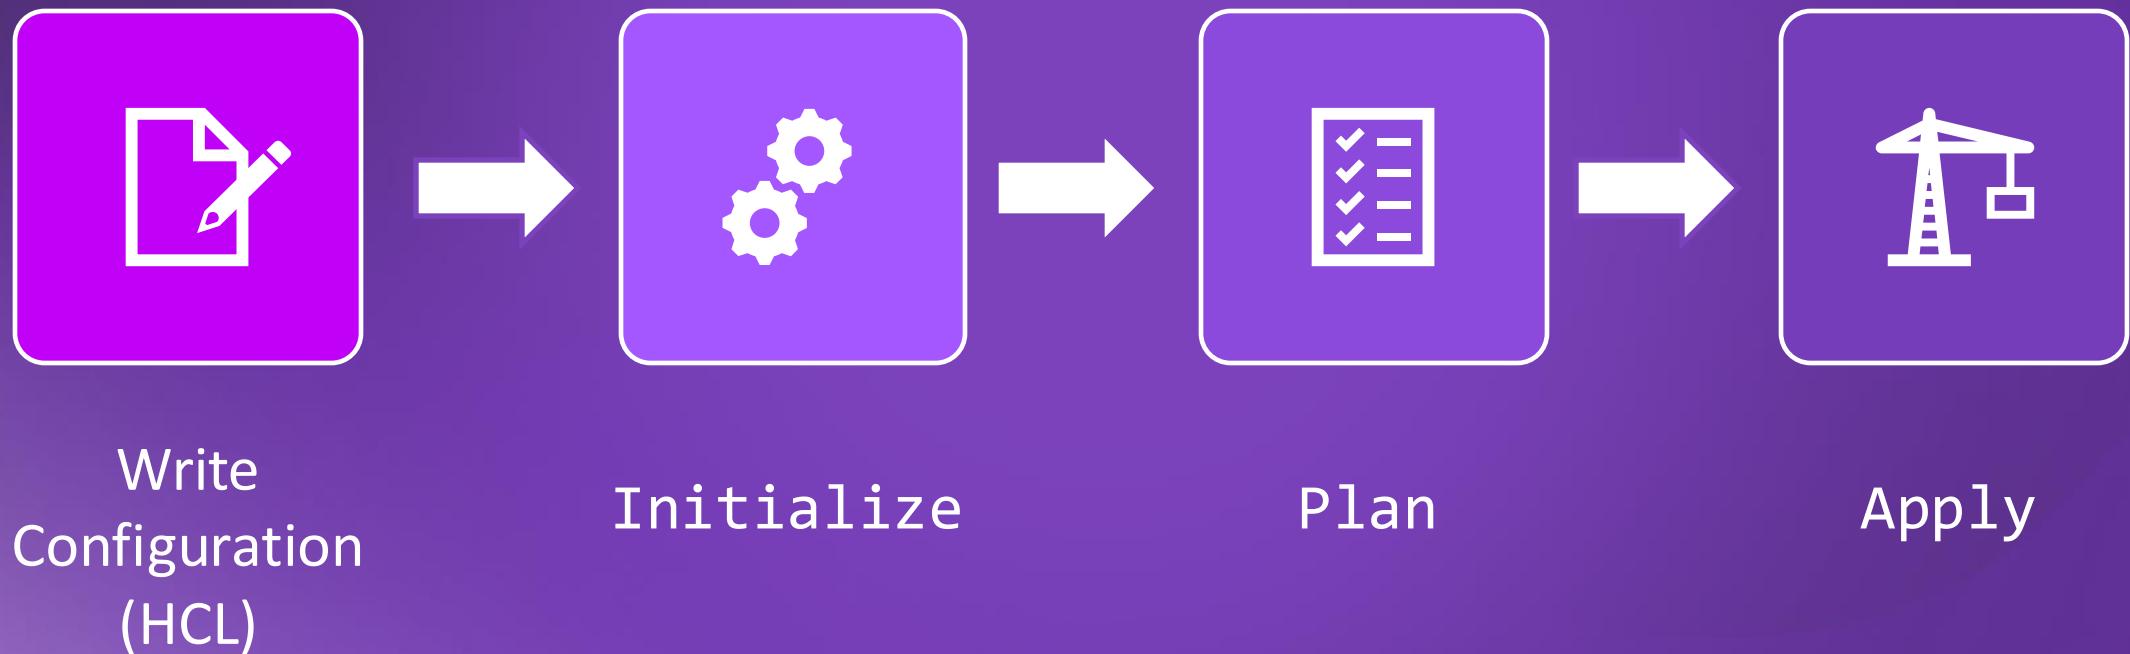


The Terraform Workflow (Official)





Terraform Workflow





Developing Terraform Configuration

IOIO
IOIO

What Does "Develop Configuration" Mean?

The Write stage involves creating one or more configuration files that define your desired state – the end result you want for your infrastructure. Terraform configurations are iterative, meaning you can build and refine your configuration incrementally.



Defining Desired State

You describe the infrastructure components (e.g., VMs, Networks) and the properties for each. Terraform ensures that the "real-world" infrastructure matches this state



terraform init

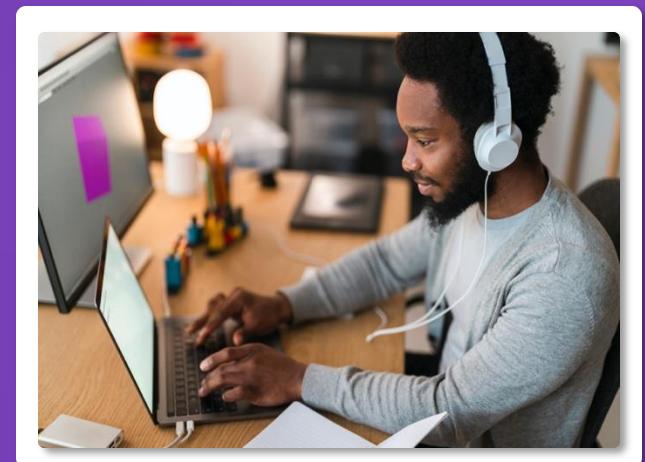
`terraform init` is the first command you run in a new Terraform project to set up your working directory

- **What Does It Do?**

It prepares the environment to run other Terraform commands. Downloads the required providers (e.g., AWS, Azure, GCP) and modules used in the Terraform configuration.

- **Backend Management**

`terraform init` also configures the backend for storing Terraform state



- **When Do You Use It?**

At the start of a new project or any time you add, remove, or upgrade providers/modules or change the backend settings in your configuration.

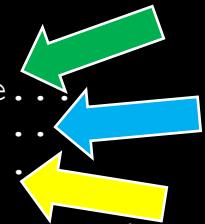


terraform init



TERMINAL

```
$ terraform init
Initializing the backend...
Initializing provider plugins...
- Finding latest version of hashicorp/vsphere...
- Finding latest version of hashicorp/random...
- Finding latest version of hashicorp/vault...
- Installing hashicorp/vault v4.5.0...
- Installed hashicorp/vault v4.5.0 (signed by HashiCorp)
- Installing hashicorp/vsphere v2.10.0...
- Installed hashicorp/vsphere v2.10.0 (signed by HashiCorp)
- Installing hashicorp/random v3.6.3...
- Installed hashicorp/random v3.6.3 (signed by HashiCorp)
Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.
```



Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.



terraform plan

terraform plan generates an execution plan showing the changes that Terraform will make to the 'real-world' resources in order to match the desired state configuration



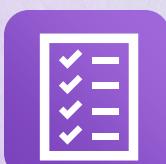
What Does It Do?

Provides you with an output of the changes that will happen if you apply the changes. It shows you the resources to be created (+), updated (~), or destroyed (-).



Perform a Dry-Run of Your Changes

terraform plan gives you the ability to perform a 'dry-run' without impacting 'real-world' resources. It helps avoid unintended changes by showing exactly what Terraform will do



Why a Plan is Important

It is crucial for reviewing changes in production environments to prevent downtime due to errors. It should be run before applying any changes but it is **NOT** required for a **terraform apply**.



terraform plan



TERMINAL

```
$ terraform plan
vault_generic_secret.example_secret: Refreshing state... [id=secret/example]
random_pet.example: Refreshing state... [id=smashing-mutt]
vault_generic_secret.example_kv: Refreshing state... [id=secret/example1]
```

Refreshes State

Terraform used the selected providers to generate the following execution plan.
Resource actions are indicated with the following symbols:

- + create
- ~ update in-place
- destroy
- /+ destroy and then create replacement

Terraform will perform the following actions:

```
# random_pet.additional_pet will be created
+ resource "random_pet" "additional_pet" {
    + id          = (known after apply)
    + length      = 3
    + separator   = "_"
}

# random_pet.example must be replaced
-/+ resource "random_pet" "example" {
    ~ id          = "smashing-mutt" -> (known after apply)
    ~ length      = 2 -> 3 # forces replacement
    # (1 unchanged attribute hidden)
}
```

Create new resource

Replaced



terraform plan

TERMINAL

```
...continued from previous output

# vault_generic_secret.example_kv will be updated in-place
~ resource "vault_generic_secret" "example_kv" {
    ~ data_json          = (sensitive value)
    ~ id                = "secret/example1"
    # (4 unchanged attributes hidden)
}

# vault_generic_secret.example_secret will be destroyed
# (because vault_generic_secret.example_secret is not in configuration)
- resource "vault_generic_secret" "example_secret" {
    - data              = (sensitive value) -> null
    - data_json         = (sensitive value) -> null
    - delete_all_versions = false -> null
    - disable_read      = false -> null
    - id                = "secret/example" -> null
    - path              = "secret/example" -> null
}

Plan: 2 to add, 1 to change, 2 to destroy.
```

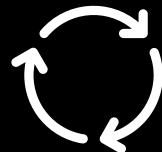
Modify (Yellow arrow pointing to the first resource block)

destroy (Red arrow pointing to the second resource block)

Overall plan (Green arrow pointing to the Plan summary)

Note: You didn't use the -out option to save this plan, so Terraform can't guarantee to take exactly these actions if you run "terraform apply" now.

SUMMARY



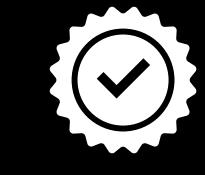
Terraform workflow begins with `terraform init`, which sets up the project and downloads necessary providers and modules.



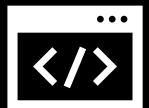
`terraform plan` generates an execution plan, allowing you to preview changes before applying them



`terraform apply` is then used to make the planned changes to your infrastructure.



By default, Terraform will ask for confirmation before making any changes to real-world infrastructure



Each step in the workflow is crucial for managing infrastructure predictably and efficiently.



Following this workflow ensures consistency and control over infrastructure changes.

Lecture

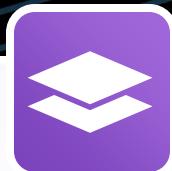
Using the Terraform CLI





What is the Terraform CLI?

The Terraform CLI is the primary way to interact with Terraform and execute the workflow steps



Unified Workflow

The CLI ensures consistency and repeatability across your workflow.
– regardless of what platform that you're using Terraform



Command Interface

A command-line tool to manage infrastructure through commands like `terraform init`, `plan`, and `apply`



Consistent Pattern

Follows a consistent pattern:
`terraform <command> [options]`
`tofu <command> [options]`



Getting to Know the Terraform CLI

Terraform commands follow this structure:

→ `terraform <subcommand> [options or flags]`

```
terraform plan -out=planfile
```



Base command to
invoke the CLI



Subcommand

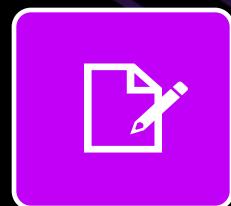


Optional Flag



Mapping the Workflow to Terraform Commands

Write



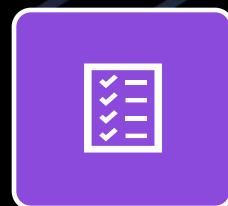
`terraform
fmt`

Initialize



`terraform
init`

Plan



`terraform
plan`

Apply



`terraform
apply`

Destroy



`terraform
destroy`



Terraform Subcommands

- apply
- console
- destroy
- fmt
- force-unlock
- get
- graph
- import

- init
- login
- logout
- modules
- output
- plan
- providers

- show
- state
- test
- validate
- version
- workspace

*not an exhaustive list

Using Environment Variables with the CLI



- Environment variables provide a way to pass configuration settings and credentials to Terraform without hardcoding them into your files.
- They increase security by keeping sensitive data like API keys out of .tf files.

The screenshot shows a terminal window with three tabs: 'TERMINAL', 'PowerShell', and 'Command Prompt'. The 'TERMINAL' tab contains the command '\$ export TF_LOG=DEBUG'. The 'PowerShell' tab contains 'PS> \$Env:TF_LOG = "DEBUG"'. The 'Command Prompt' tab contains three commands: 'C:\> setx TF_LOG = "DEBUG"', 'C:\> setx TF_VAR_svr_name = "prod-db-01"', and 'C:\> setx AWS_ACCESS_KEY_ID = "1a2b3c4d"'. This illustrates how environment variables can be set in different environments (CLI, PowerShell, and Command Prompt) to be used by Terraform.



Making the Most of the Terraform CLI

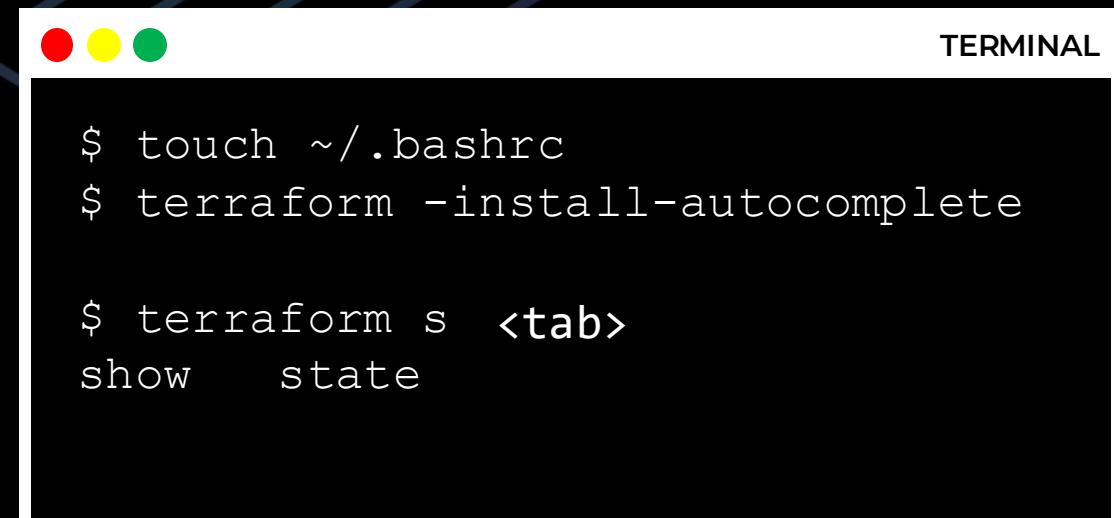


Auto-Complete

Tab completion helps you quickly fill in subcommands, flags, and file paths by pressing the **Tab** key.

It reduces typing errors and ensures you use valid options for Terraform commands.

Available for most shell environments (like Bash or Zsh) after enabling it with `terraform -install-autocomplete`.

A screenshot of a Mac OS X terminal window titled "TERMINAL". The window has red, yellow, and green close buttons at the top. Inside the terminal, the following command sequence is shown:

```
$ touch ~/.bashrc
$ terraform -install-autocomplete
$ terraform s <tab>
show state
```

The word "state" is partially typed, and the terminal is waiting for the user to press the Tab key to complete the command.



Making the Most of the Terraform CLI

Using the help feature

TERMINAL

```
$ terraform --help
Usage: terraform [global options] <subcommand> [args]

The available commands for execution are listed below.
The primary workflow commands are given first, followed by
less common or more advanced commands.

Main commands:
  init          Prepare your working directory for other commands
  validate      Check whether the configuration is valid
  plan          Show changes required by the current configuration
  apply         Create or update infrastructure
  destroy       Destroy previously-created infrastructure
  ...


```



Making the Most of the Terraform CLI

Using the `help` feature



TERMINAL

```
$ terraform plan --help
```

Usage: terraform [global options] plan [options]

Generates a speculative execution plan, showing what actions Terraform would take to apply the current configuration. This command will not actually perform the planned actions.

You can optionally save the plan to a file, which you can then pass to the "apply" command to perform exactly the actions described in the plan.

Plan Customization Options:

The following options customize how Terraform will produce its plan. You can also use these options when you run "terraform apply" without passing it a saved plan, in order to plan and apply in a single command.

`-destroy`

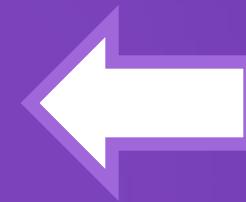
Select the "destroy" planning mode, which creates a plan to destroy all objects currently managed by this



Making the Most of the Terraform CLI

Easily Format Your Code To a Canonical Format and Style

```
1 resource "aws_instance" "example" {
2   ami           = "ami-12345678"
3   instance_type = "t2.micro"
4   tags = {
5     Name = "example"
6   }
7 }
```



`terraform fmt`

Format all Terraform files in
the current directory

Use `-recursive` to format
files in subdirectories as well



Making the Most of the Terraform CLI

Validate Your Configurations

The `terraform validate` command checks your Terraform configuration files for syntax errors and ensures they are structurally valid. It does not check whether the configuration will successfully deploy but ensures your code is properly written and references resources and variables correctly.

This is an essential step to catch errors early before running commands like `plan` or `apply`.

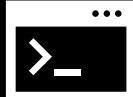
SUMMARY



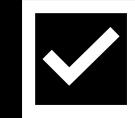
Use `terraform <subcommand>` **to manage each step of the workflow**



Enable `terraform -install-autocomplete` **to speed up typing and reduce errors**



Customize commands with options and flags depending on the subcommand you're using.



Run `terraform validate` **to catch syntax errors early**



Use `terraform fmt` **to ensure clean and consistent formatting**



Pass sensitive data and reusable defaults securely via `TF_VAR_*` **and other variables.**

LECTURE

Understanding Terraform State

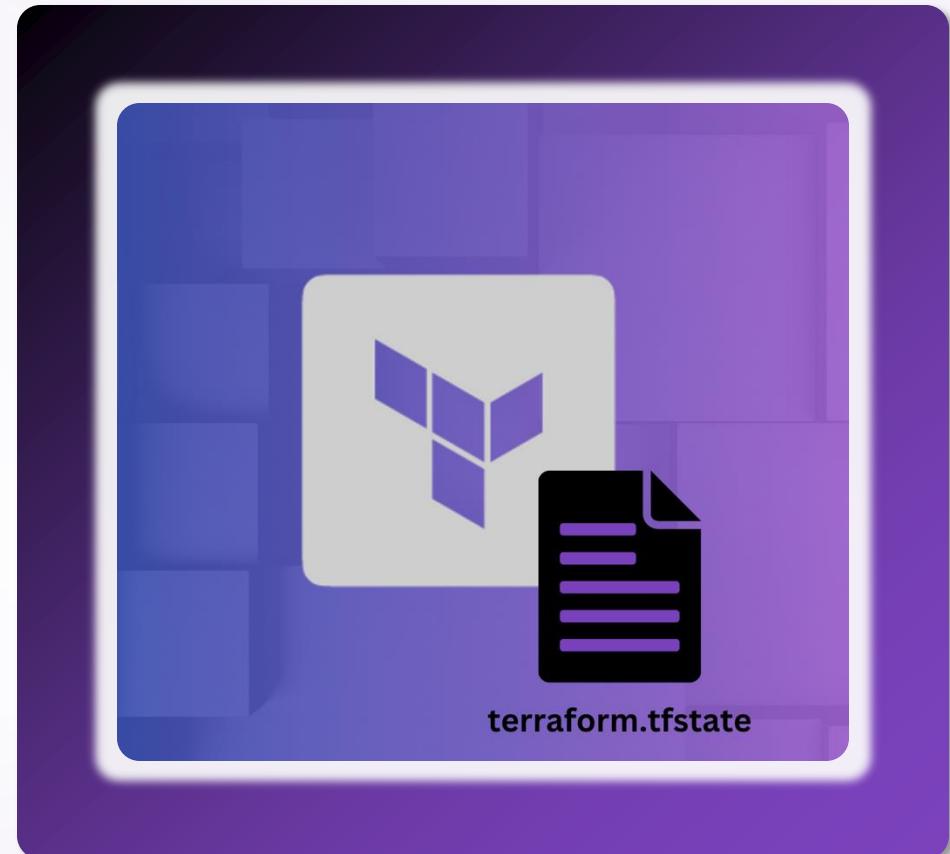




Introduction to Terraform State

Terraform state is a file that tracks the resources managed by Terraform, including their IDs and relationships.

- It acts as a single source of truth for your infrastructure. Without the state file, Terraform would not know its current state.
- Terraform cannot operate without state
- Before any operation, Terraform does a refresh to update the state with the real-world infrastructure



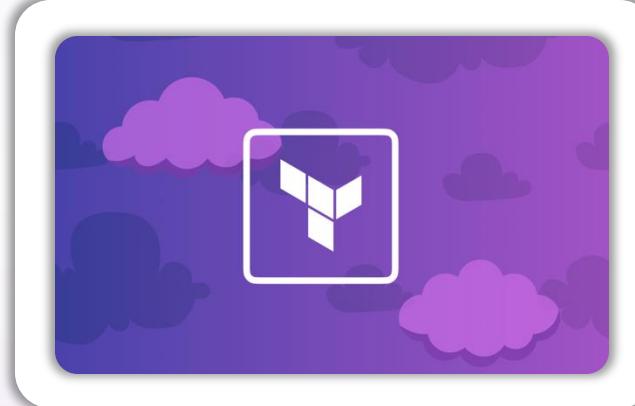


Why Terraform State is Important



Resource Management

Enables Terraform to determine the difference between the desired and current state.



Dependency Management

Tracks resource relationships to apply updates in the correct order.
Terraform does a refresh to update the state before any operations



Code Collaboration

Remote state enables secure, shared workflows for teams.



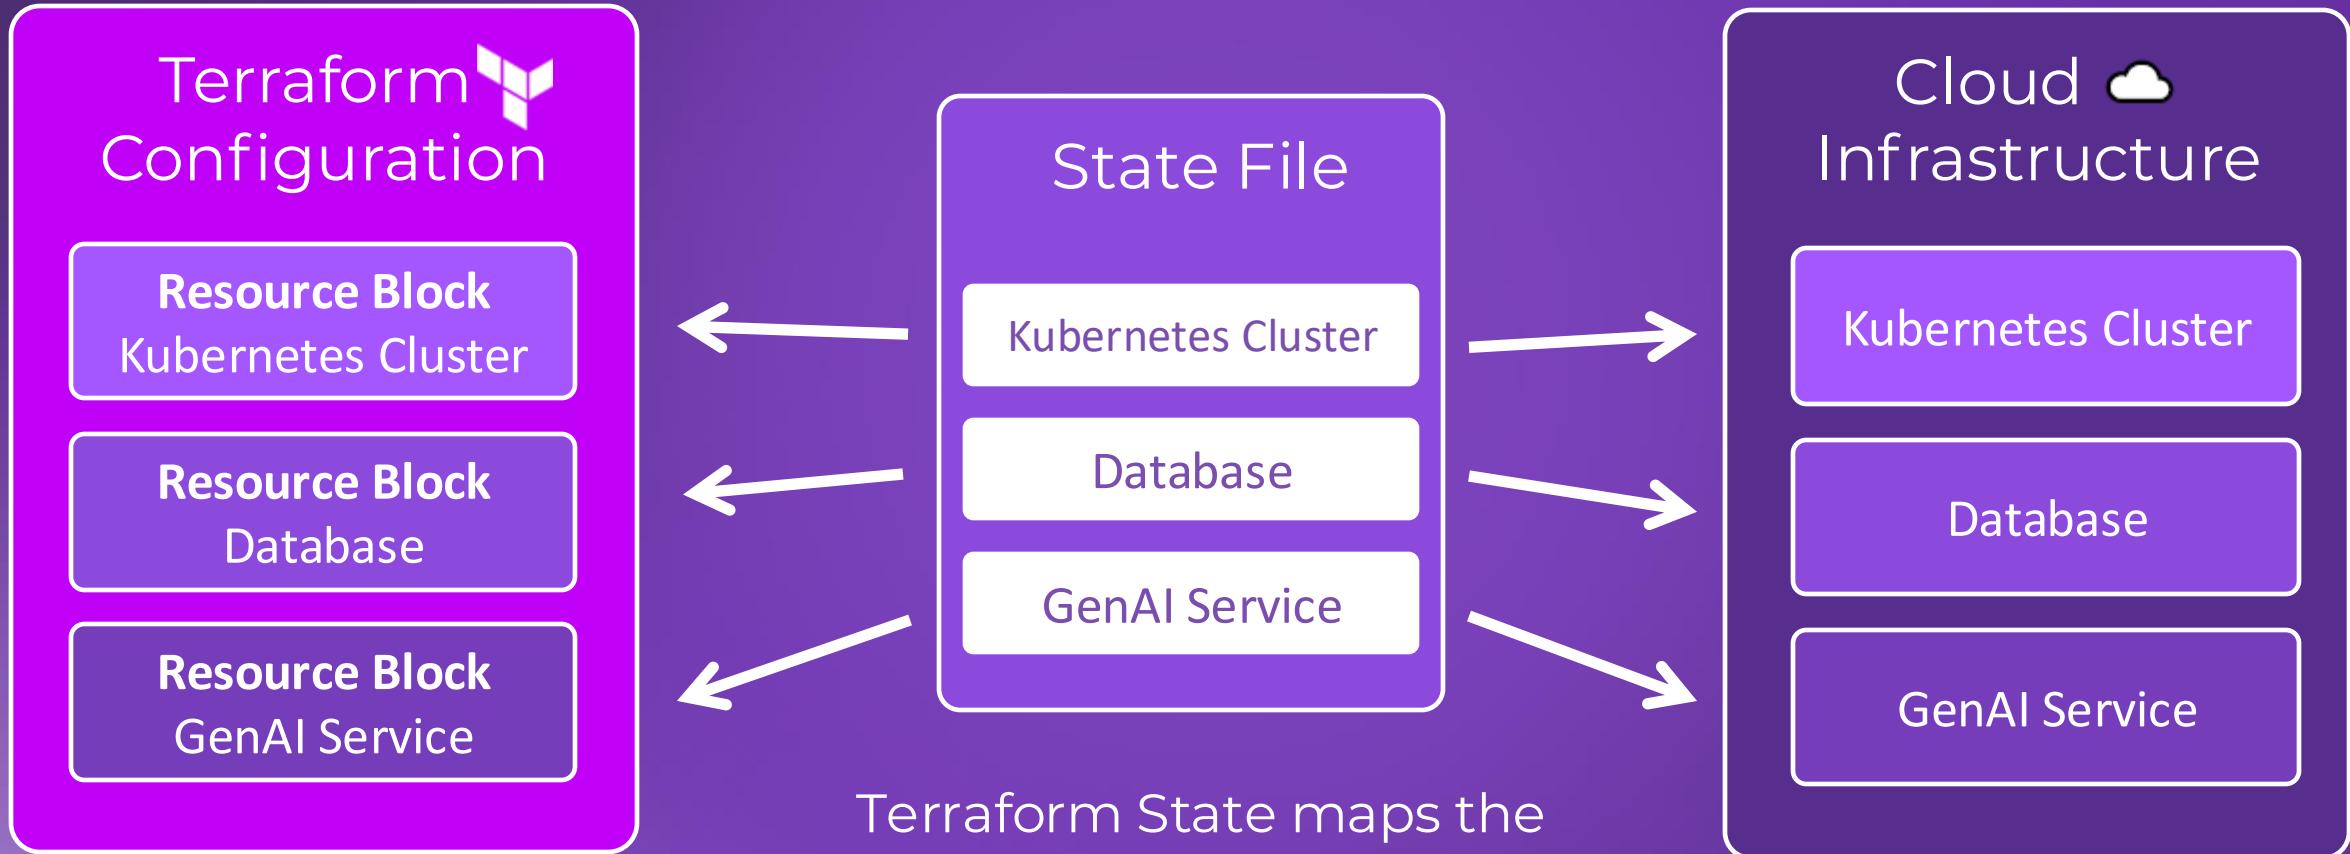
State File

- Y Terraform stores state in file, which acts as a sort of database to map Terraform configurations to real-world resources
- Y Alongside the mapping between resources and remote objects, Terraform also tracks metadata for internal purposes.
- Y This metadata allows Terraform to track order of operations, the provider used, etc.

```
1  {
2    "module": "module.nomad-client[\"nomad-client-3\"]",
3    "mode": "managed",
4    "type": "aws_instance",
5    "name": "this",
6    "provider": "provider[\"registry.terraform.io/hashicorp/aws\"]",
7    "instances": [
8      {
9        "index_key": 0,
10       "schema_version": 1,
11       "attributes": {
12         "ami": "ami-03a6eaae9938c858c",
13         "arn": "arn:aws:ec2:us-east-1:634211456147:instance/i-0eb4591447d4fdf7c",
14         "associate_public_ip_address": true,
15         "availability_zone": "us-east-1c",
16         "capacity_reservation_specification": [
17           {
18             "capacity_reservation_preference": "open",
19             "capacity_reservation_target": []
20           }
21         ],
22         "cpu_core_count": 1,
23         "cpu_options": [
```



Terraform State



Terraform State maps the resources in your config to the real-world infrastructure

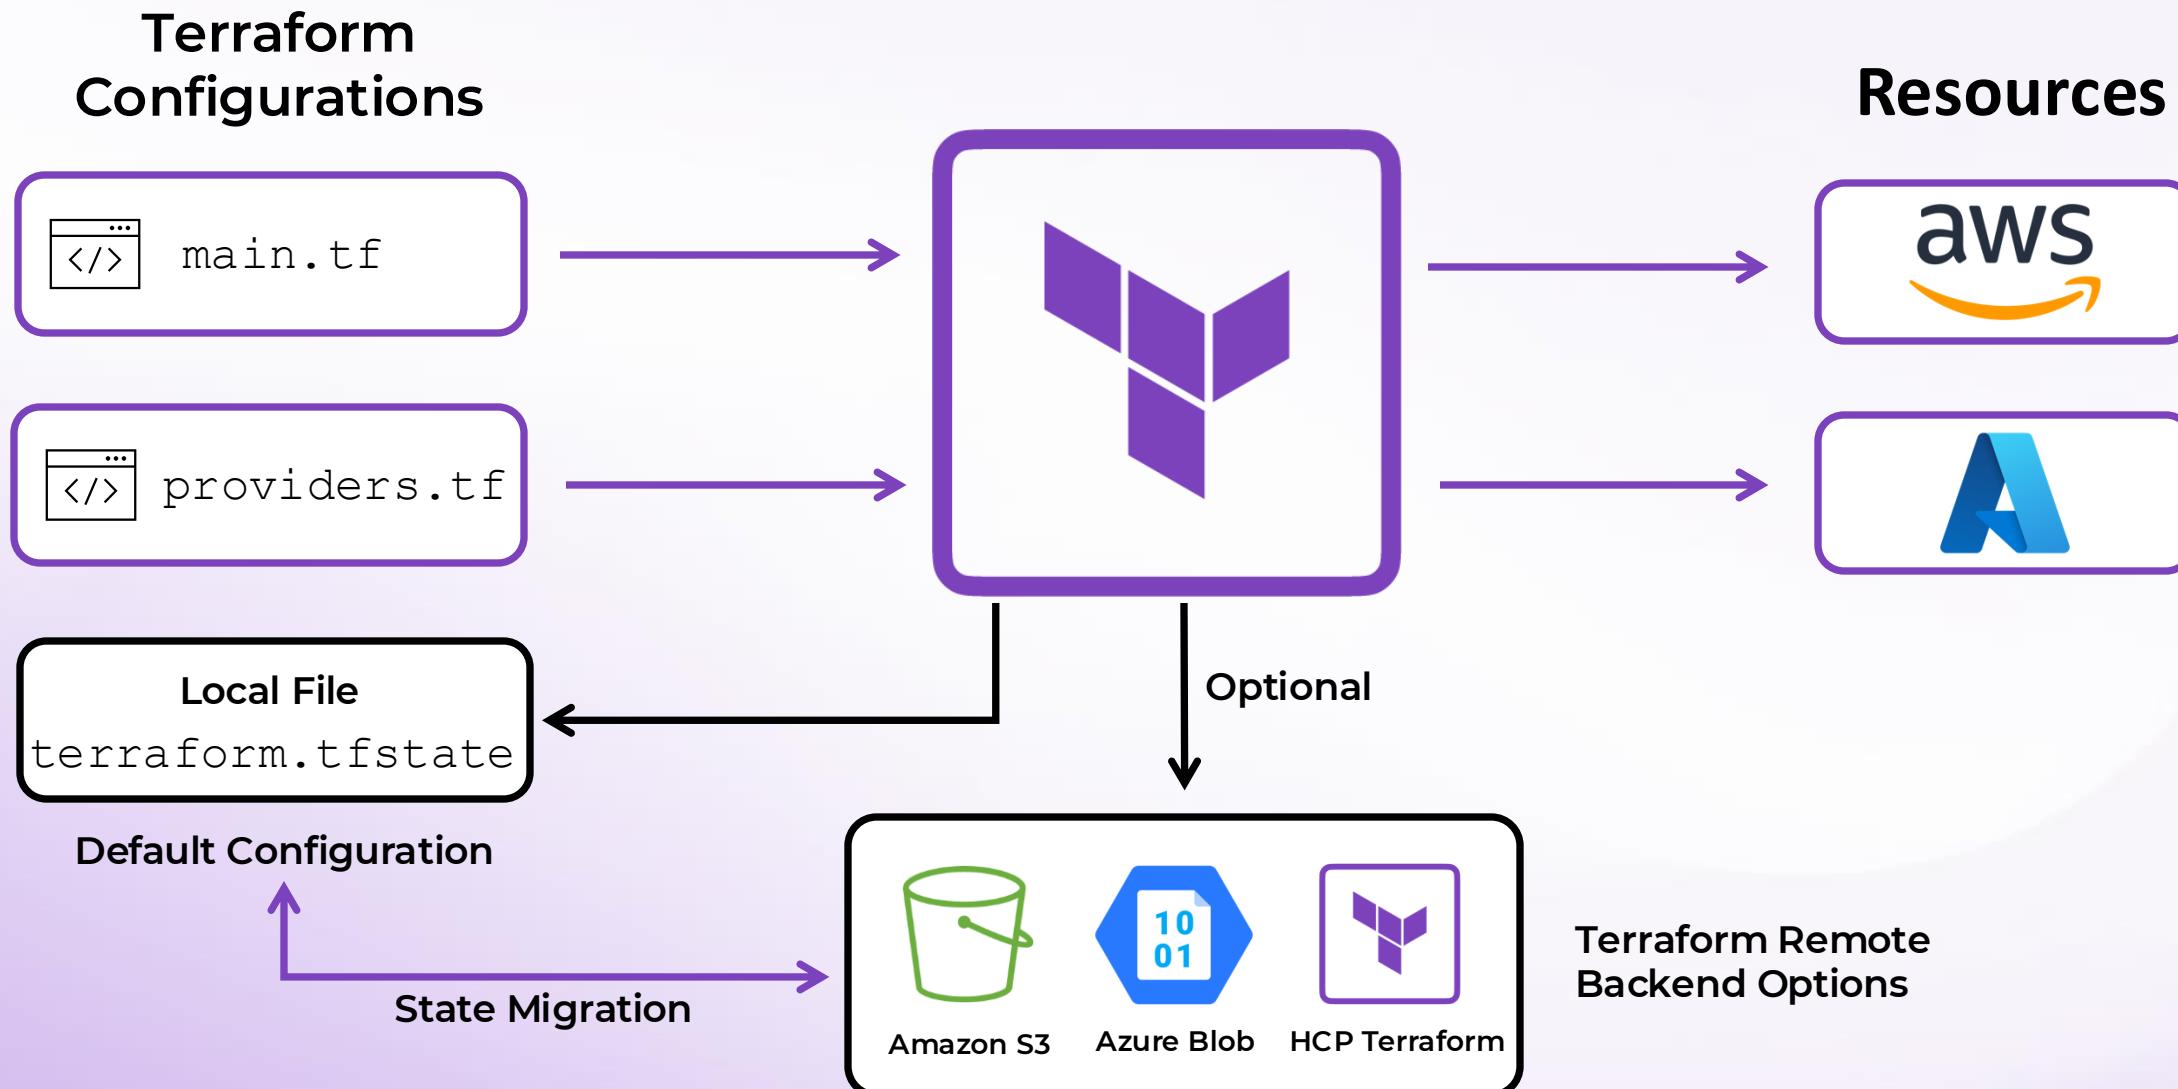


State Storage

- Y By default, Terraform stores state in the current working directory in a JSON-formatted file named `terraform.tfstate`
- Y Terraform needs to "lock" the state file when performing any operations that might write to the file (to avoid conflicts)
- Y Terraform can be configured to store state in a shared location. A backend configuration will determine where state is stored

```
1  {
2    "module": "module.nomad-client[\"nomad-client-3\"]",
3    "mode": "managed",
4    "type": "aws_instance",
5    "name": "this",
6    "provider": "provider[\"registry.terraform.io/hashicorp/aws\"]",
7    "instances": [
8      {
9        "index_key": 0,
10       "schema_version": 1,
11       "attributes": {
12         "ami": "ami-03a6eaae9938c858c",
13         "arn": "arn:aws:ec2:us-east-1:634211456147:instance/i-0eb4591447d4fdf7c",
14         "associate_public_ip_address": true,
15         "availability_zone": "us-east-1c",
16         "capacity_reservation_specification": [
17           {
18             "capacity_reservation_preference": "open",
19             "capacity_reservation_target": []
20           }
21         ],
22         "cpu_core_count": 1,
23         "cpu_options": [
```

Terraform – Remote State



LECTURE

Making Code Reusable

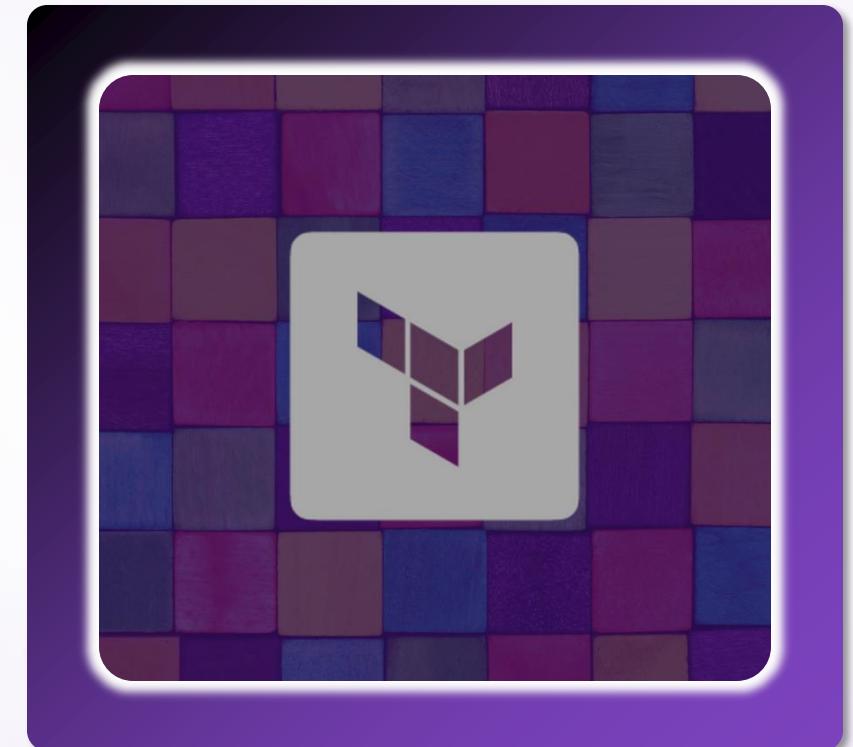




The "Why" Behind Reusability

Building scalable, maintainable infrastructure starts with reusing and standardizing your code. This ensures you can adapt quickly and minimize errors as your projects grow.

- Eliminating repetition (the DRY principle) keeps your code lean and updates simple.
- Centralizing common values reduces confusion and ensures consistency.
- Shared patterns help teams collaborate more effectively and onboard new members quickly.

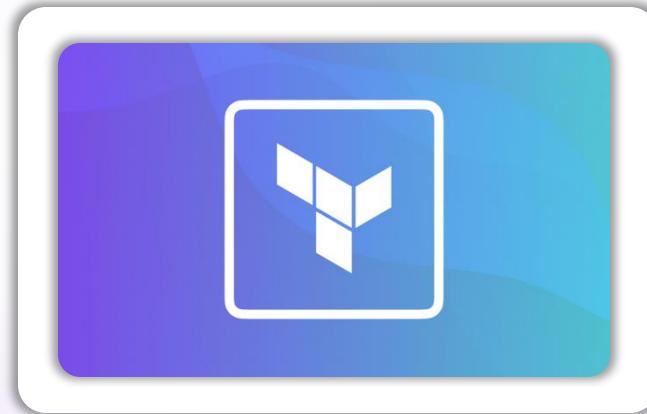




LET'S MAKE YOUR CODE BETTER



The DRY Principle (Don't Repeat Yourself)



Minimize Repeated Code

Copying similar blocks of code can create confusion and tech debt.

DRY keeps your Terraform files clean and easy to maintain



Centralize Common Logic

Using variables, locals, or modules allow you to define values or resource patterns once and reuse them over and over again



Real-World Benefits

Helps standardize names, tags, or resources across the board.

A single source of truth reduces errors and speeds up changes.



The Benefits of Reusable Code

- **Maintainability** – Easier to update and debug
- **Scalability** – Quickly extend or modify infrastructure
- **Consistency** – Uniformity across resources minimizes errors
- **Collaboration** – A clear structure makes teamwork simpler



Reusability = Flexibility



Buying Pizza
(Hardcoding)

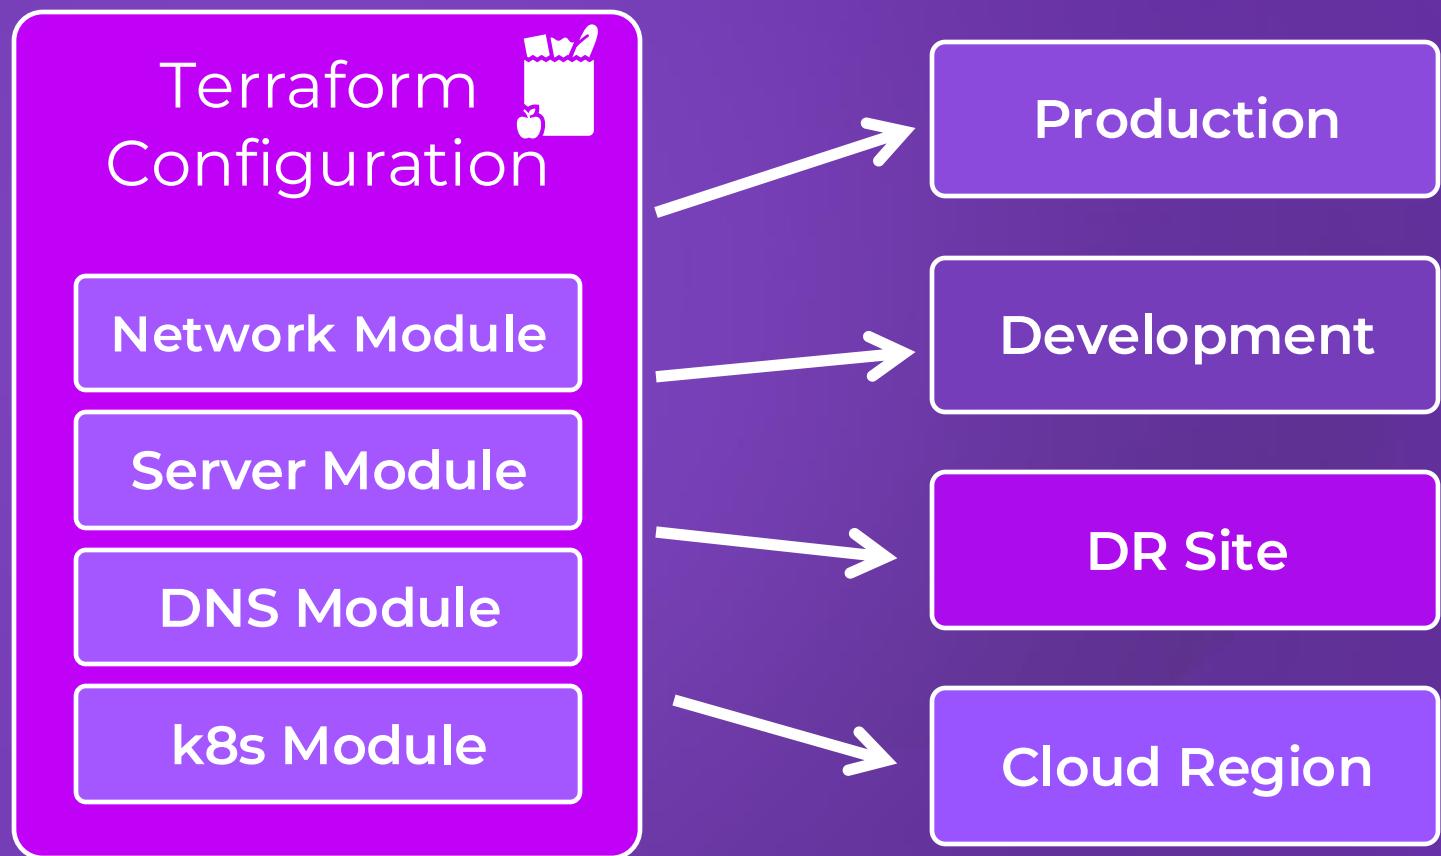
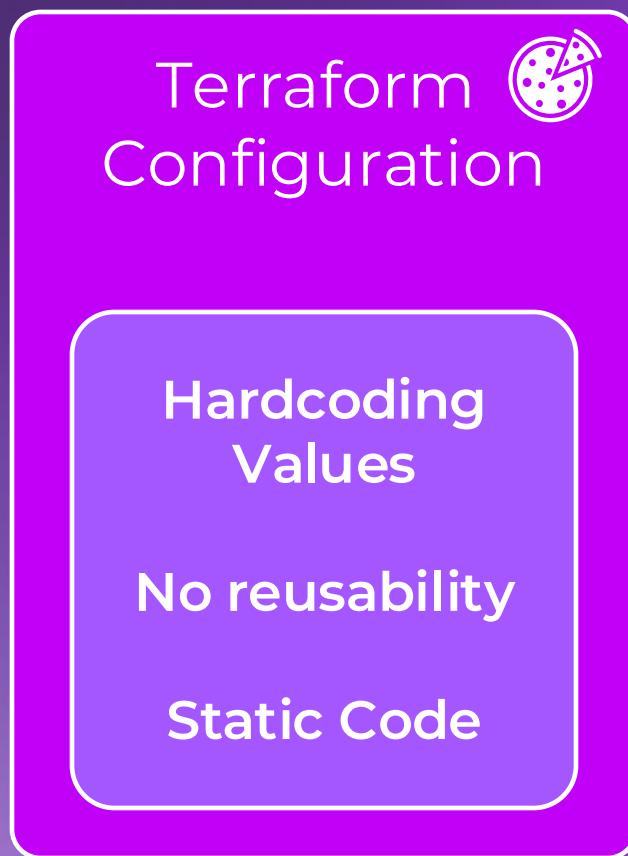


Buying Groceries
(Modular Code)





Repeatability = Flexibility



Lecture

Enhancing Code with Dynamic Values



Variables Recap



Dynamic Inputs for Terraform Code

Input variables let you customize Terraform configurations without changing the source code or hardcoding values



Variable Types

Variable types include string, number, bool, map, set, list.

Set values via defaults, ENV, .tfvars, and command line



Variables Keep Your Code DRY

Pass in different values to create unique resources



Data Sources - Reap

Pull External Information into Terraform



Provider-Specific Info

Data sources (blocks) query provider-specific information about existing resources so we can use it in our code



Common Use Cases

Dynamic lookups, reference existing infrastructure, and avoid hardcoding or repeating information



No Changes

Data sources don't create or make any changes to infrastructure. They only read information and pull it into Terraform



Combining Variables and Data Sources Can Create More Flexible Code

Dynamic code doesn't use static values – it uses interpolation, meta-arguments, functions, local values, and other methods to adapt to changing needs.



Static Inputs & Variables



```
main.tf

resource "aws_instance" "example" {
    ami           = "ami-12345678"
    instance_type = "t2.micro"
}

resource "aws_vpc" "example" {
    cidr_block = "10.0.0.0/16"
    tags = {
        Name = "hardcoded-vpc"
    }
}
```



```
main.tf

resource "github_repository" "example" {
    name          = "my-static-repo"
    description   = "A static name"
    visibility    = "public"
}

resource "github_team" "example" {
    name          = "hardcoded-team"
    description   = "A team with a name"
}
```



Dynamic Code



```
resource "aws_instance" "example" {
    ami = data.aws_ami.ubuntu.id
    instance_type = local.instance_type

    tags = {
        Name = "${var.app}-${local.env}-server"
    }
}

resource "aws_vpc" "example" {
    cidr_block = var.vpc_cidr

    tags = {
        Name = "vpc-${local.env}-${data.aws_region}-${var.network}"
    }
}
```



Interpolation

This \${...} sequence is called an interpolation. It evaluates the expression between the markers, converts the result, and inserts it into the string



main.tf

```
variable "name" {  
    type    = string  
    default = "Bryan"  
}
```

```
"Hello, ${var.name}!"  
Hello, Bryan!
```

```
"My name is ${var.name}, I'm your instructor!"  
My name is Bryan, I'm your instructor!
```



Interpolation

The screenshot shows a code editor window with a dark theme. At the top, there are three colored circles (red, yellow, green) on the left and the file name "main.tf" on the right. The code itself is written in Terraform, a configuration language for infrastructure. It defines variables, data sources, and resources, including interpolation expressions for resource names and bucket names.

```
variable "environment" {
    type    = string
    default = "dev"
}

data "aws_region" "current" {}
data "aws_caller_identity" "current" {}

name = "server-${data.aws_region.current.name}-${var.environment}"
server-us-east-1-dev

bucket_name = "s3-${data.aws_caller_identity.current.account_id}-backups"
s3-9876543210-backups
```

LECTURE

Using Locals to Avoid Code Duplication



When Code Duplication Strikes



```
1 resource "some_resource" "server-abc" {
2   argument      = var.argument
3   argument_type = var.argument_type
4
5   tags = {
6     Name          = "${var.app_name}"
7     ManagedBy    = "Terraform"
8     Team          = var.team
9     Environment  = var.environment
10    ID            = "server-${local.env}-${
11  }
12 }
```



```
1 resource "some_resource" "server-xyz" {
2   argument      = var.argument
3   argument_type = var.argument_type
4
5   tags = {
6     Name          = "${var.app_name}-${var.server}-xyz"
7     ManagedBy    = "Terraform"
8     Team          = var.team
9     Environment  = var.environment
10    ID            = "server-${local.env}-${data.aws_region}-${var.server[1]}"
11  }
12 }
```



What are Local Values?

Named values that you can reference in other parts of your Terraform code

- Most often referred to as "locals," think of locals like short-term memory for your Terraform code
- They centralize and store values you use often, like names, environments, or shared tags
- Local values can be helpful to avoid repeating the same values or expressions multiple times in a configuration

```
1  locals {  
2      app_name      = "my-app"  
3      environment  = "dev"  
4  }  
5
```

Basic Syntax - Local Values



```
● ● ●  
1 # Define shared values in locals  
2 locals {  
3   environment = "dev"  
4   prefix      = "myorg"  
5 }  
6  
7 resource "github_repository" "dev-repo" {  
8   name        = "${local.prefix}-${local.environment}-repo"  
9   visibility  = "private"  
10  description = "Terraform-managed repo for dev environment"  
11 }  
12  
13 resource "github_team" "awesome-people" {  
14   name        = "${local.prefix}-${local.environment}-team"  
15   description = "My Awesome Team"  
16   privacy     = "closed"  
17 }
```

locals block

**Using a
locals block**

Basic Syntax - Local Values



```
1 locals {  
2   app_team = "customer-experience"  
3 }
```

```
4  
5 locals {  
6   # Common tags to be used across all resources  
7   common_tags = {  
8     Name      = "${var.app_name}",  
9     Owner     = var.owner,  
10    App       = var.app_name,  
11    Service   = "${var.service_name}",  
12    AppTeam   = local.app_team,  
13    CreatedBy = data.aws_account_info.account_id,  
14    Image     = data.aws_lambda_image.image_id  
15  }  
16 }
```

```
1 resource "aws_instance" "web_server" {  
2   ami           = data.aws_ami.ubuntu.name  
3   instance_type = "t3.micro"  
4   subnet_id    = var.aws_subnet.public_subnets[0].id  
5  
6   tags          = local.common_tags  
7 }
```

locals block

Benefits of Locals



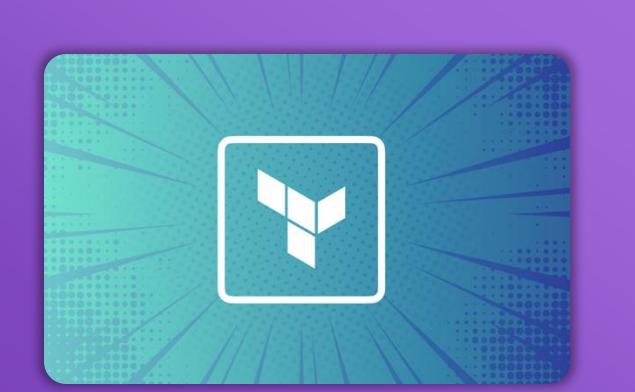
Centralized Logic

The ability to easily change the value in a central place is the key advantage of local values



Clearer Code

Reduces clutter and makes it obvious where a value comes from



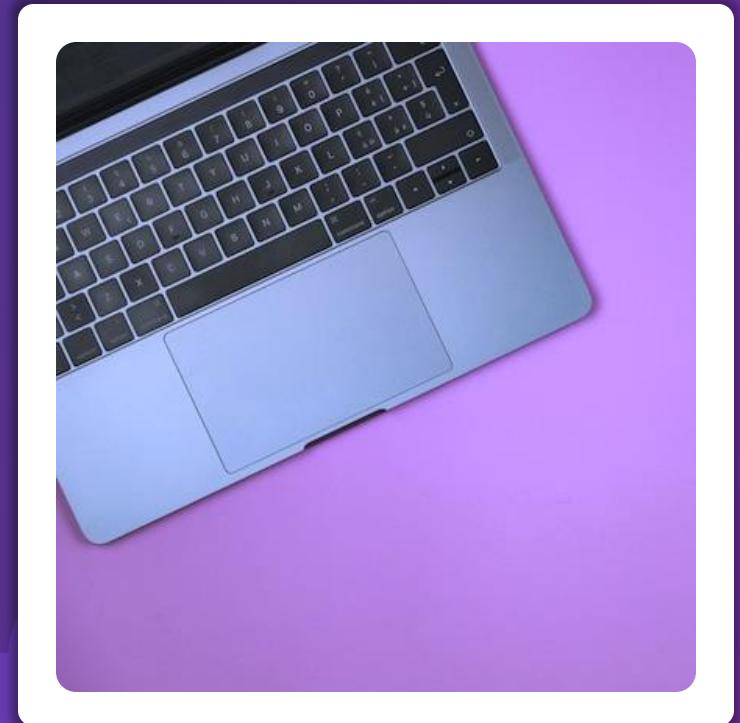
Less Risk

Changing a local is less error-prone than updating multiple hardcoded values



When Should I Use Local Values?

- When you need to combine variables or data sources to create dynamic values, locals help keep your code readable and consistent.
- If the same value appears multiple times (like a prefix, environment name, or common tag), define it once in locals instead of repeating it across multiple resources.
- Instead of embedding lengthy calculations or conditions directly in resource definitions, you can compute them in locals first and reference them where needed.



Lecture

Leveraging Meta-Arguments to Reduce Repetition





What are Meta-Arguments?

Meta-arguments are special arguments that modify how Terraform creates, manages, and structures multiple resources.

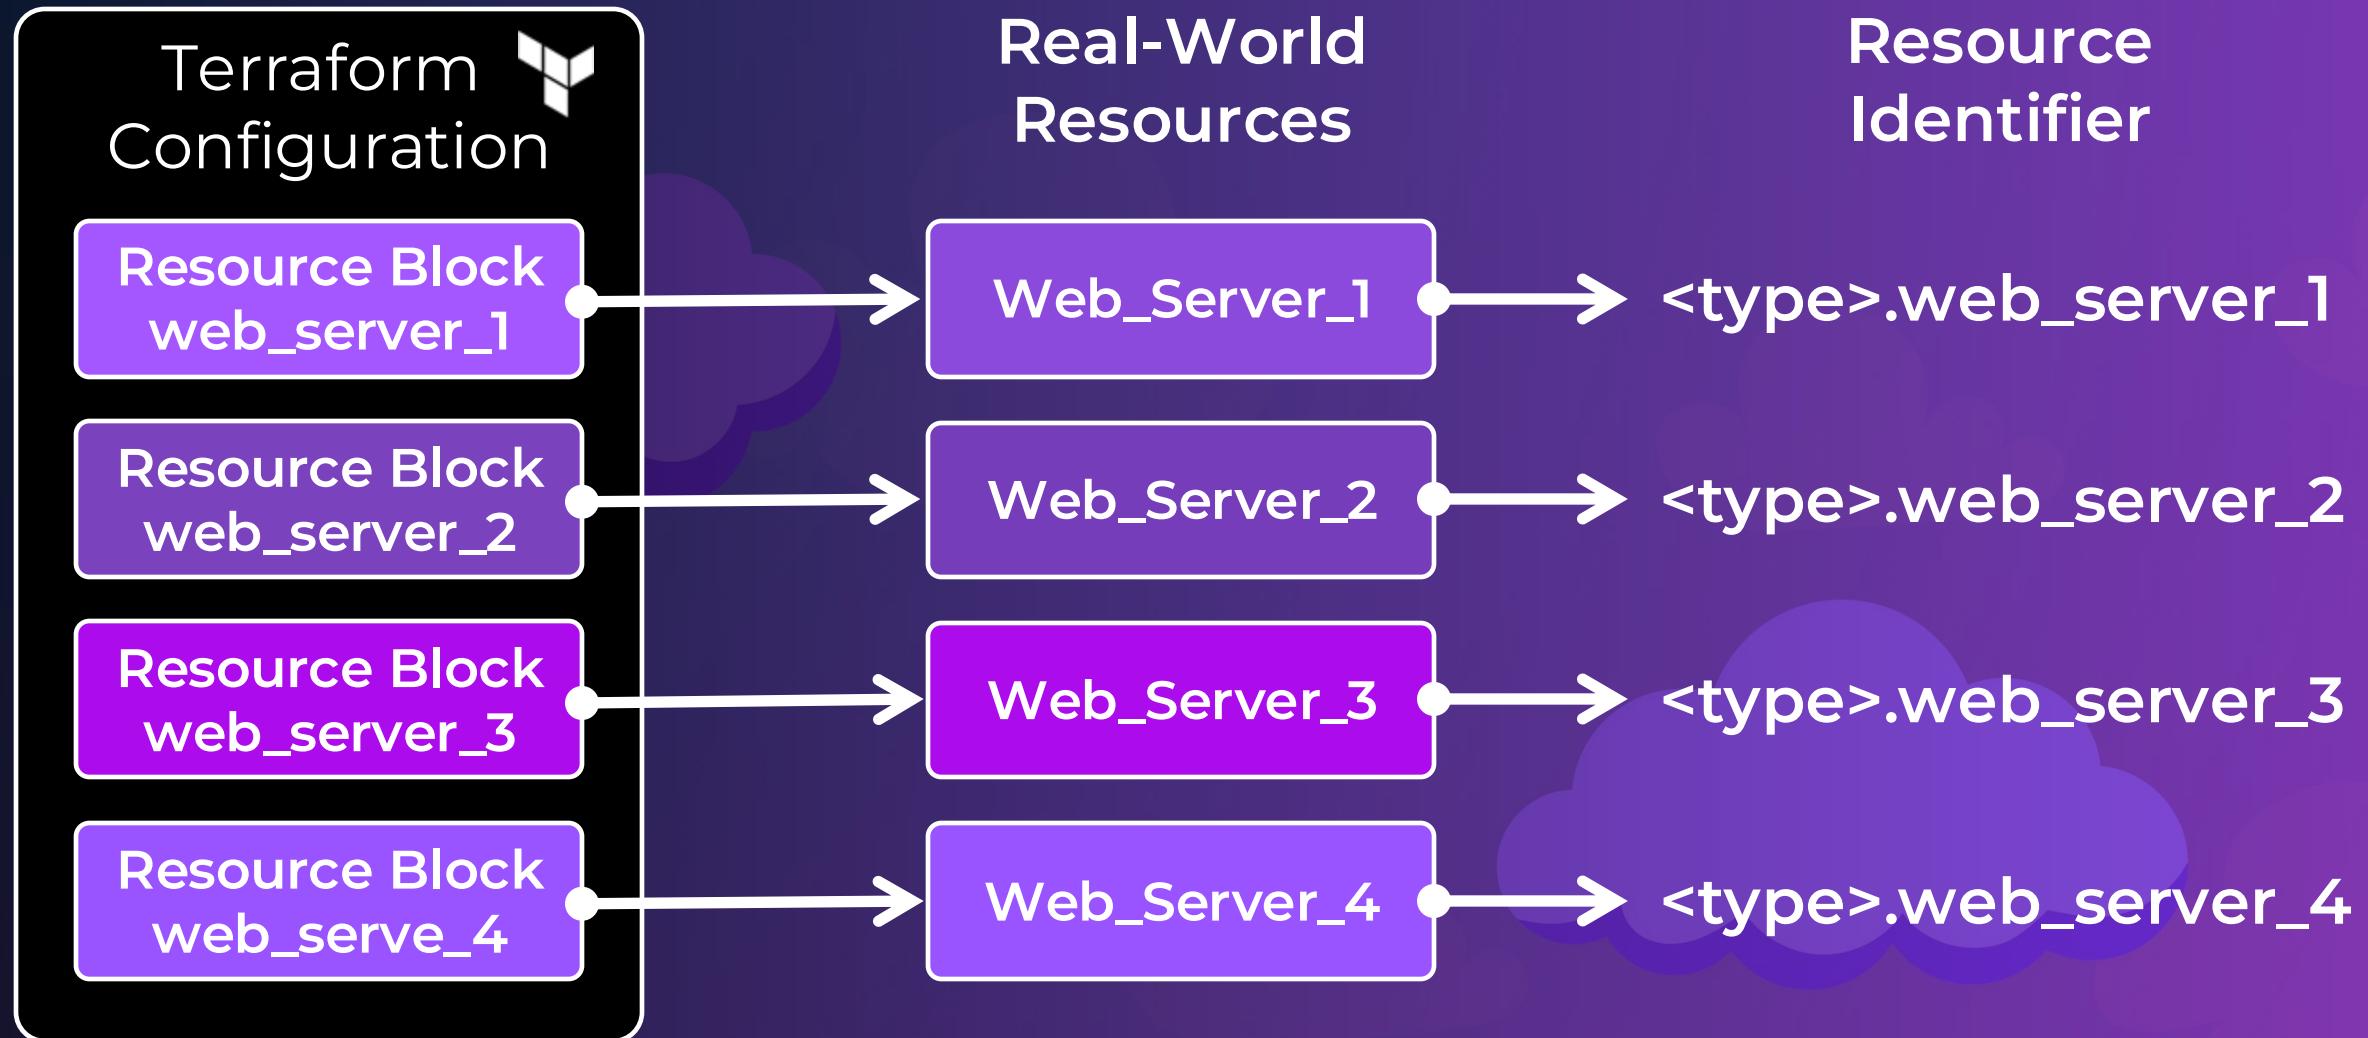
- Reduce redundancy
- Make configurations more scalable
- Simplify resource dependencies

Think of meta-arguments as Terraform's built-in automation tools. Instead of copy-pasting similar resources, you can use meta-arguments to create them dynamically based on input data.



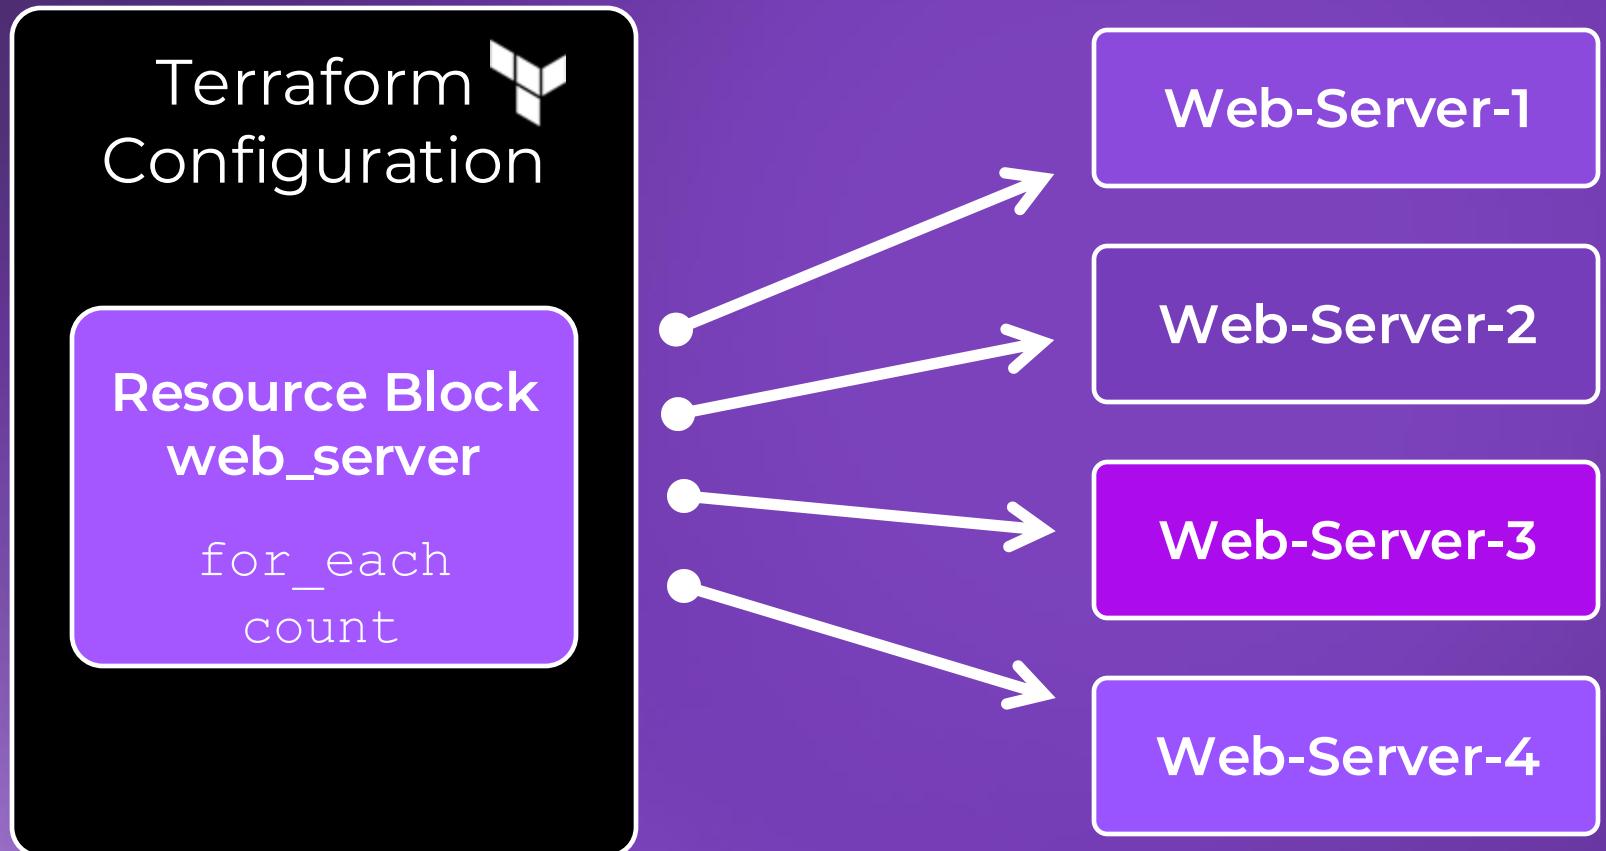


Don't Repeat Yourself





Use Meta-Arguments Instead



`<type>.web_server[0]`

`<type>.web_server[1]`

`<type>.web_server[2]`

`<type>.web_server[3]`

Meta-Arguments

Write Less Code, Deploy More Efficiently

count

3

Allows you to specify a number of instances for a resource.

Example: Deploying five virtual machines with the same configurations



for_each

Lets you create multiple resources dynamically from a map or set.

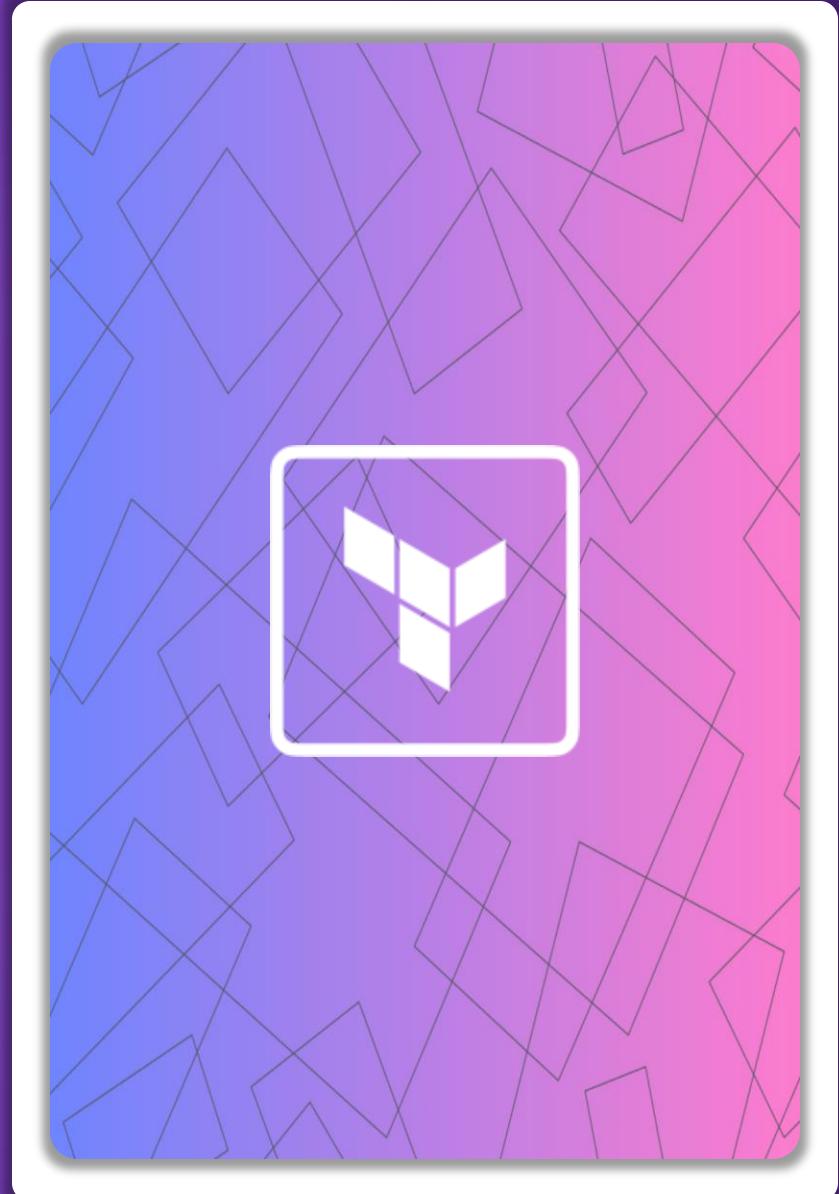
Example: Deploying three VMs with different names instead of the same



depends_on

Explicitly define that a resource must be created before another.

Example: Ensure a database is provisioned before an app server



Meta-Arguments

Write Less Code, Deploy More Efficiently



provider

Specify what provider configuration to use

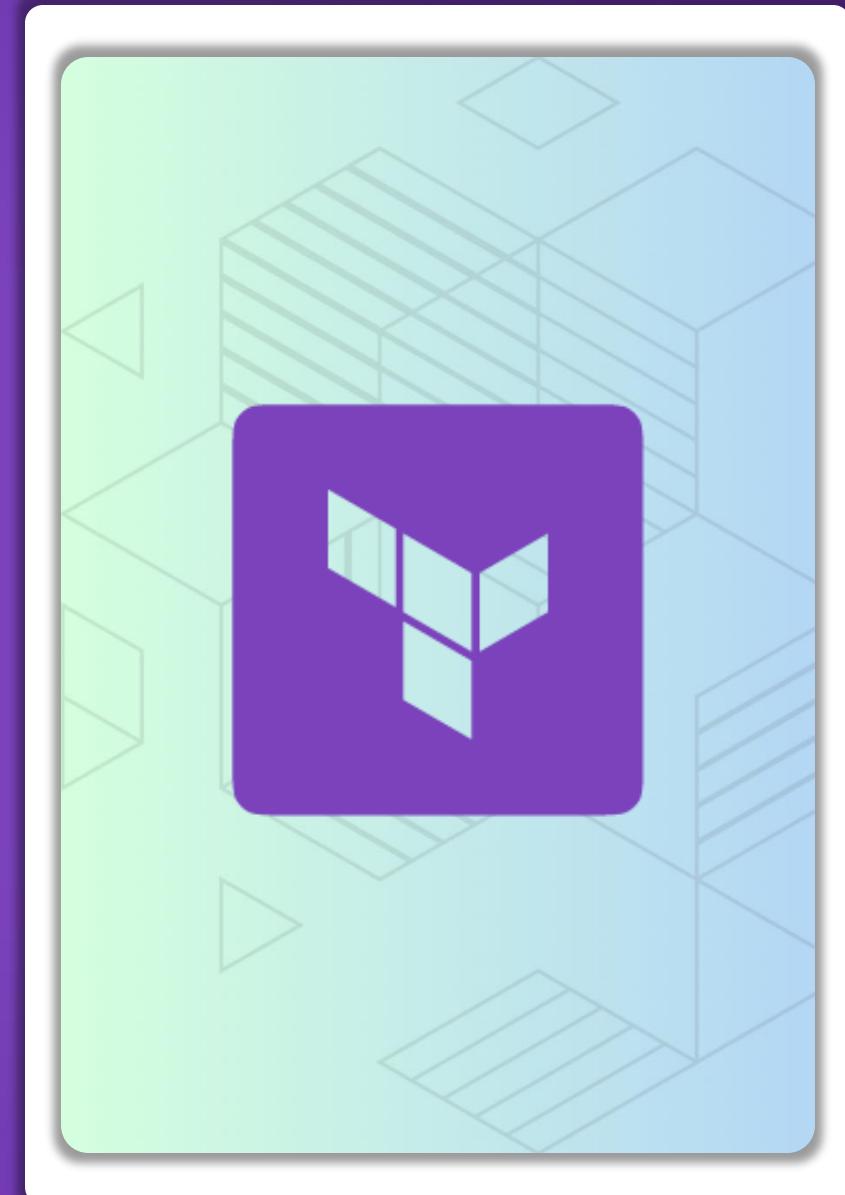
Example: Assign resources to use a different provider block mapped to a second region



lifecycle

Control how Terraform manages the lifecycle of the resource

Example: Protect critical infrastructure or ensure that Terraform doesn't replace resources unnecessarily





count Meta-Argument

Specify a number of instances for a resource

- Allows Terraform to create multiple identical resources using a single block
- Assigns an index (starts at 0) to reference each instance
- Use count to create resources that are identical but need to be repeated or when a variable controls the number of resources
- Reference the resource using <type>.<name>[index]

```
● ● ●  
1 variable "vm_count" {  
2   type    = number  
3   default = 3  
4 }  
5  
6 resource "azurerm_virtual_machine" "example" {  
7   count          = var.vm_count  
8   name           = "vm-${count.index}"  
9   location        = "East US"  
10  resource_group_name = "my-resource-group"  
11  vm_size         = "Standard_D2s_v3"  
12  
13  os_profile {  
14    computer_name  = "vm-${count.index}"  
15    admin_username = "adminuser"  
16  }  
17 }  
azurerm_virtual_machine.example[0]
```

Variable controlling the number of resources



for_each Meta-Argument

A More Flexible Way to Manage Multiple Resources



Creates Multiple Resources

Pulls in values from a map or set rather than a fixed number



Resources Get a Unique Key

Rather than an index, each resource gets a unique key as a reference, making it more stable than count



Ideal for Managing Unique Configurations

Use it when resources need different names, tags, or settings

Variable controlling the resources

```
1 variable "vms" {  
2   type = map(string)  
3   default = {  
4     "vm1" = "Standard_D2s_v3"  
5     "vm2" = "Standard_D4s_v3"  
6     "vm3" = "Standard_B2ms"  
7   }           Key          Value  
8 }  
9  
10 resource "azurerm_virtual_machine" "example" {  
11   for_each            = var.vms  
12   name                = each.key  
13   location             = "East US"  
14   resource_group_name = "my-resource-group"  
15   vm_size              = each.value  
16  
17   os_profile {  
18     computer_name    = each.key  
19     admin_username   = "adminuser"  
20   }  
21 }  
azurerm_virtual_machine.example["vm1"]
```



depends-on Meta-Argument

Ensuring Resources Are Created in the Right Order

- **Explicitly Define Dependencies**

Terraform may not automatically detect or know about dependencies needed for your deployment

- **Ensure Correct Provisioning Order**

Useful for resources that rely on others to exist first

- **Prevents Race Conditions**

Complex deployments may require one resource to be created before another



```
1  resource "aws_db_instance" "db" {
2    identifier      = "mydb"
3    engine          = "mysql"
4    instance_class = "db.t3.micro"
5    allocated_storage = 20
6    username        = "admin"
7    password        = "SuperSecurePassword123!"
8    skip_final_snapshot = true
9  }
10
11 resource "aws_instance" "app_server" {
12   ami            = "ami-0abcdef1234567890"
13   instance_type = "t2.micro"
14
15   depends_on = [aws_db_instance.db]
16 }
```



provider Meta-Argument

Explicitly Assign Resources to a Specific Provider

- ▶ Explicitly assigns a resource to a specific provider configuration.
- ▶ Most often used for multi-cloud or multi-account deployments.
- ▶ Use it when deploying resources across multiple cloud accounts or regions or managing various instances of the same provider (e.g., separate AWS accounts for dev and prod).
- ▶ Uses the `provider = <provider_alias>` syntax

```
● ● ●  
1 provider "aws" {  
2   region = "us-east-1"  
3 }  
4  
5 provider "aws" {  
6   alias  = "prod"  
7   region = "us-west-1"  
8 }  
9  
10 resource "aws_s3_bucket" "dev_bucket" {  
11   bucket  = "my-dev-bucket"  
12 }  
13  
14 resource "aws_s3_bucket" "prod_bucket" {  
15   provider = aws.prod  
16   bucket  = "my-prod-bucket"  
17 }
```



lifecycle Meta-Argument

Controlling How Terraform Manages Resources

- **Customizes how Terraform Manages Resources**

Change the default behavior during your Terraform workflow.

- **You can protect resources or avoid changes**

Customizes how Terraform creates, updates, and destroys specific resources.

- **Uses the lifecycle block inside of your resource**

Options include prevent_destroy, create_before_destroy, and ignore_changes



```
1  resource "azurerm_resource_group" "example" {
2      name      = "my-resource-group"
3      location = "East US"
4
5      lifecycle {
6          prevent_destroy = true
7      }
8
9
10 resource "azurerm_virtual_machine" "example" {
11     name           = "my-vm"
12     location       = "East US"
13     resource_group_name = "my-resource-group"
14     vm_size        = "Standard_D2s_v3"
15
16     lifecycle {
17         ignore_changes = [vm_size]
18     }
19 }
```

Lecture

Using Built-In Functions to Standardize Code



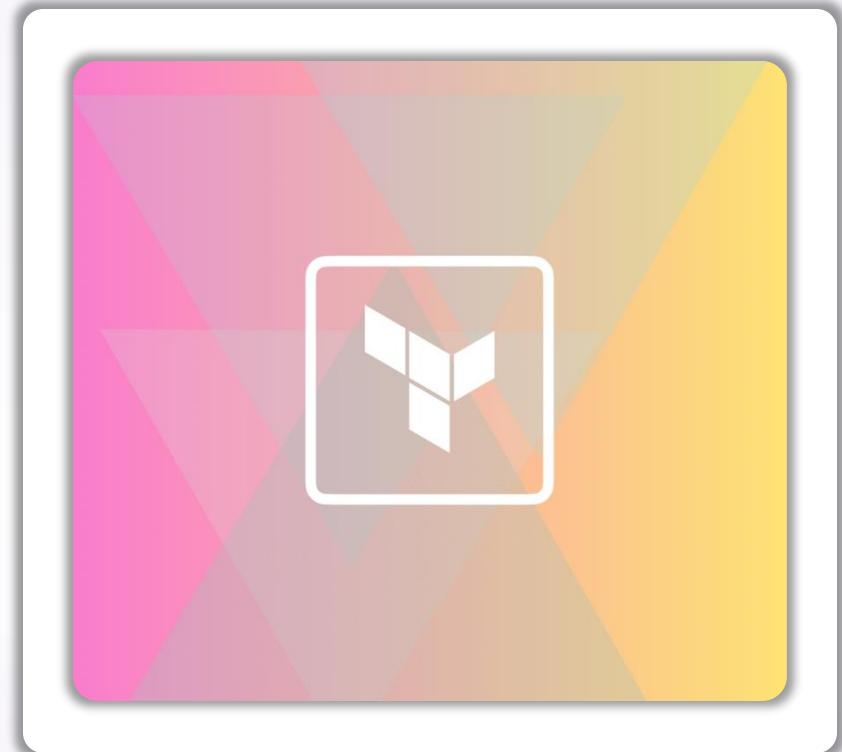


Built-In Functions

Simplify infrastructure code by providing built-in helpers to manipulate strings, numbers, lists, and more.

- Reduces repetitive tasks
- Simplifies logic
- Ensures consistent patterns in deployments

Mastering core functions creates more maintainable, reliable, and readable Terraform configurations.



Core Terraform Functions



Numeric Functions

Quickly evaluate numeric values (e.g., finding the smallest subnet size).

Examples:

`mix` & `max`



String Functions

Concatenate or encode data for resource names, user data, etc.

Examples:

`join` & `split`



Type Conversions

Ensure consistent data structures across your configuration and avoid type mismatch issues.

Examples:

`toset` & `tolist`



General Syntax for Terraform Built-In Functions

Syntax: `function_name(argument1, argument2, ...)`

- ↳ **Example 1:** `upper ("hello")` → **returns** "HELLO"
- ↳ **Example 2:** `min (4, 7, 2, 9, 5)` → **returns** 2
- ↳ **Example 3:** `join ("-", "hello", "terraform")` → "hello-terraform"

Numeric Functions



Determining the largest or smallest value among different variables/outputs

- **Dynamic Resource Allocation:** Adjust CPU, memory, or storage sizes without hardcoding.
- **Automated Calculations:** Remove guesswork by letting Terraform handle numeric logic.



```
1  # Max: Returns the highest value in a list of numbers
2  max(10, 4, 7) → 10
3
4  # Min: Returns the lowest value in a list of numbers
5  min(10, 4, 7) → 4
6
7  # #####
8
9  variable "number" {
10    type      = number
11    default   = 15
12  }
13
14  max(10, 4, var.number) → 15
```

String Functions



Manipulate and transform textual data — enabling consistent naming, simplified automation, and more maintainable configurations.



```
1 # Join: Concatenates a list of strings into a single string.  
2 join("-", ["prod", "web", "us-west-1"]) → prod-web-us-west-1  
3  
4 # upper / lower: Purpose: Converts a string to uppercase/lowercase  
5 upper("example") → EXAMPLE  
6 lower("HELLO") → hello  
7  
8 # replace: Substitutes part of a string with another value.  
9 replace("123-abc", "abc", "xyz") → 123-xyz  
10  
11 # base64encode: Encodes string to Base64, often required for user data  
12 base64encode("my-secret-data") → "bXktc2VjcmV0LWRhdGE="
```

Network Functions



Specialized functions to work with IP addresses and subnets

- **Automate CIDR Calculations:**
Build subnet blocks and host addresses seamlessly.
- **Reduce Manual Errors:**
Eliminate hand-calculating IP ranges.
- **Scalable Network Design:**
Easily create multiple subnets in large or complex deployments.

```
● ● ●  
1 # Create a VPC with a /16 CIDR  
2 resource "aws_vpc" "main" {  
3   cidr_block = "10.0.0.0/16"  
4 }  
5  
6 # Create a public subnet in the first available range  
7 resource "aws_subnet" "public" {  
8   vpc_id        = aws_vpc.main.id  
9   availability_zone = "us-east-1a"  
10  cidr_block     = cidrsubnet(aws_vpc.main.cidr_block, 3, 0)  
11 }  
12  
13 # Create a private subnet in the second available range  
14 resource "aws_subnet" "private" {  
15   vpc_id        = aws_vpc.main.id  
16   availability_zone = "us-east-1b"  
17   cidr_block     = cidrsubnet(aws_vpc.main.cidr_block, 3, 1)  
18 }
```

®Copyright Bryan Krausen – DO NOT DISTRIBUTE

Type Conversion Functions



Convert collections of data into different types for inputs or outputs

Y toset / tolist

Convert collections into lists or sets

Y tomap

Convert a list or object to a map

Y tostring

Convert the argument to a string

```
● ● ●  
1 variable "availability_zones" {  
2   type    = list(string)  
3   default = ["us-east-1a", "us-east-1a", "us-east-1b"]  
4 }  
5  
6 locals {  
7   unique_zones = toset(var.availability_zones)  
8 }  
9  
10 resource "aws_subnet" "example" {  
11   for_each = local.unique_zones  
12  
13   vpc_id      = aws_vpc.main.id  
14   availability_zone = each.value  
15   cidr_block   = cidrsubnet(aws_vpc.main.cidr_block, 2, each.key)  
16 }
```

LECTURE

Intro to Modules





What are Modules?

A module is a container for related resources that are used together.
They are the primary way to package and reuse configurations.



Improved Organization

Break large configurations into logical units to easily reuse over and over again.



Easier Collaboration

Teams can share and maintain standardized building blocks or publish approved modules for organizational use.

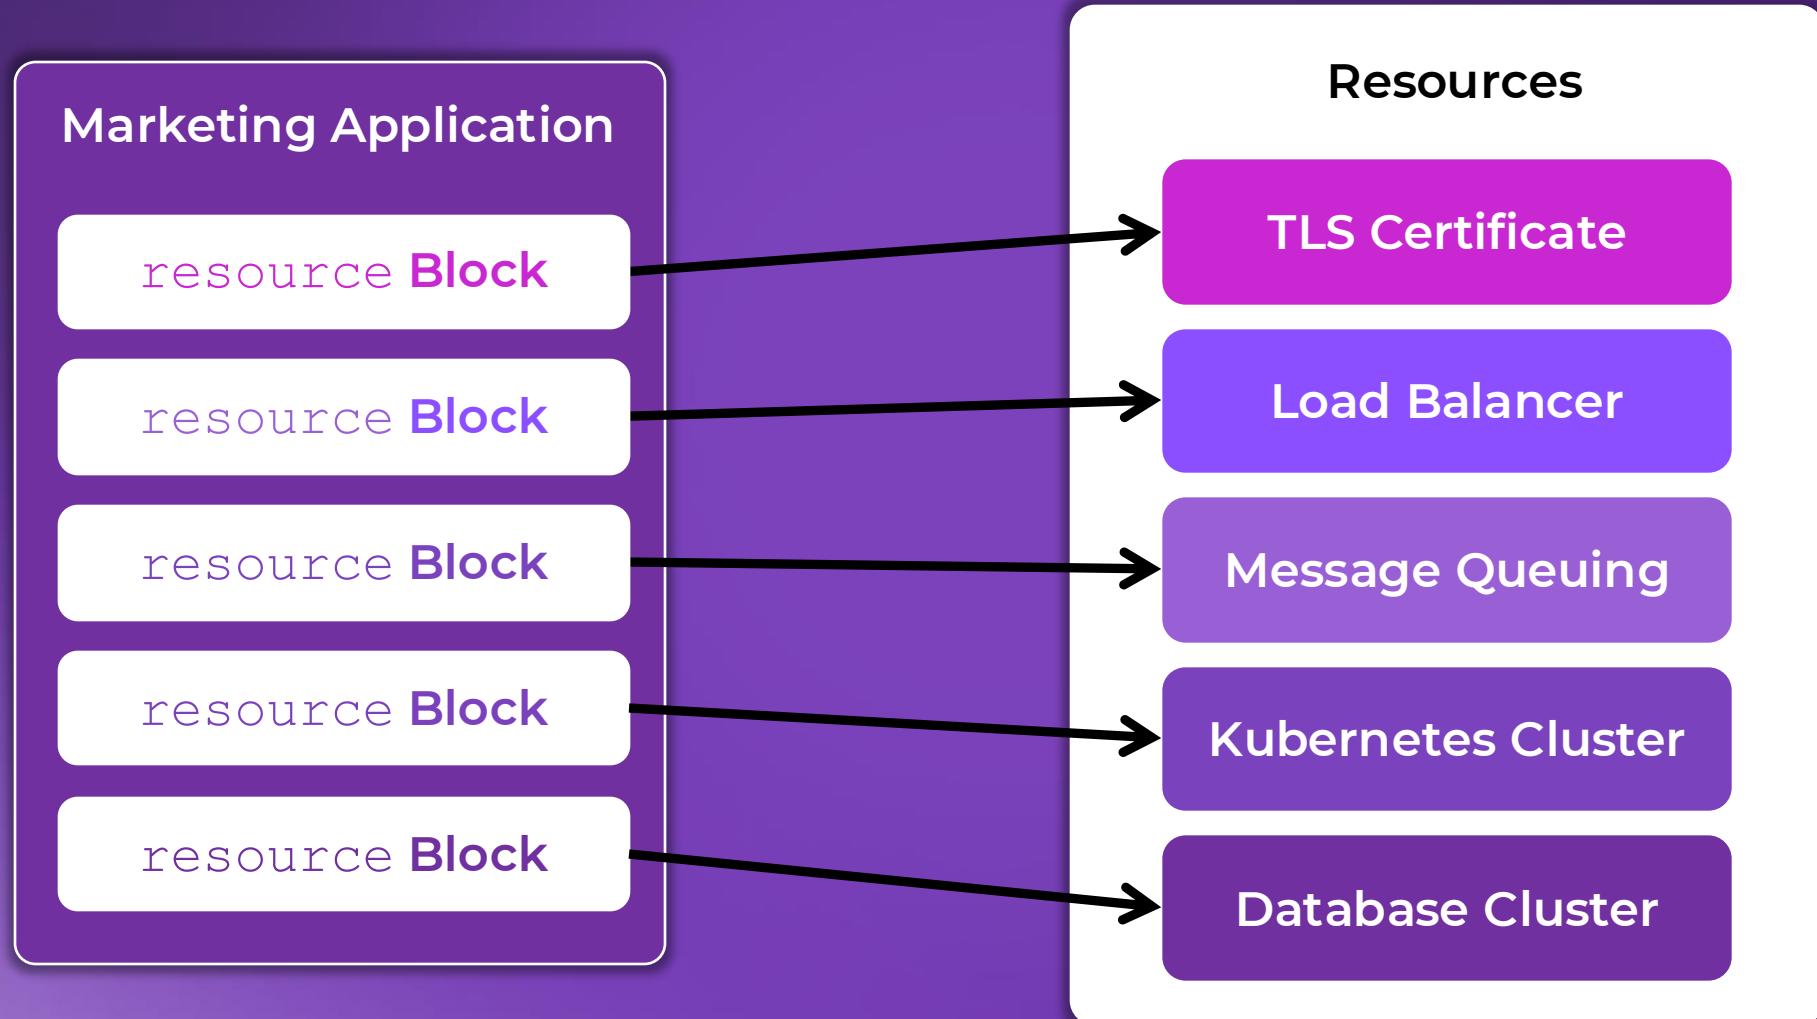


Consistent Patterns

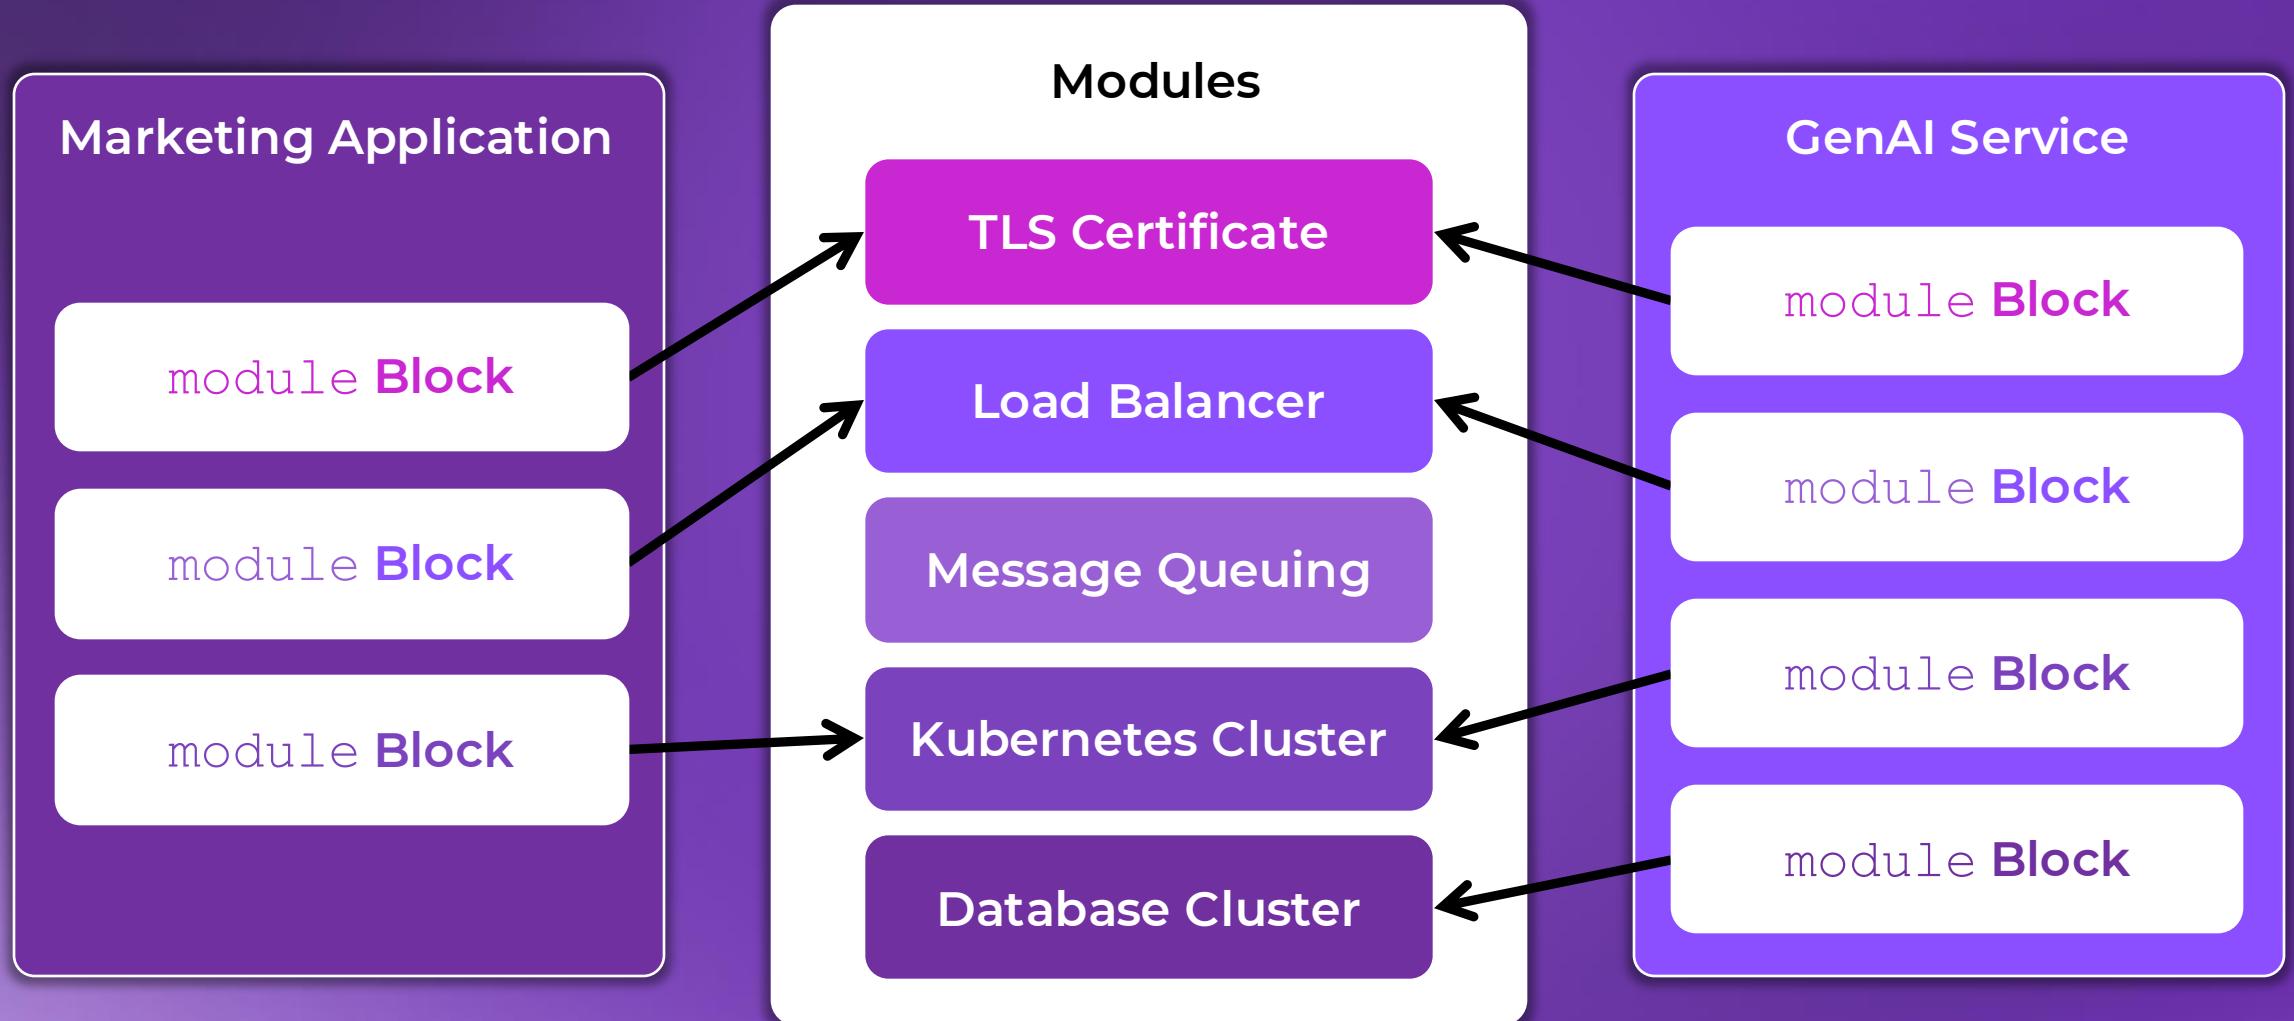
Enforce naming conventions, security controls, and best practices across deployments.



Terraform Building Blocks



Terraform Building Blocks





Terraform Building Blocks

Modules

TLS Certificate

Load Balancer

Message Queuing

Kubernetes Cluster

Database Cluster

Module

main.tf

variables.tf

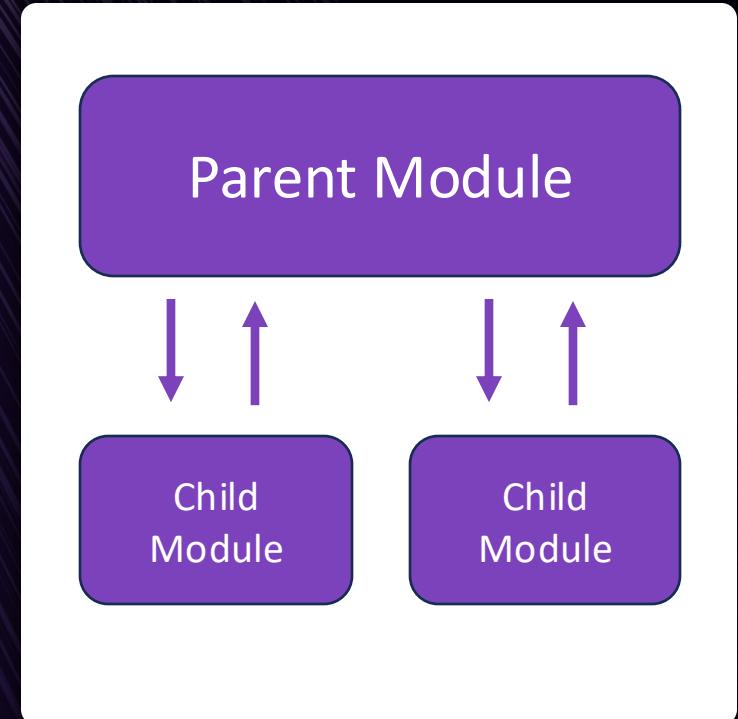
outputs.tf

Module Block



Parent Module vs. Child Module

- A parent (calling) module references and configures other modules, while a child module is the reusable infrastructure component being called.
- Modules can be retrieved from a registry, a repo, or stored locally.





Parent/Root/Calling Module

```
1 module "vpc" { ←
2   source = "terraform-aws-modules/vpc/aws" ←
3
4   name = "my-vpc"
5   cidr = "10.0.0.0/16"
6
7   azs          = data.aws_availability_zone.east.name_suffix
8   private_subnets = var.private_subnets
9   public_subnets  = var.public_subnets
10
11  enable_nat_gateway = true
12  enable_vpn_gateway = true
13
14  tags = {
15    CreatedBy    = "terraform"
16    Environment = "prod"
17  }
18 }
```



Module Blocks Call Child Modules

Defines that this module is the calling module of the child module



Must Define Where the Module Lives

Path to local directory or remote module source to download from



Variables Defined in Root Module

Values must be passed to the module to be used



Very Simple Child Module

```
1 variable "cidr_block" {  
2   type    = string  
3   default = "10.0.0.0/16"  
4 }  
5  
6 variable "name" {  
7   type    = string  
8   default = ""  
9 }  
10  
11 resource "aws_vpc" "vpc" {  
12   name      = var.name  
13   cidr_block = var.cidr_block  
14  
15   tags = {  
16     Name = var.name  
17   }  
18 }  
19  
20 output "vpc_id" {  
21   value = aws_vpc.vpc.id  
22 }
```

Modules Should Look Familiar



Modules are just regular Terraform configurations that are variablized



Must Define One or More Resources

The child module is where the resource blocks are defined



Child Modules Have Outputs

Outputs allow you to pass information from one module to another



Child Module

Root Module

```
1 variable "cidr_block" {  
2   type    = string  
3   default = "192.168.0.0/16"  
4 }  
5  
6 module "prod_vpc" {  
7   source = "terraform-aws-modules/vpc/aws"  
8  
9   name = "my-vpc"  
10  cidr = var.cidr_block  
11 }  
12  
13 output "vpc_id" {  
14   module.prod_vpc.vpc_id  
15 }
```

Values we want to pass to
our module

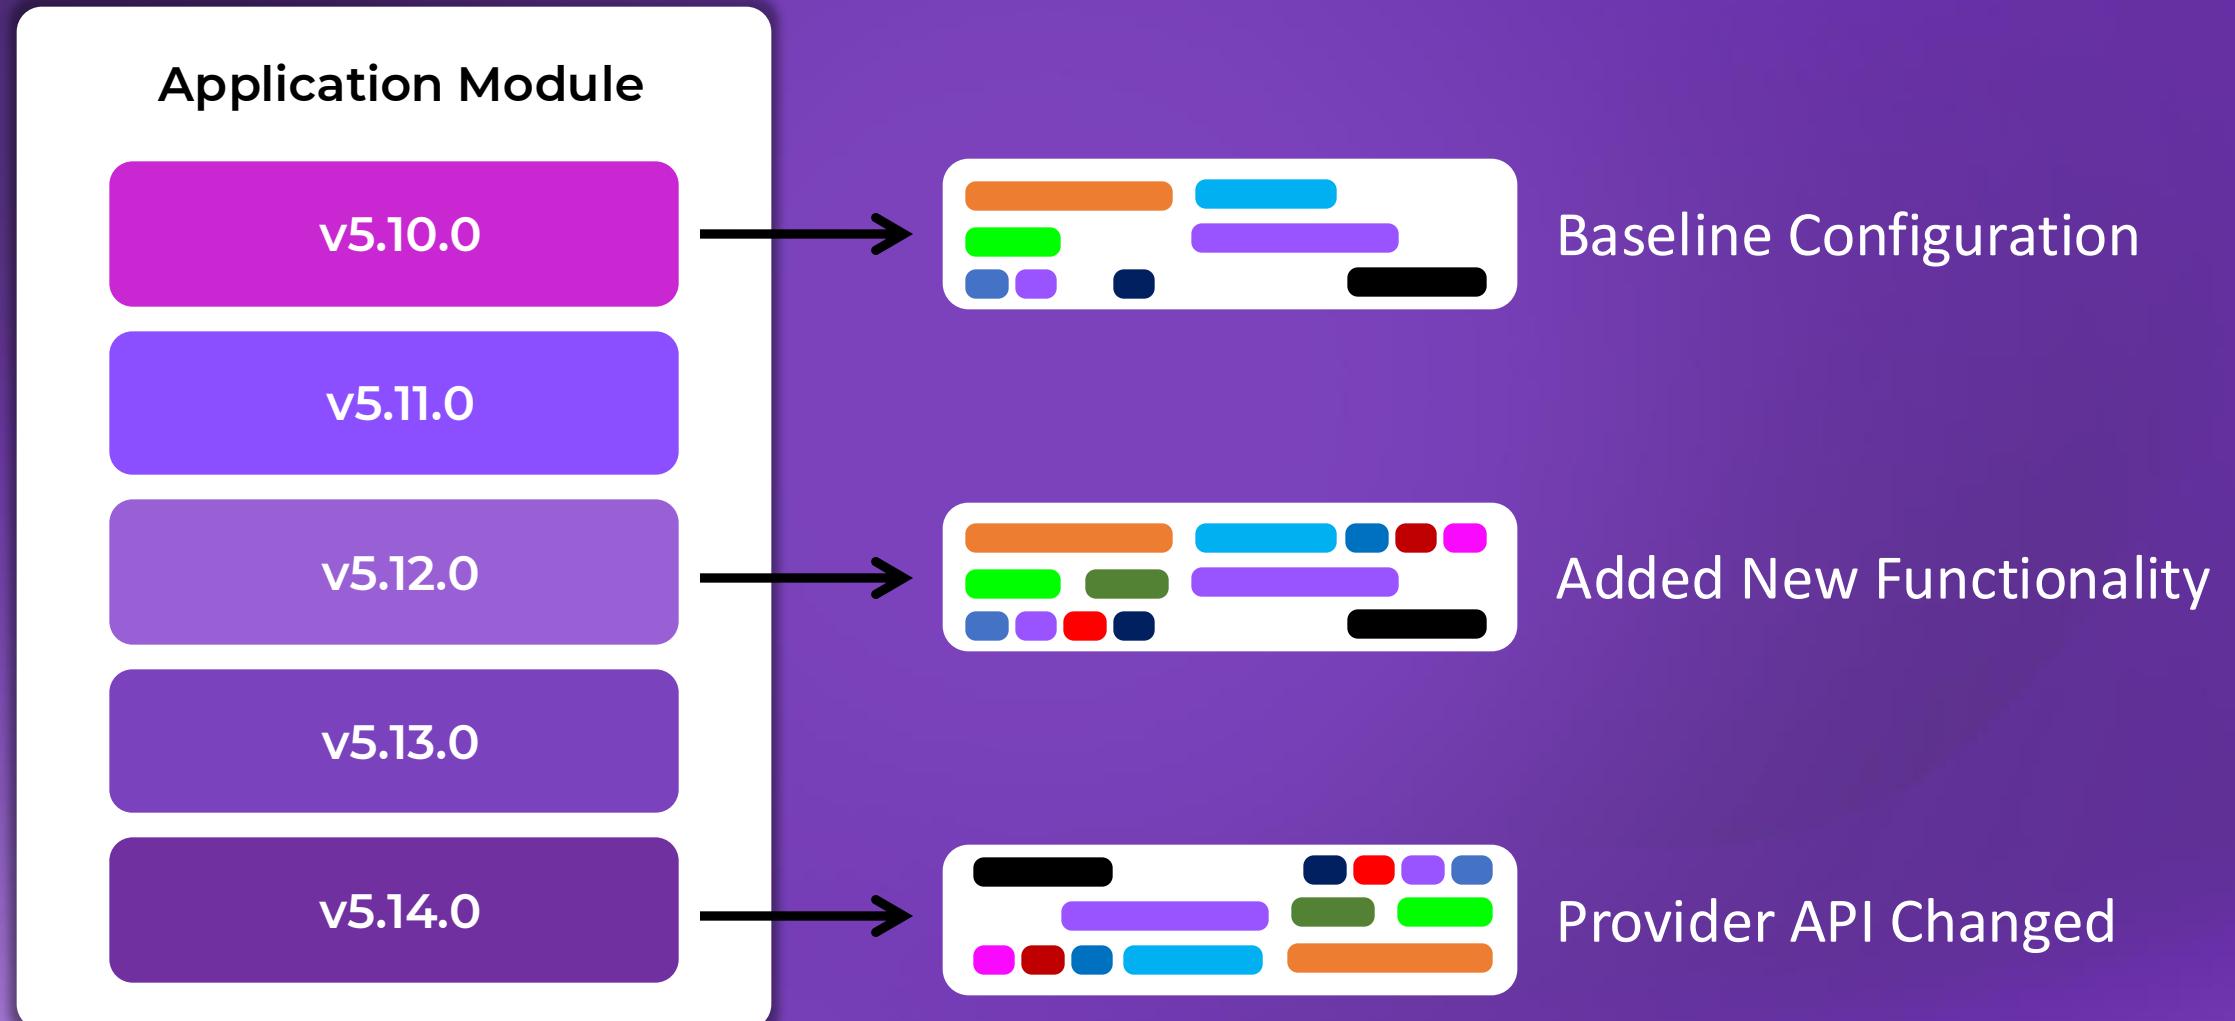
Export values to be accessed
by calling the module

```
1 variable "cidr" {  
2   type    = string  
3   default = "10.0.0.0/16"  
4 }  
5  
6 variable "name" {  
7   type    = string  
8   default = ""  
9 }  
10  
11 resource "aws_vpc" "vpc" {  
12   name      = var.name  
13   cidr_block = var.cidr  
14  
15   tags = {  
16     Name = var.name  
17   }  
18 }  
19  
20 output "vpc_id" {  
21   value = aws_vpc.vpc.id  
22 }
```





Module Versioning



Pinning to a Specific Module Version

```
● ● ●  
1 variable "cidr_block" {  
2   type    = string  
3   default = "192.168.0.0/16"  
4 }  
5  
6 module "prod_vpc" {  
7   source = "terraform-aws-modules/vpc/aws"  
8   version = "5.12.0"  
9  
10  name = "my-vpc"  
11  cidr = var.cidr_block  
12 }  
13  
14 output "vpc_id" {  
15   module.prod_vpc.vpc_id  
16 }
```



SUMMARY



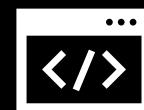
HCL is a declarative language that is easy to read and write, and is Terraform's primary interface.



A single Terraform file can have one or many blocks to define your infrastructure



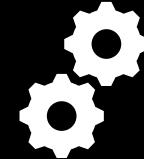
HCL uses blocks for data, resources, etc., and these blocks can be nested.



Use the '.tf' file extension so Terraform and related tools can recognize and apply configurations properly.



Single-line and block comments are used in HCL.



Use linters, extensions, and built-in tools like 'terraform fmt' to enhance code quality.